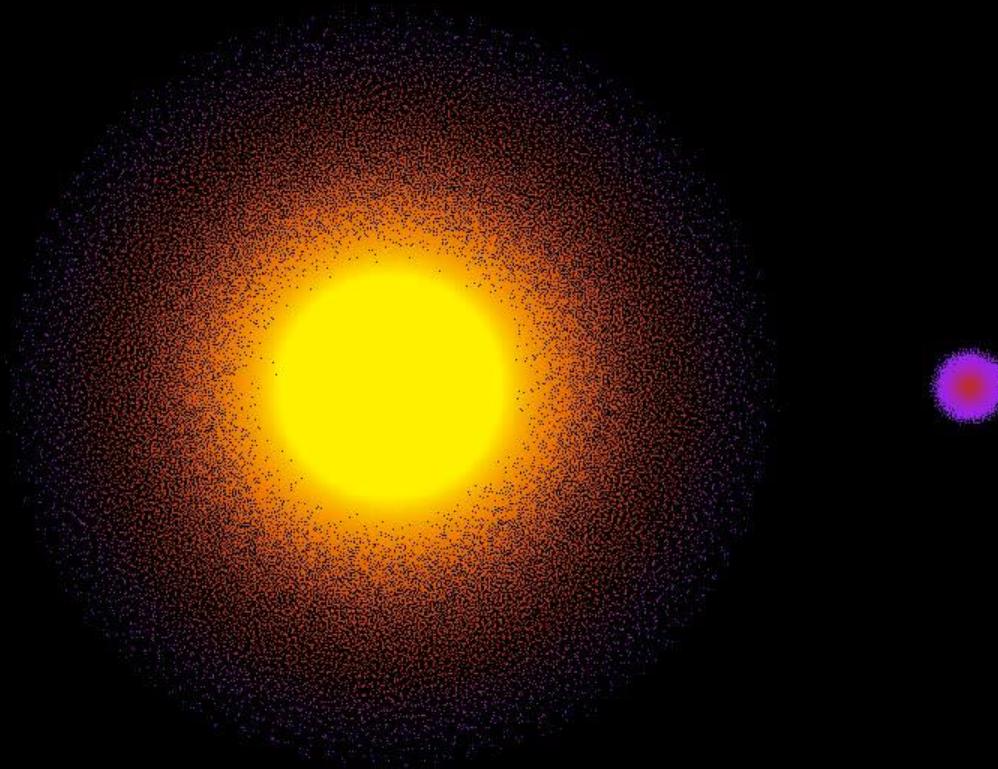
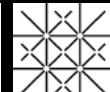


# Introduction to Parallel Computing



Dr. Rubén M. Cabezón  
Dr. Aurélien Cavelan

**sciCORE**  
Center for scientific computing



University  
of Basel

OpenMP

# Outline

Introduction

OpenMP

- Scheduling
- Variables scope
- Reduction
- Atomic
- Collapse
- Orphaned directives

Python

- Multiprocessing
- Numba

OpenMP and GPU offloading

MPI (just a smidge)

Launching in clusters (SLURM)



09:00 – 10:30: Lecture

10:30 – 10:45: Coffee break

10:45 – 12:00: Lecture

12:00 – 13:30: Lunch

13:30 – 15:00: Lecture

15:00 – 15:15: Coffee break

15:15 – 17:00: Lecture



# First thing to do:

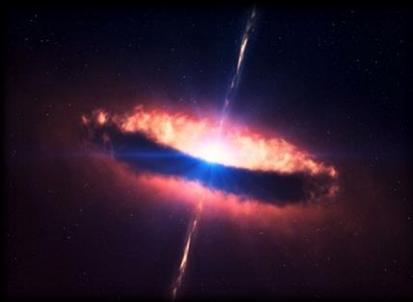
Download the slides of the course:

From the browser: [https://bit.ly/scicore\\_openmp](https://bit.ly/scicore_openmp)

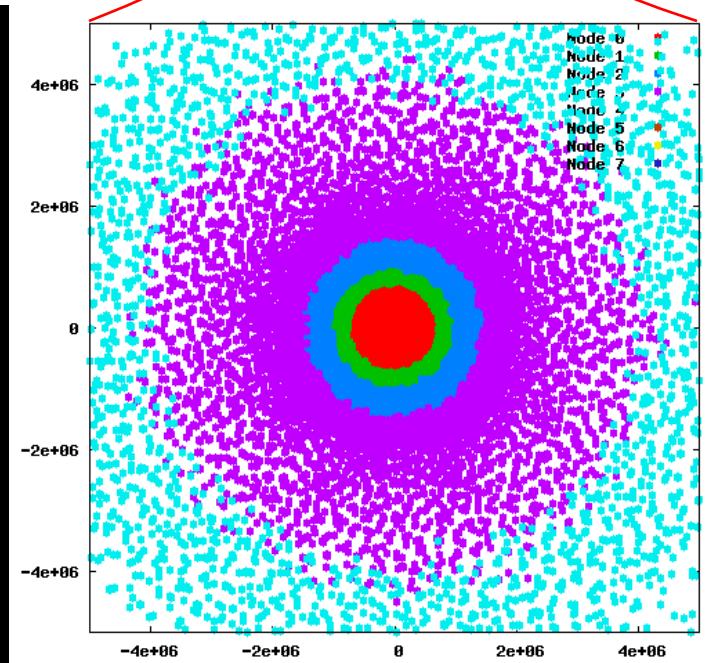
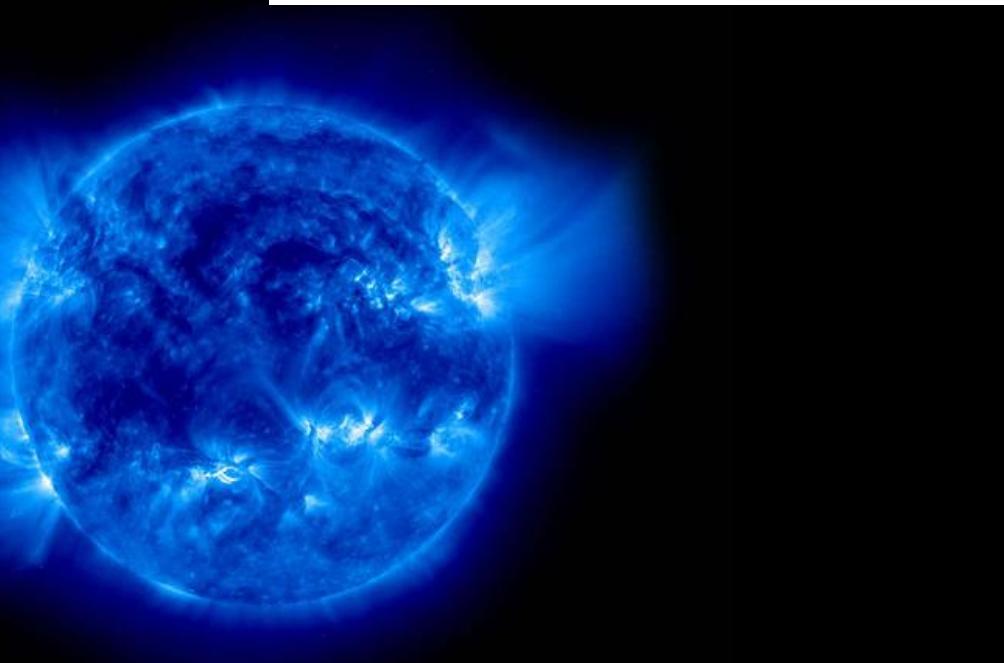
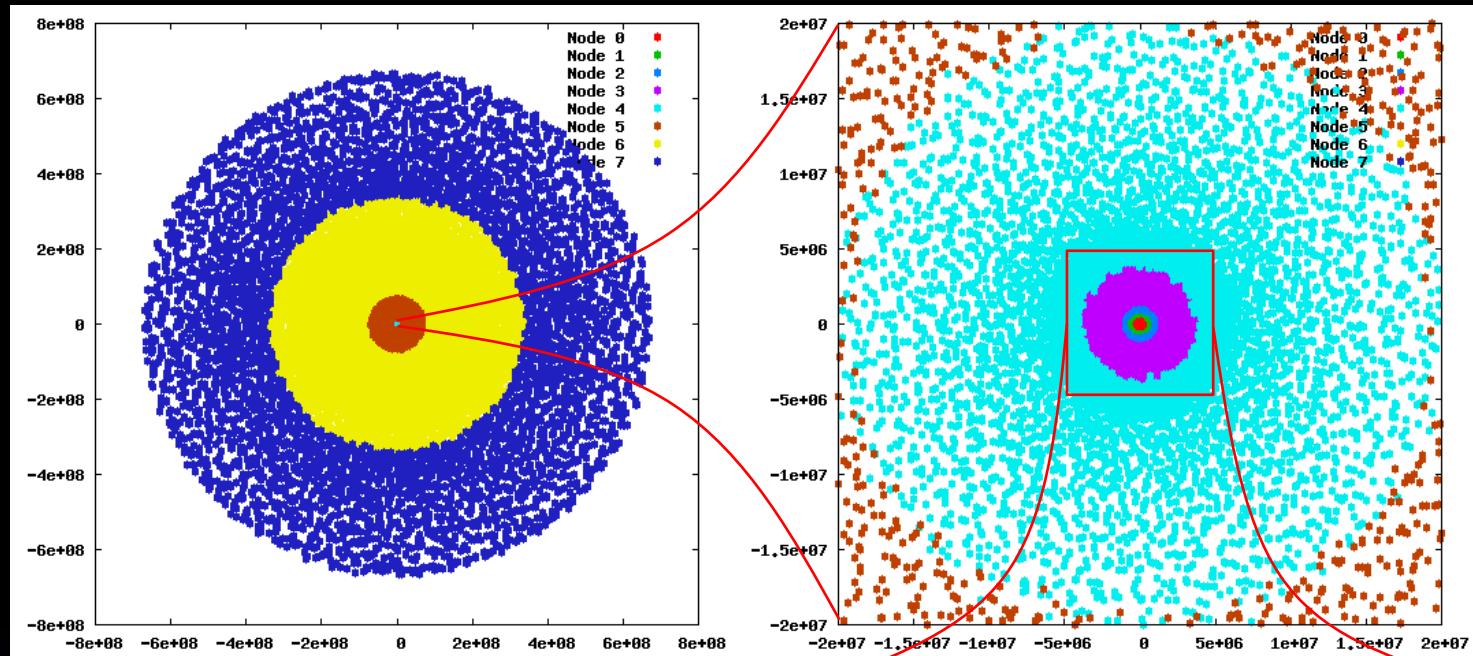
From the terminal: `wget https://bit.ly/scicore_openmp`

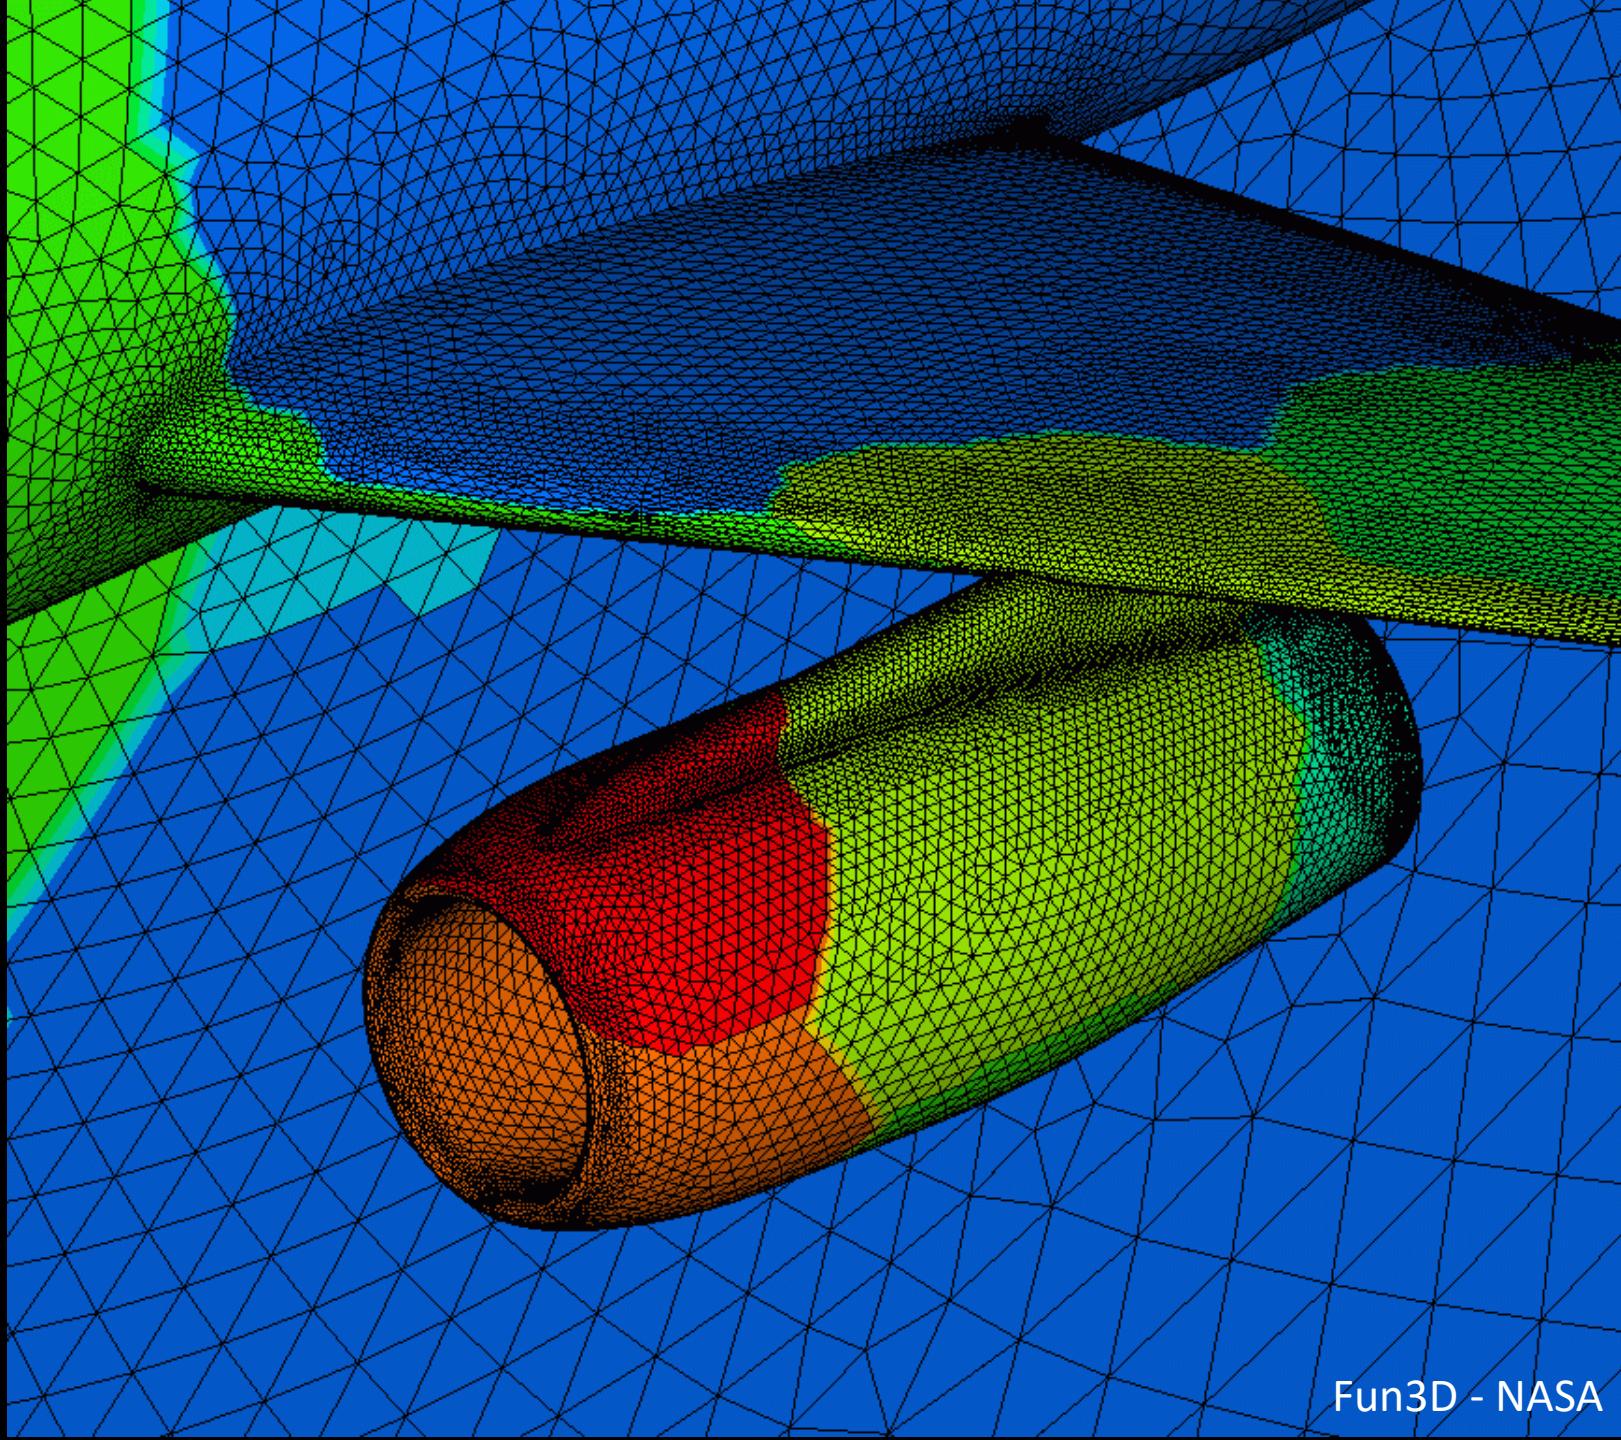
You can find there the slides of the course (exercises are redacted)

Multiple executions of the same code, applied to different sections of the input data.



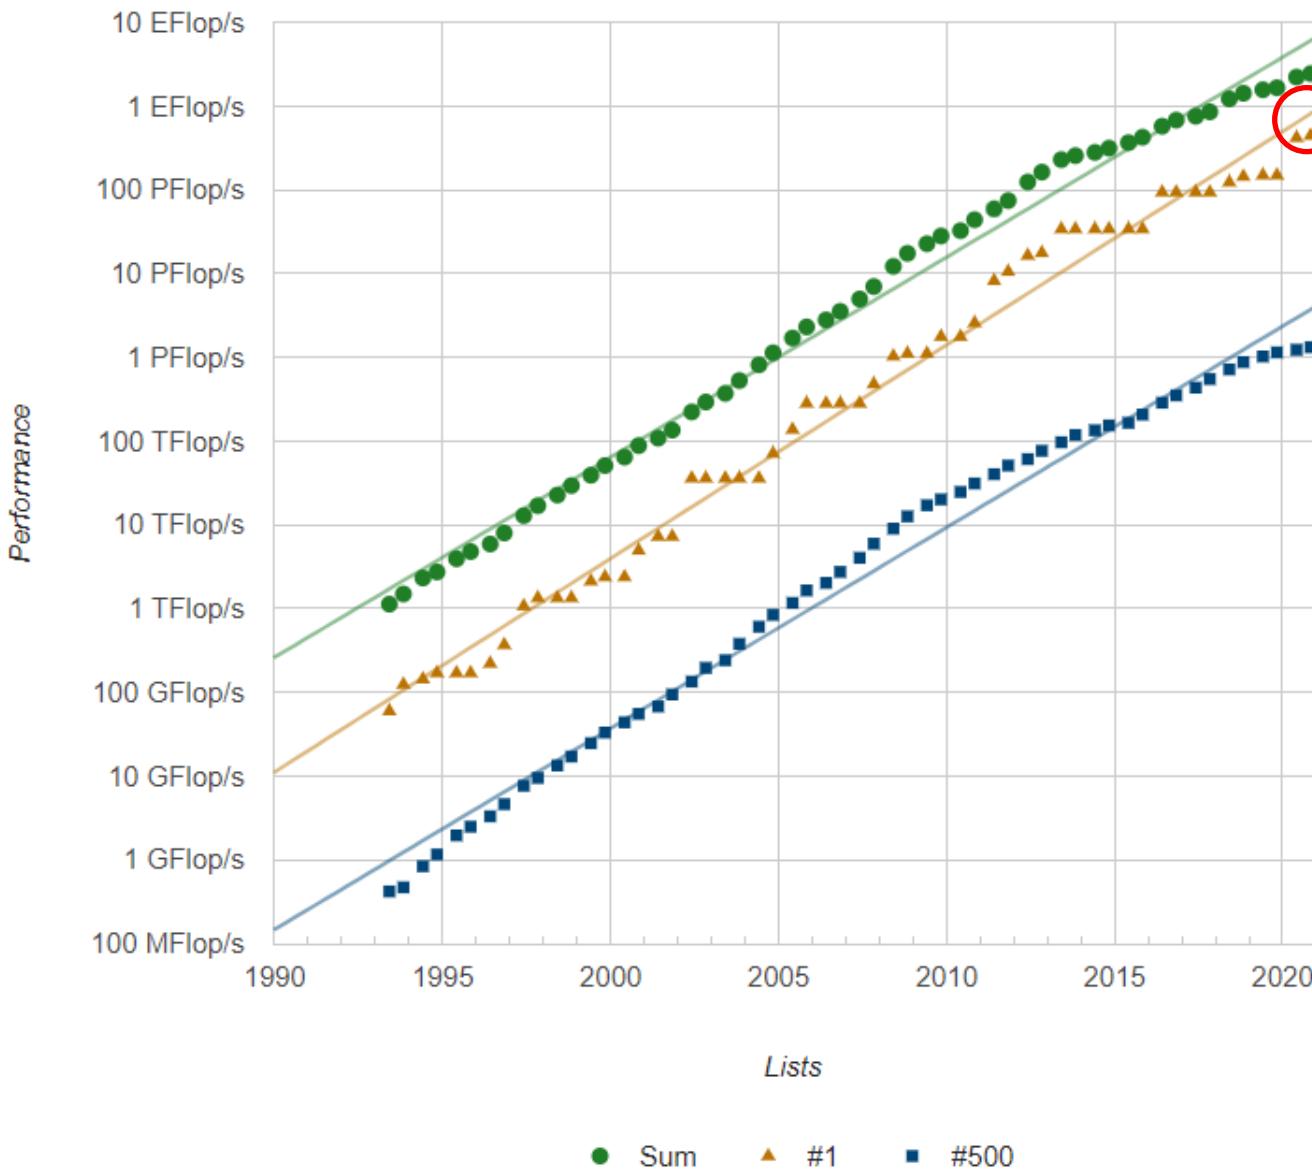
Split the problem in subdomains that are partially solved independently.





Fun3D - NASA

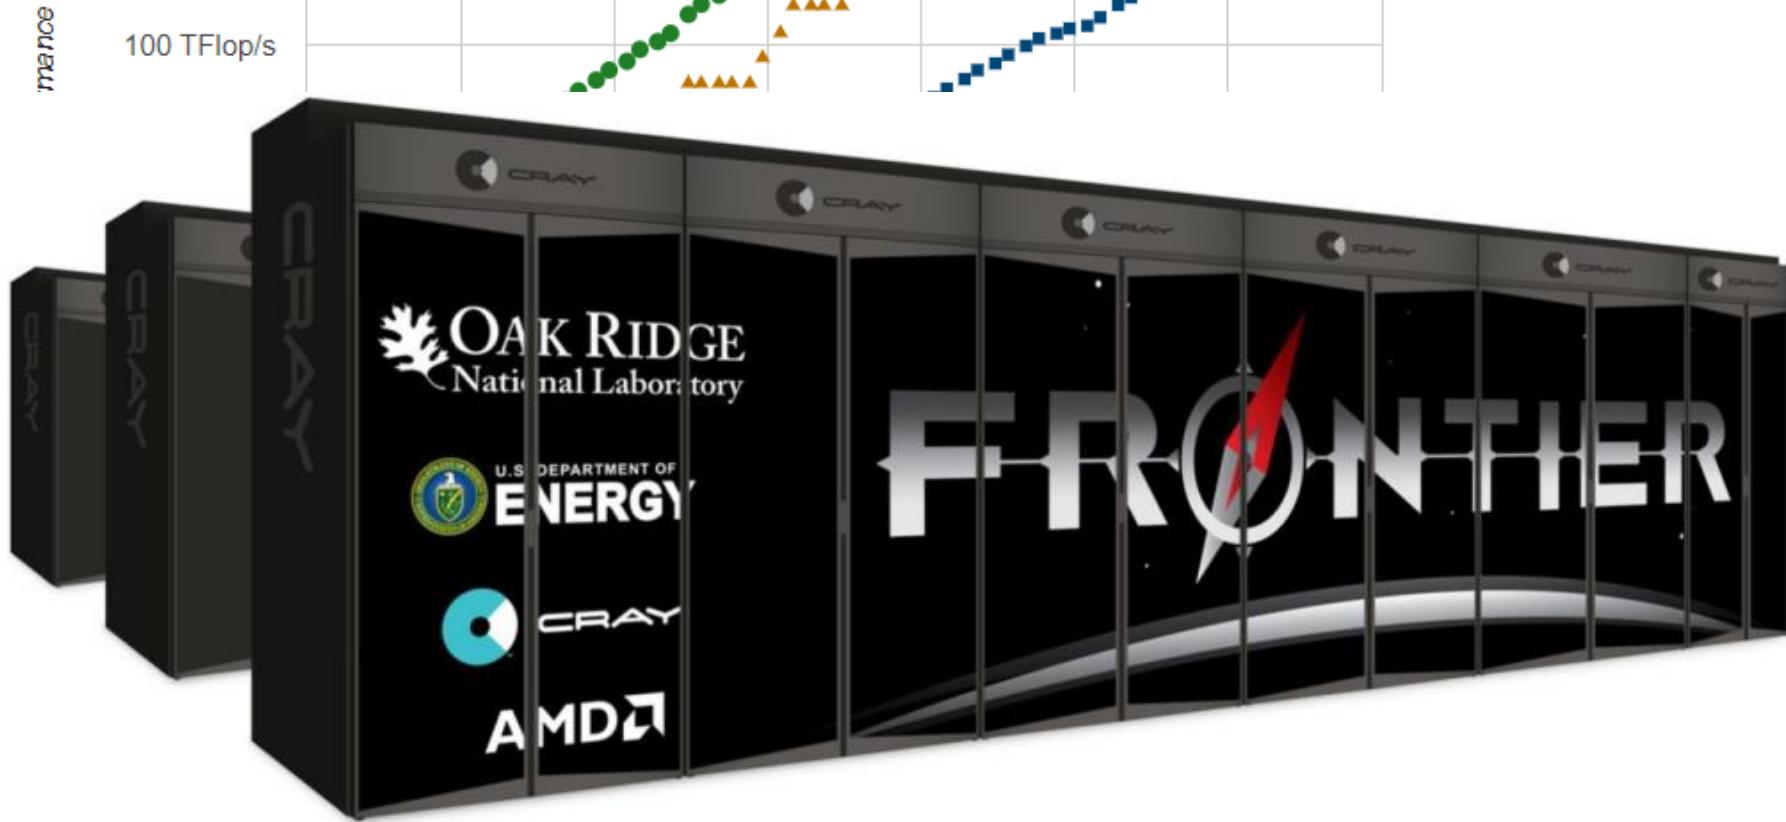
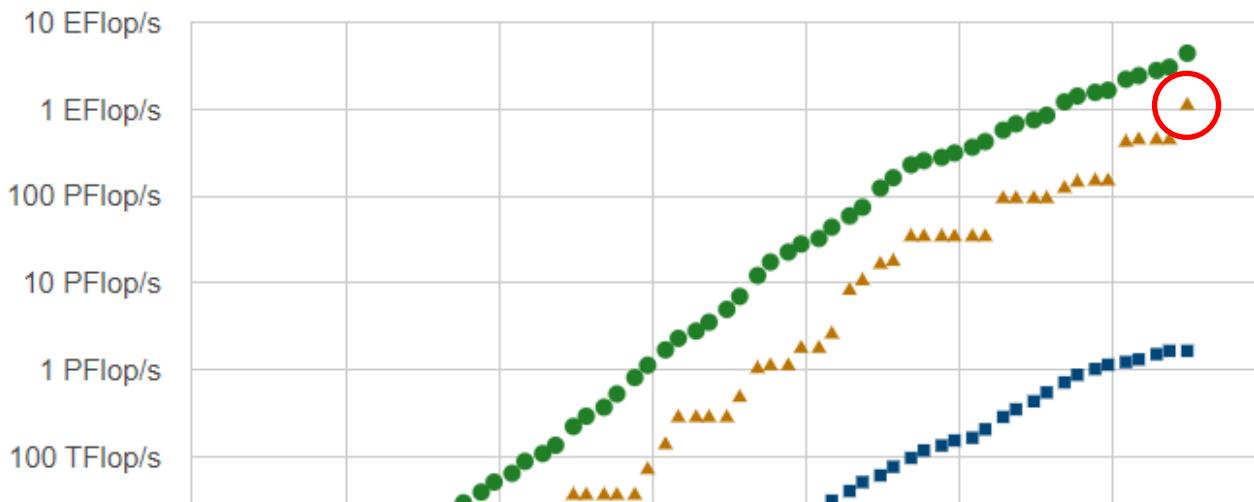
## Projected Performance Development

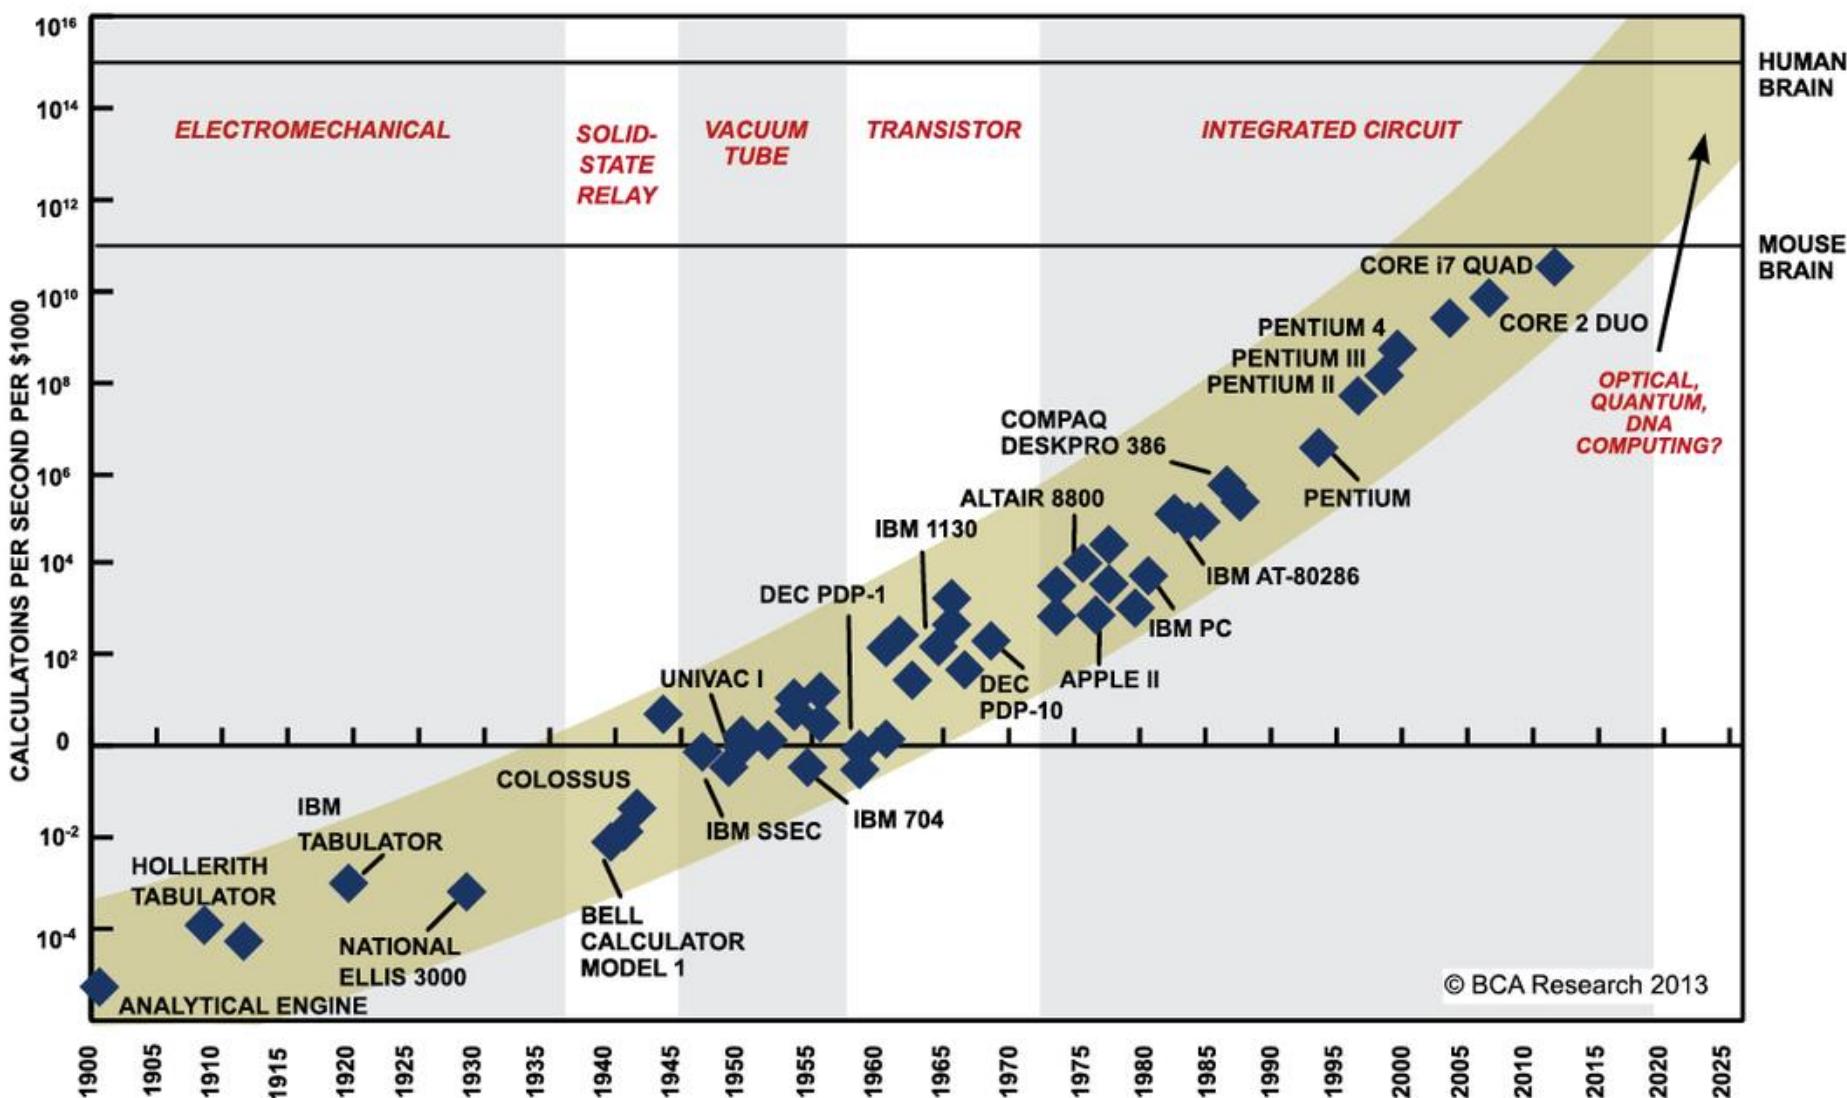


Exascale will be reached close to 2022

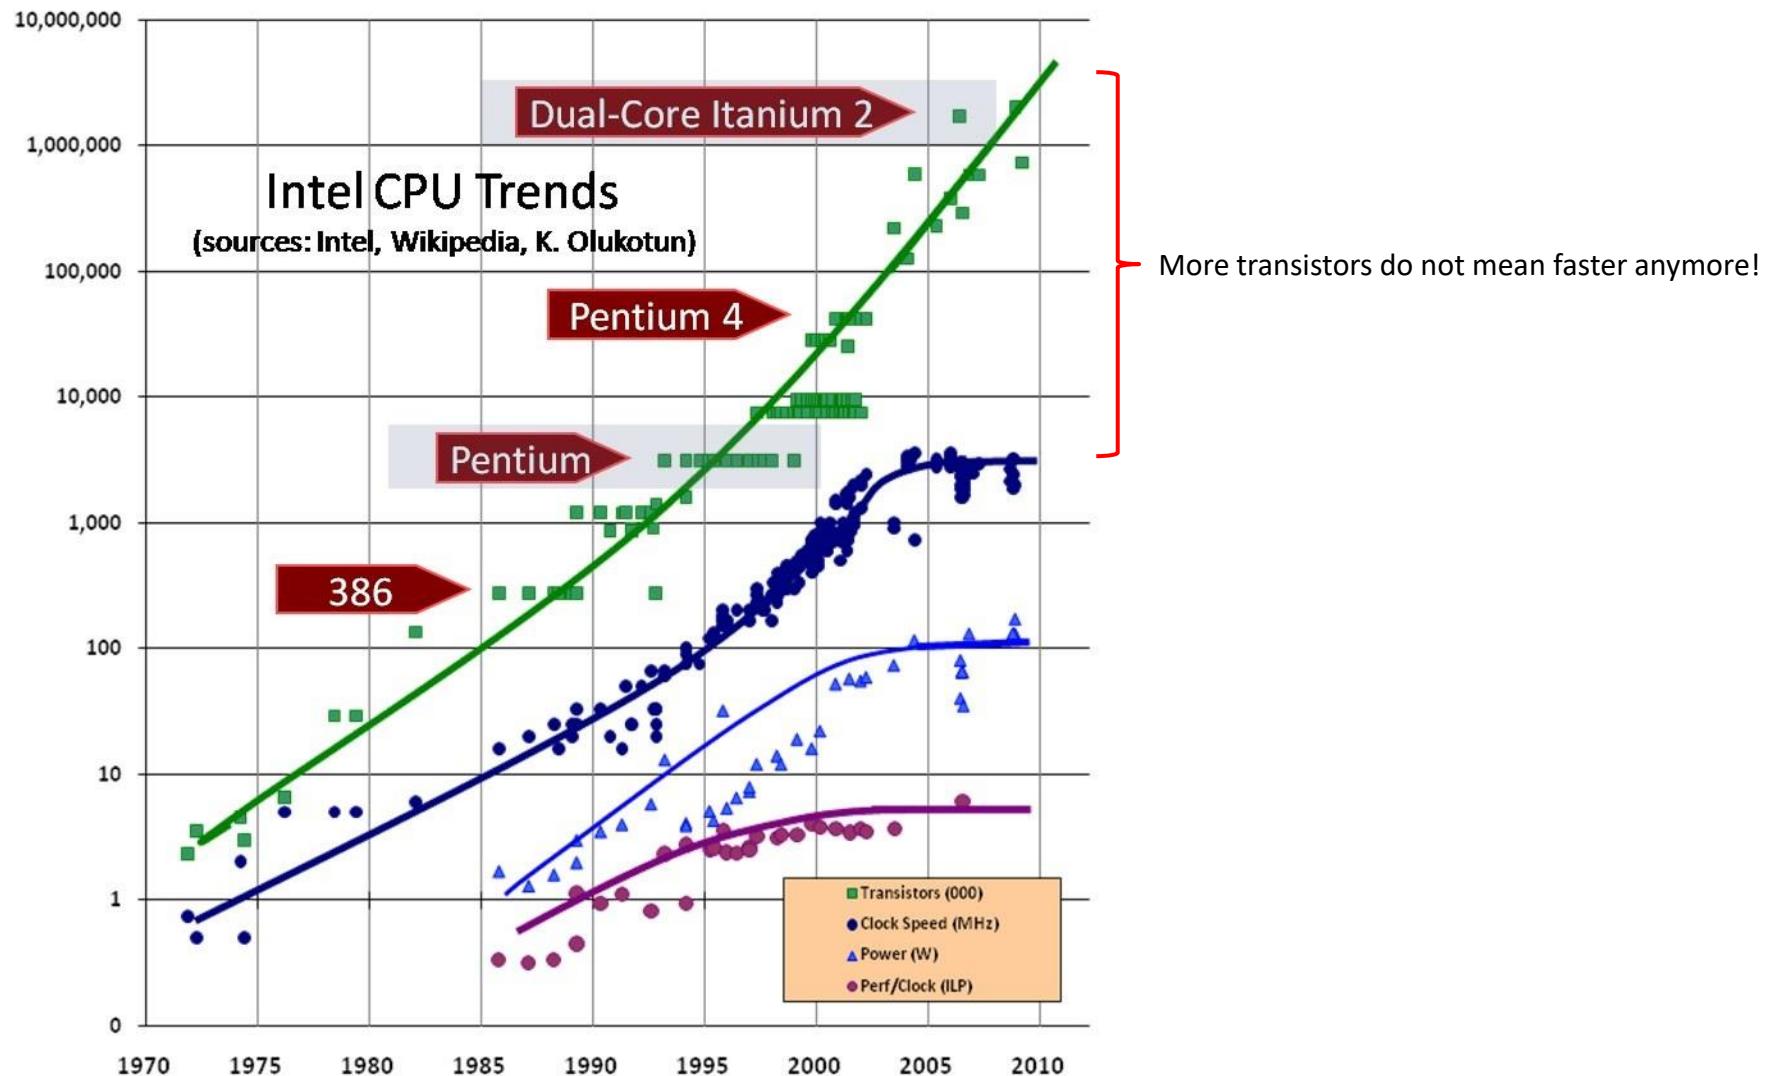
- Aerospace, Airframes, Jet Turbines
- Astrophysics
- Biological and medical systems
- Climate and weather
- Combustion
- Materials science
- Fusion energy
- National security
- Nuclear engineering

## Performance Development

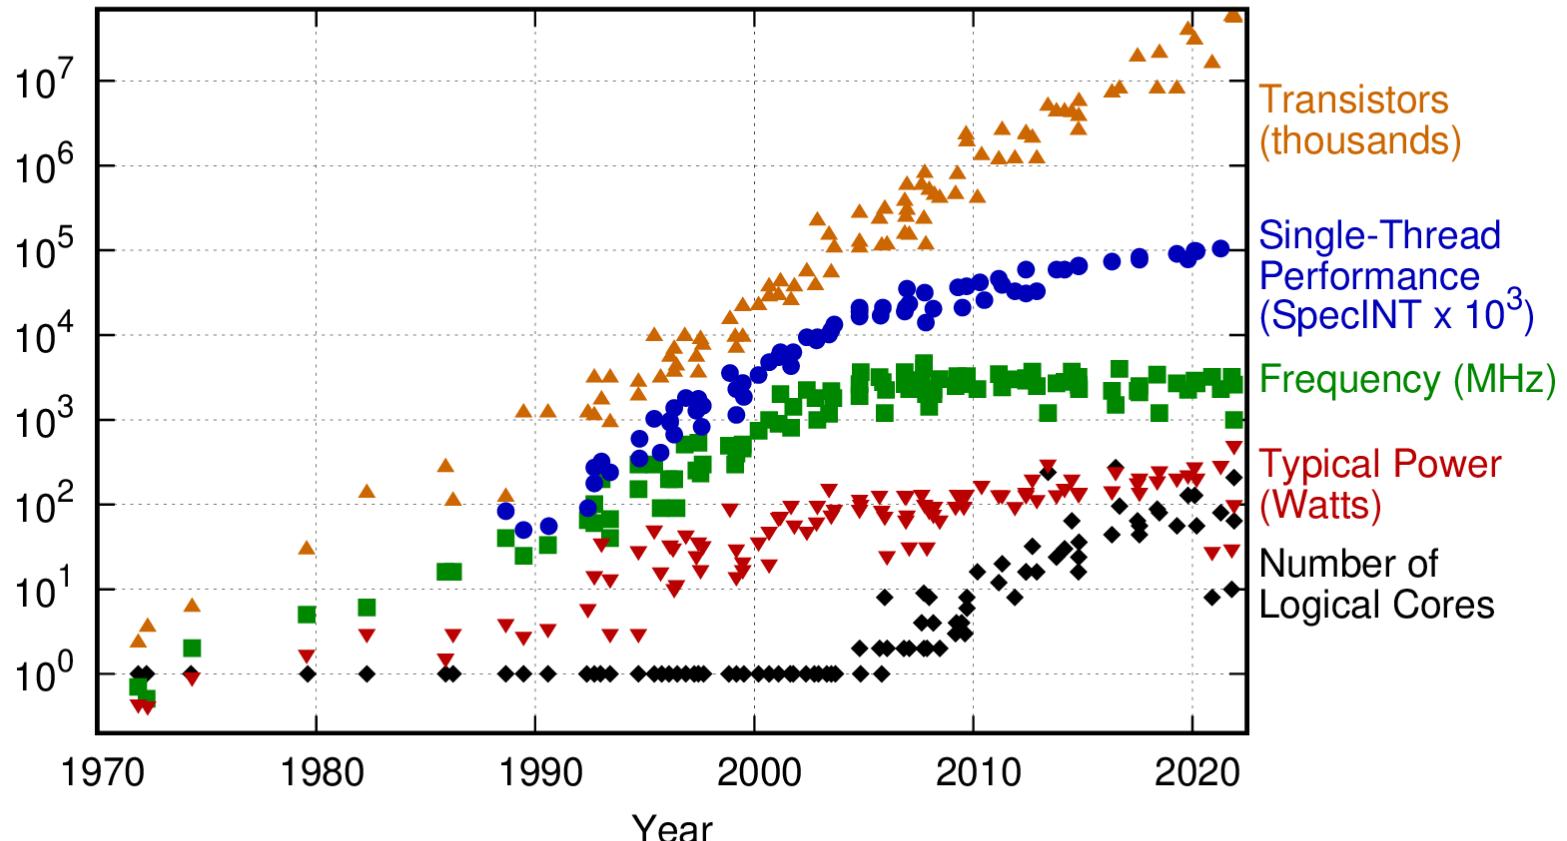




SOURCE: RAY KURZWEIL, "THE SINGULARITY IS NEAR: WHEN HUMANS TRANSCEND BIOLOGY", P.67, THE VIKING PRESS, 2006. DATAPoints BETWEEN 2000 AND 2012 REPRESENT BCA ESTIMATES.



## 50 Years of Microprocessor Trend Data

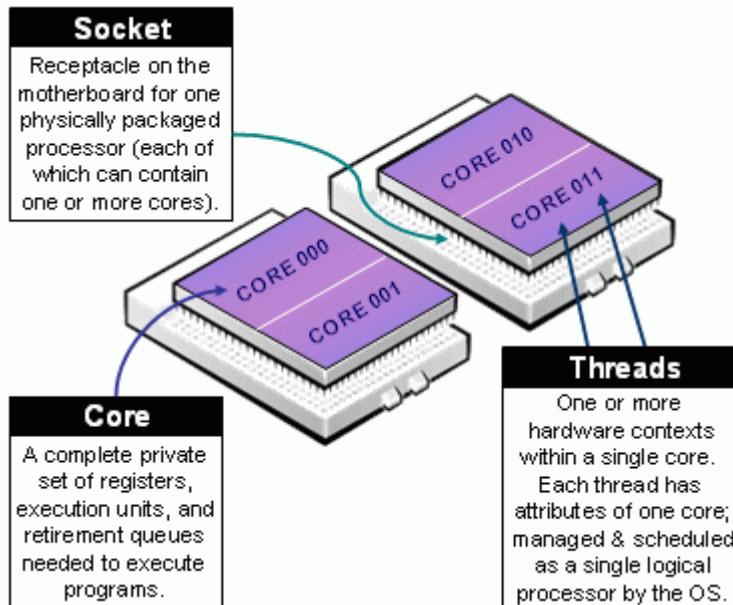


Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten  
New plot and data collected for 2010-2021 by K. Rupp

Increase in performance is done nowadays via  
**parallelization**

Usually a hybrid MPI+X approach is taken

(X=OpenMP, OpenMP offloading, OpenACC, CUDA)



Hyperthreading is usually deactivated  
in HPC facilities by default.  
Ask the sysadmin!

# logical cores = # physical cores \* # threads/core

node 0	process 0	process 1
	thr0 thr1 thr2 thr3 thr4 thr5 thr6 thr7	thr0 thr1 thr2 thr3 thr4 thr5 thr6 thr7
node 1	process 2	process 3
	thr0 thr1 thr2 thr3 thr4 thr5 thr6 thr7	thr0 thr1 thr2 thr3 thr4 thr5 thr6 thr7
node 2	process 4	process 5
	thr0 thr1 thr2 thr3 thr4 thr5 thr6 thr7	thr0 thr1 thr2 thr3 thr4 thr5 thr6 thr7
node 3	process 6	process 7
	thr0 thr1 thr2 thr3 thr4 thr5 thr6 thr7	thr0 thr1 thr2 thr3 thr4 thr5 thr6 thr7

In general, the problem can be reduced to a logical repetitive structure:

```
do i = 1, n  
<calculations>  
end do
```

Are these <calculations> completely **independent** from each other?

Yes



No

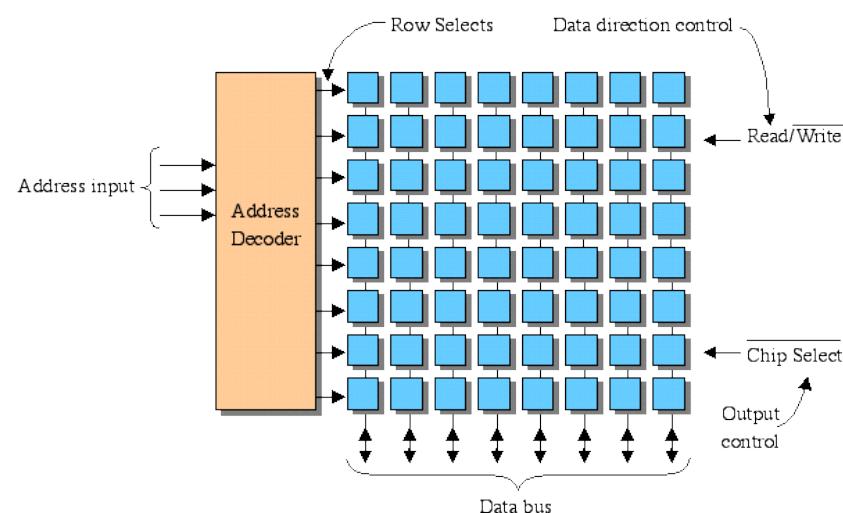
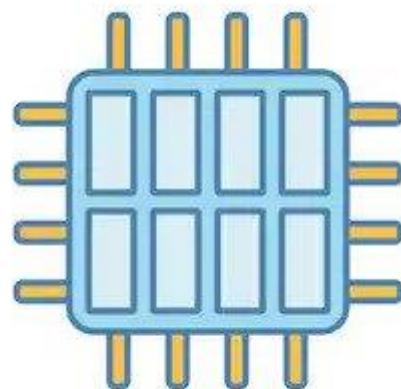
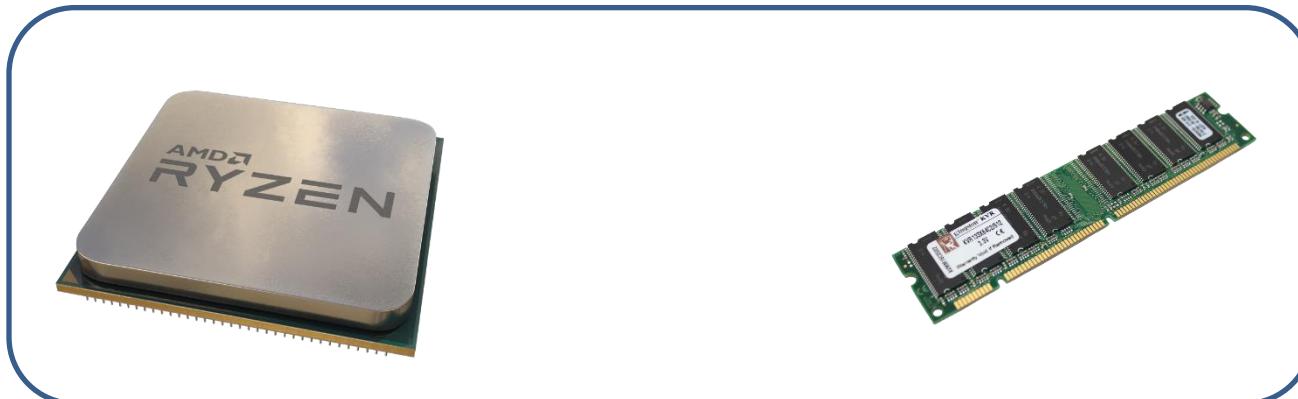
- What's the degree of dependency?
- Can I rewrite/factorize to eliminate the dependency?
- Do variables share contents? Is it necessary?



# OpenMP™

Open Multi-Processing

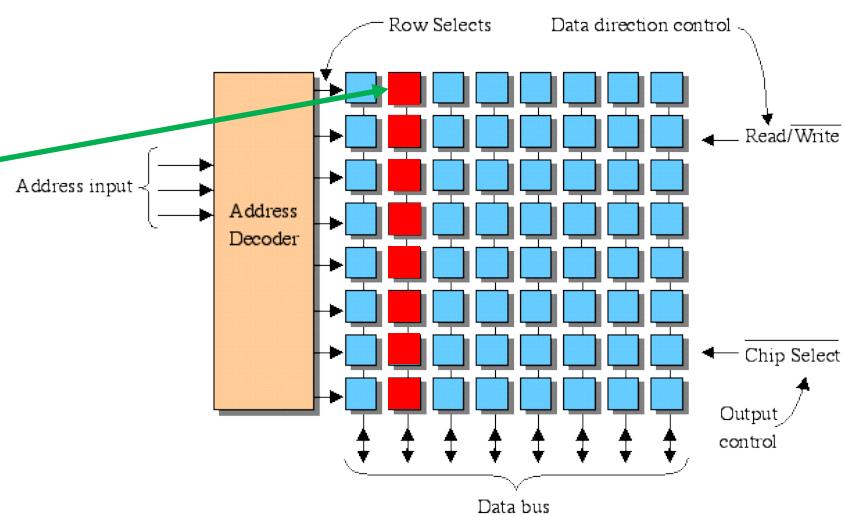
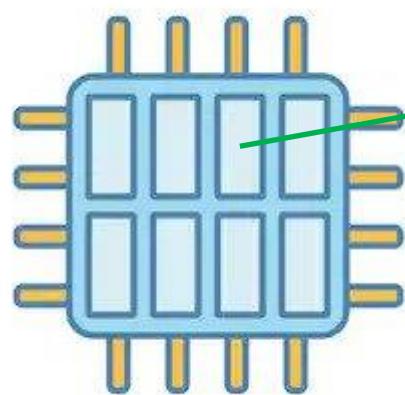
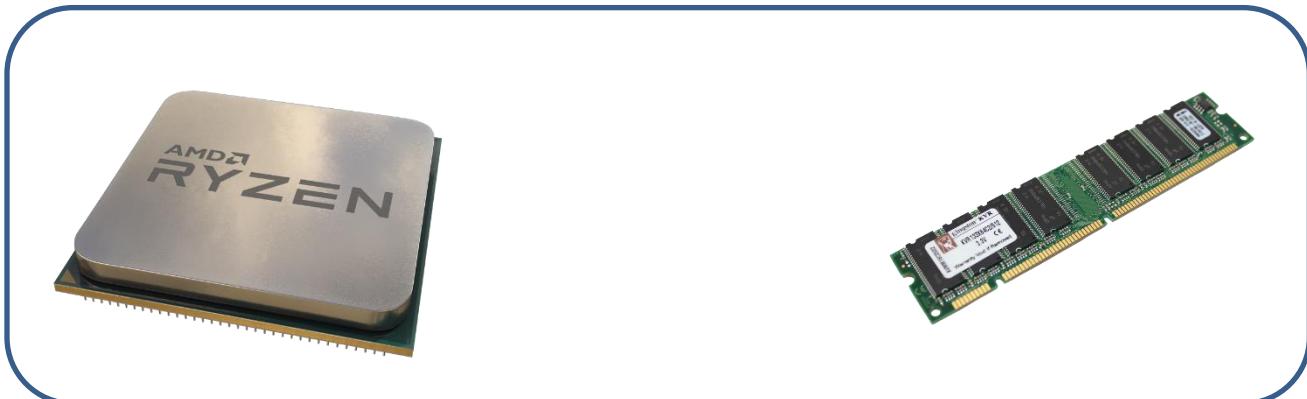
Works in **shared** memory systems  
Can use as many cores as there are in the **compute node**  
Uses **threads** to split the work  
Threads have both **private** and **shared** variables  
Uses **preprocessor directives** in commented lines



# OpenMP™

Open Multi-Processing

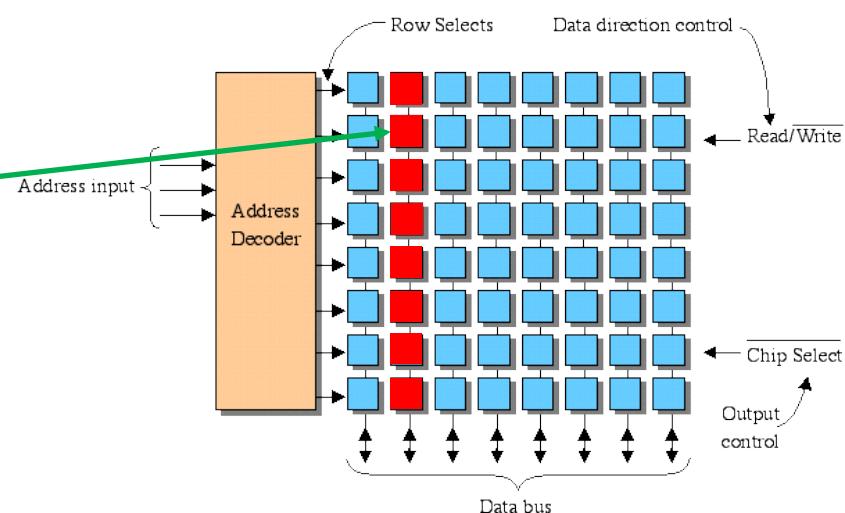
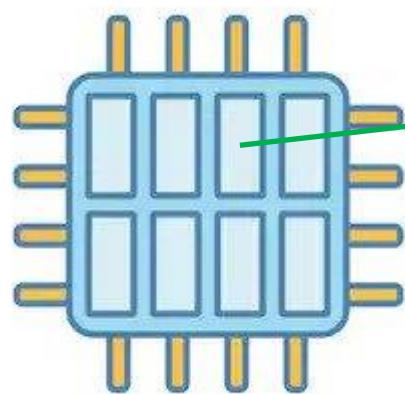
Works in **shared** memory systems  
Can use as many cores as there are in the **compute node**  
Uses **threads** to split the work  
Threads have both **private** and **shared** variables  
Uses **preprocessor directives** in commented lines



# OpenMP™

Open Multi-Processing

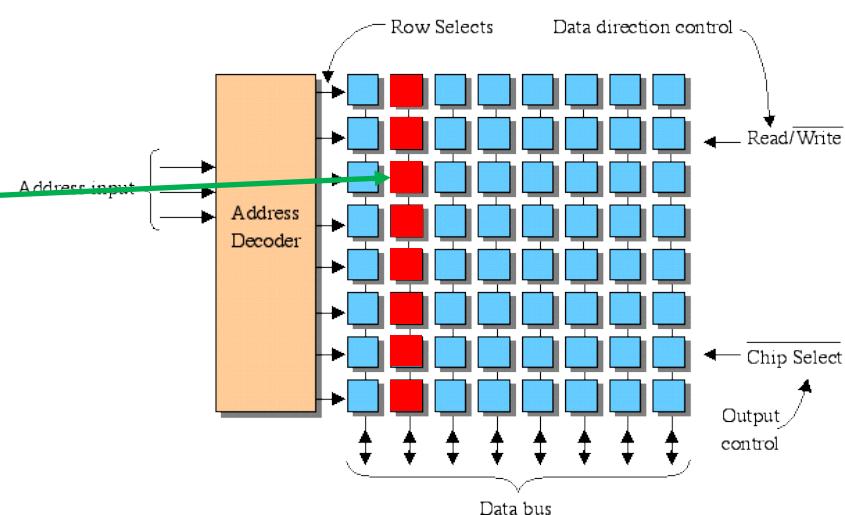
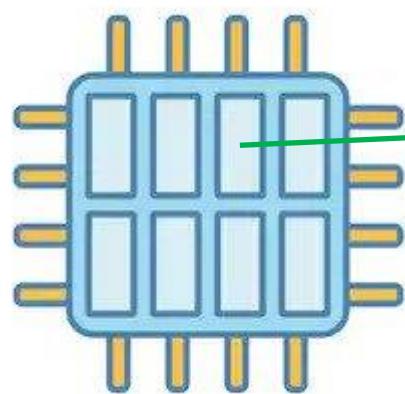
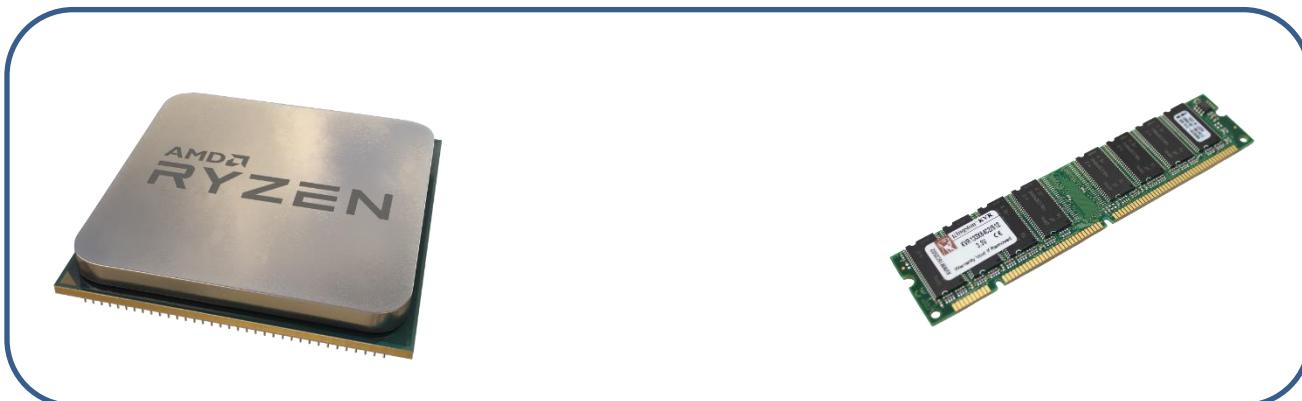
Works in **shared** memory systems  
Can use as many cores as there are in the **compute node**  
Uses **threads** to split the work  
Threads have both **private** and **shared** variables  
Uses **preprocessor directives** in commented lines



# OpenMP™

Open Multi-Processing

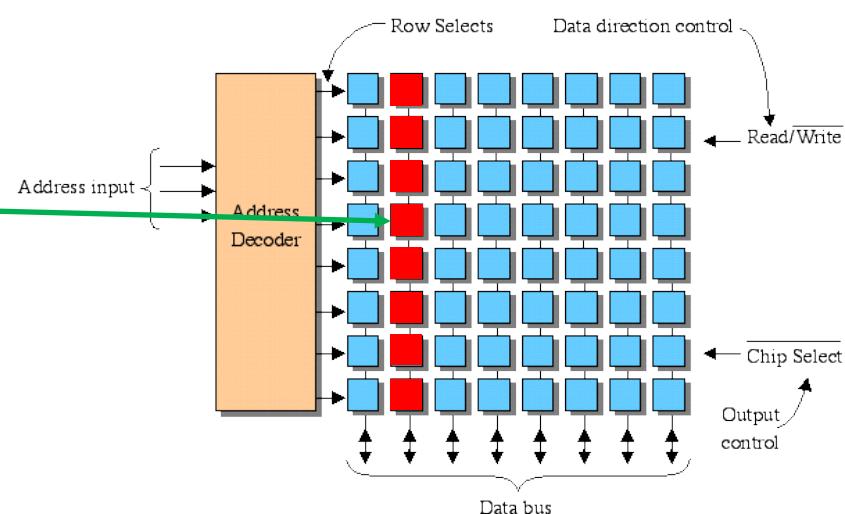
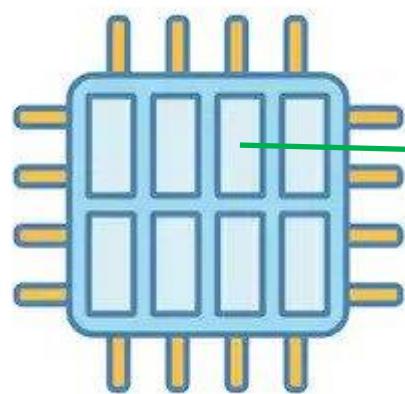
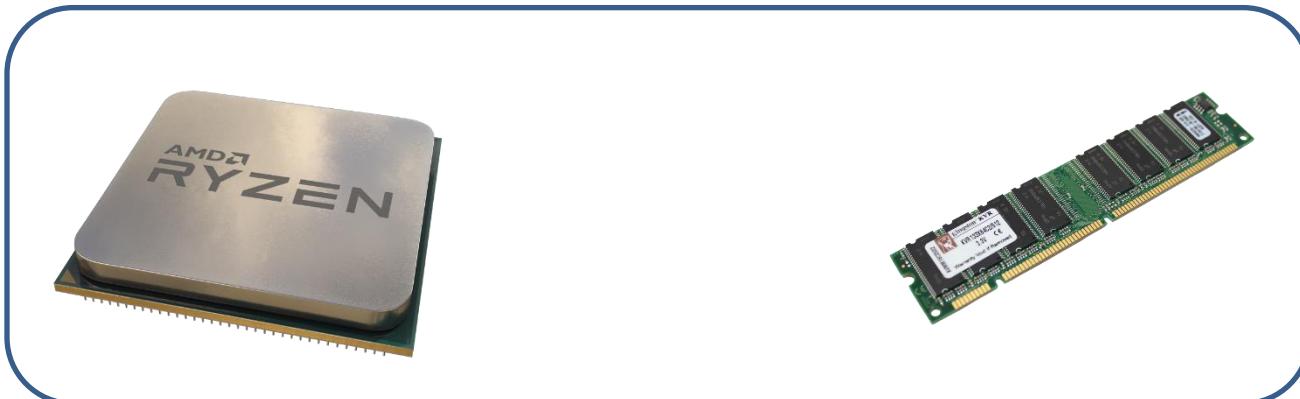
Works in **shared** memory systems  
Can use as many cores as there are in the **compute node**  
Uses **threads** to split the work  
Threads have both **private** and **shared** variables  
Uses **preprocessor directives** in commented lines



# OpenMP™

Open Multi-Processing

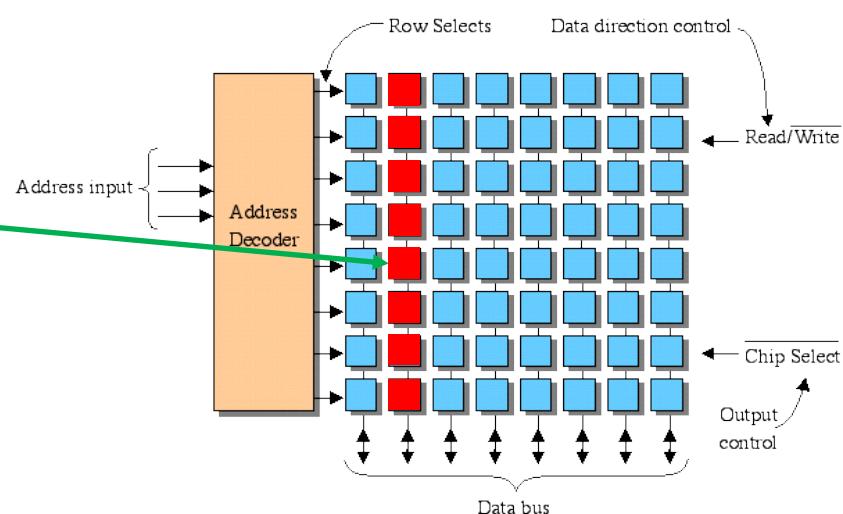
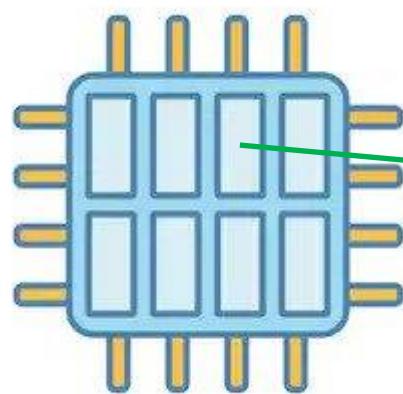
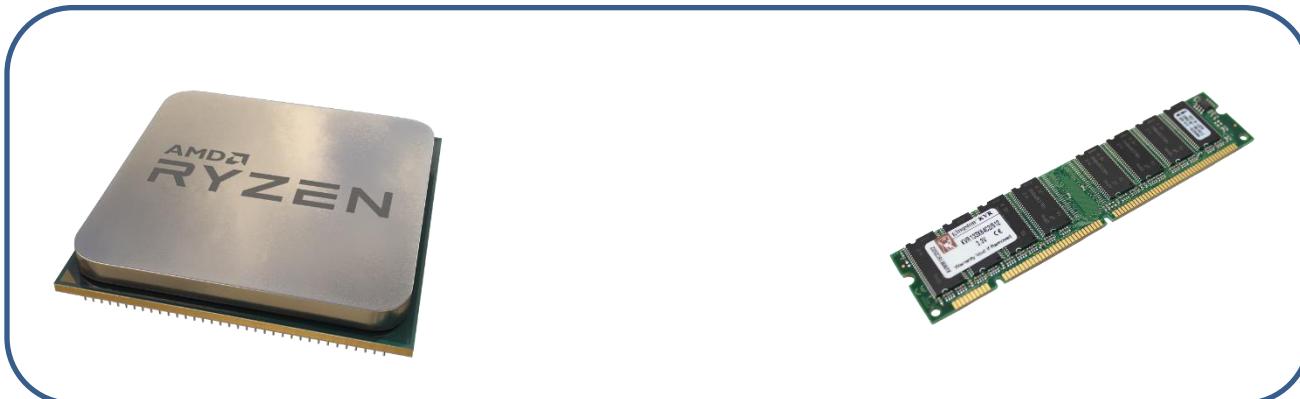
Works in **shared** memory systems  
Can use as many cores as there are in the **compute node**  
Uses **threads** to split the work  
Threads have both **private** and **shared** variables  
Uses **preprocessor directives** in commented lines



# OpenMP™

Open Multi-Processing

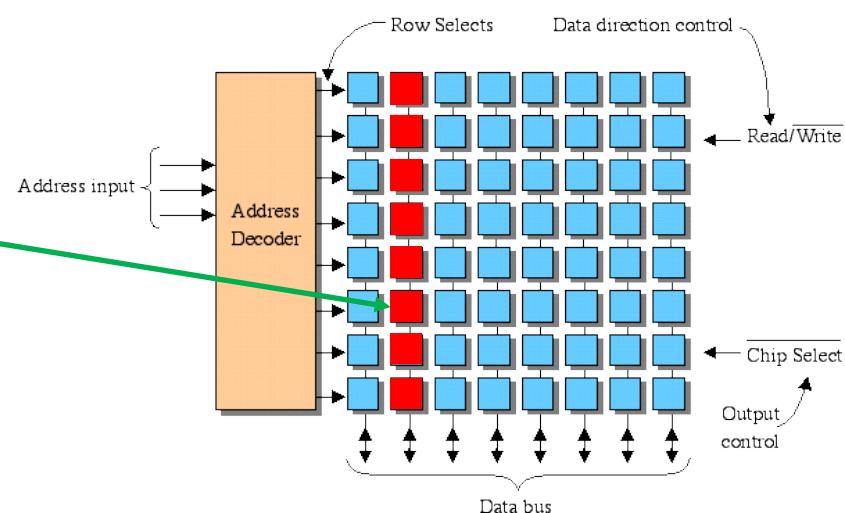
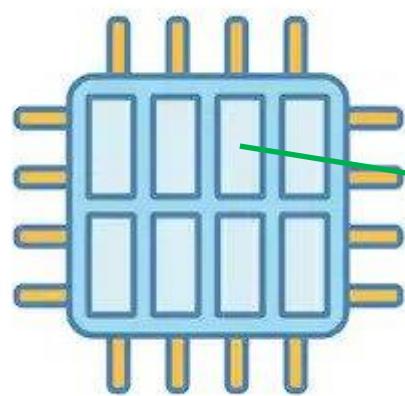
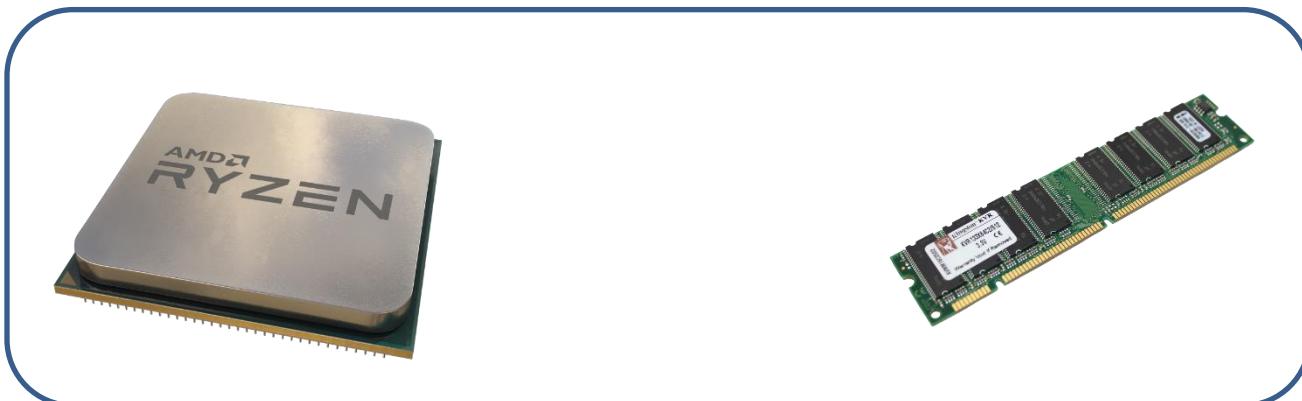
Works in **shared** memory systems  
Can use as many cores as there are in the **compute node**  
Uses **threads** to split the work  
Threads have both **private** and **shared** variables  
Uses **preprocessor directives** in commented lines



# OpenMP™

Open Multi-Processing

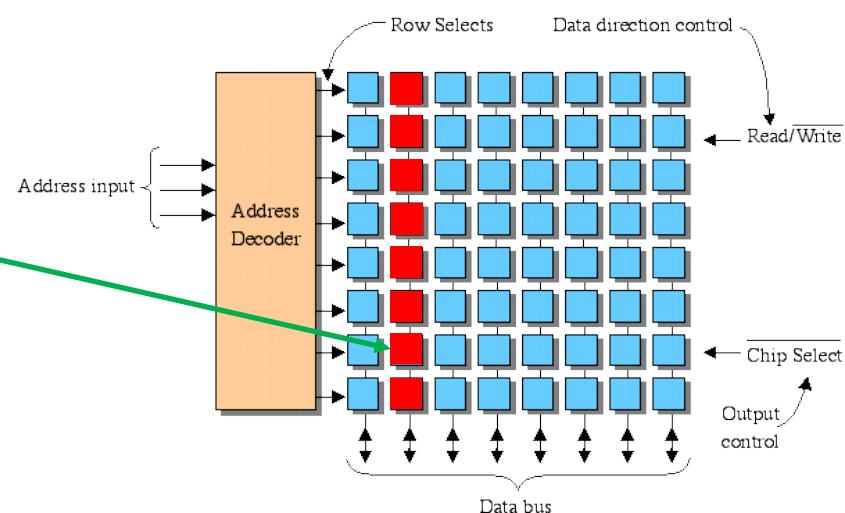
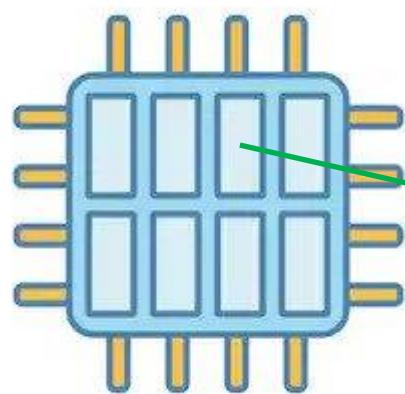
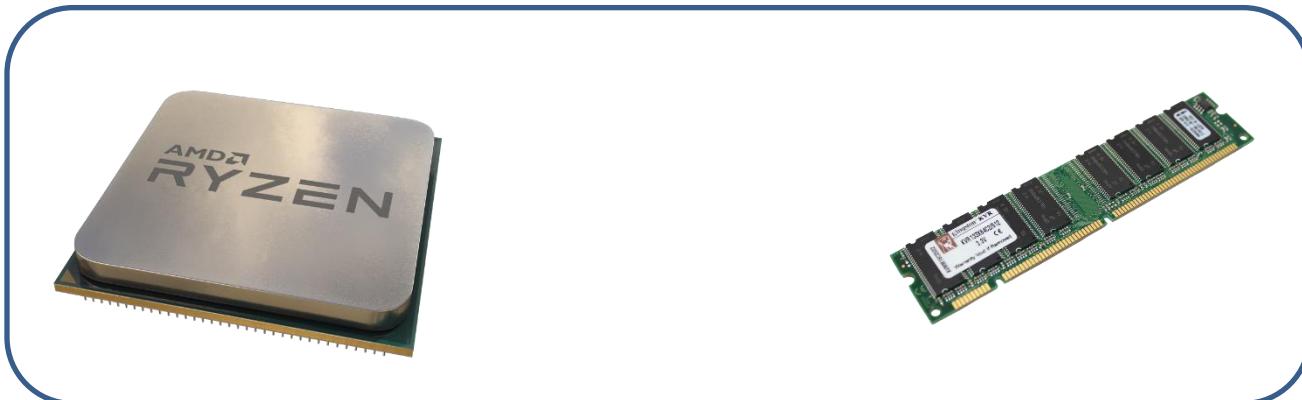
Works in **shared** memory systems  
Can use as many cores as there are in the **compute node**  
Uses **threads** to split the work  
Threads have both **private** and **shared** variables  
Uses **preprocessor directives** in commented lines



# OpenMP™

Open Multi-Processing

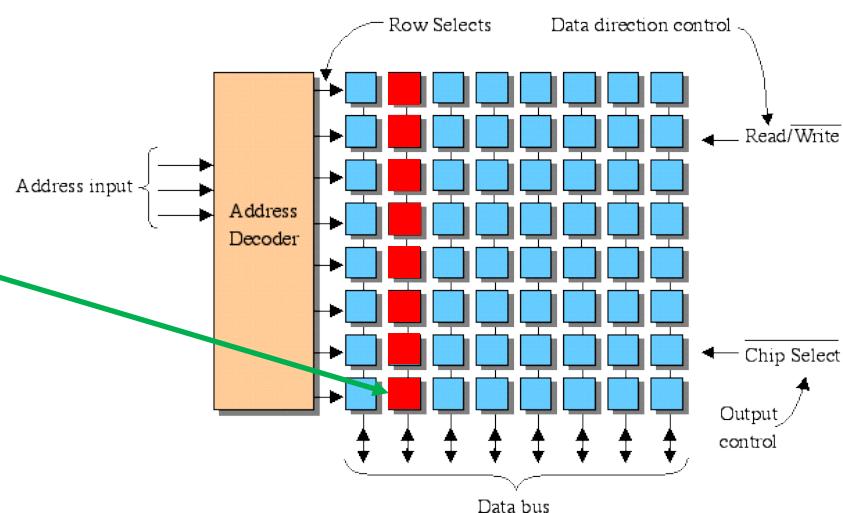
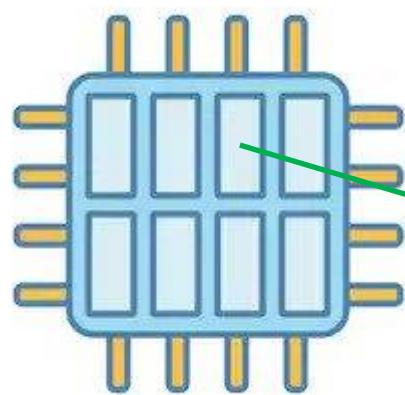
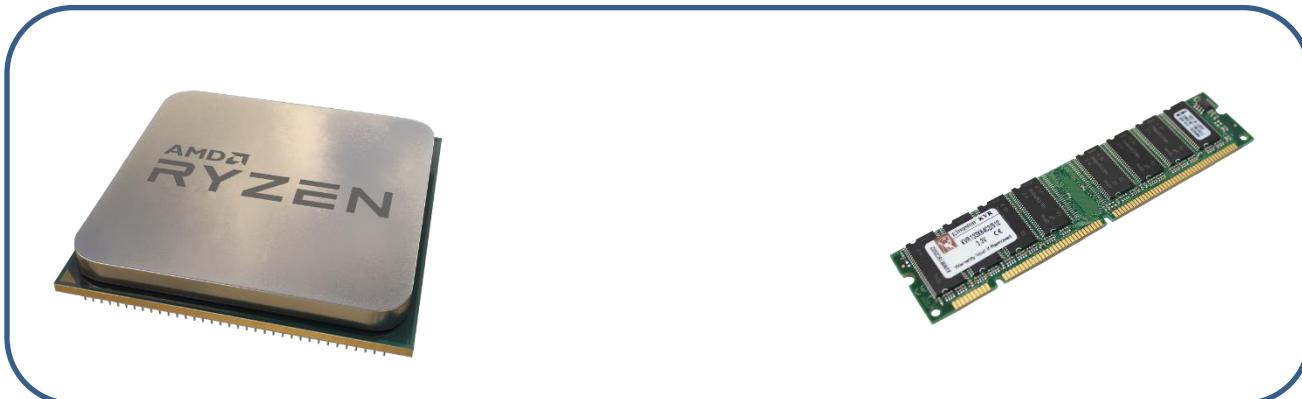
Works in **shared** memory systems  
Can use as many cores as there are in the **compute node**  
Uses **threads** to split the work  
Threads have both **private** and **shared** variables  
Uses **preprocessor directives** in commented lines



# OpenMP™

Open Multi-Processing

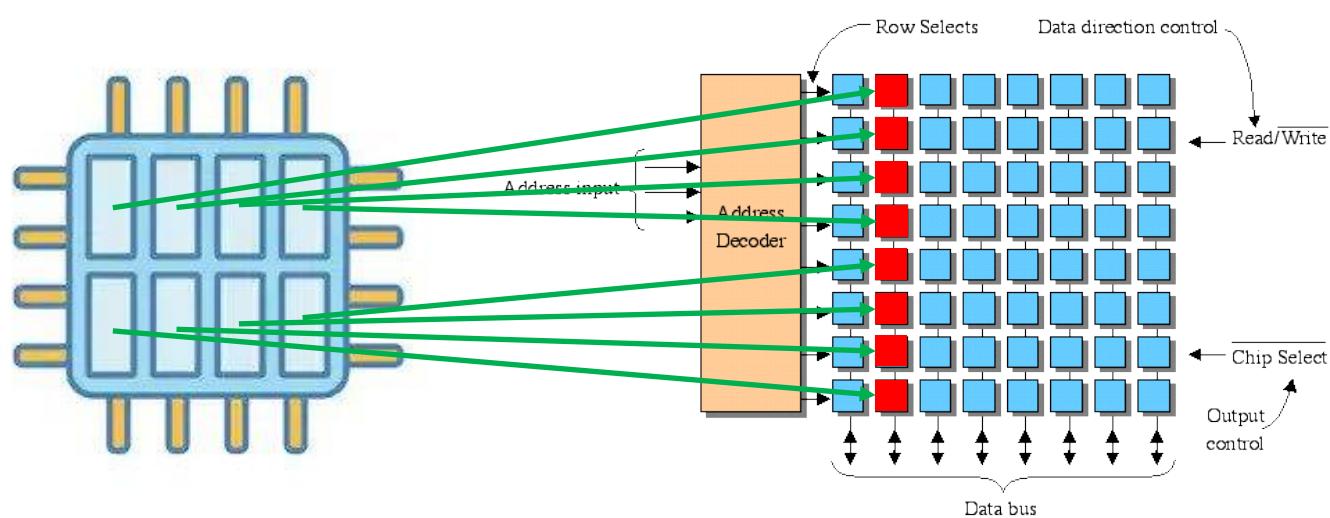
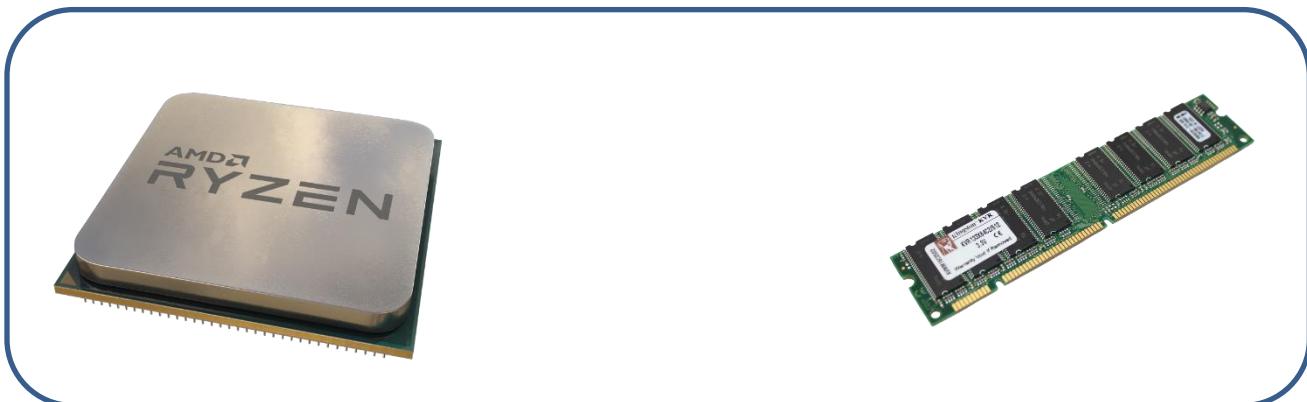
Works in **shared** memory systems  
Can use as many cores as there are in the **compute node**  
Uses **threads** to split the work  
Threads have both **private** and **shared** variables  
Uses **preprocessor directives** in commented lines



# OpenMP™

Open Multi-Processing

Works in **shared** memory systems  
Can use as many cores as there are in the **compute node**  
Uses **threads** to split the work  
Threads have both **private** and **shared** variables  
Uses **preprocessor directives** in commented lines



# Some resources

- Main references:

<https://www.openmp.org/wp-content/uploads/OpenMPRefCard-5-2-web.pdf>  
<https://www.openmp.org/community/>

- Examples & tutorials:

<https://www.openmp.org/wp-content/uploads/openmp-examples-5.1.pdf>  
<https://hpc-tutorials.llnl.gov/openmp/>  
<https://www.youtube.com/playlist?list=PLLX-Q6B8xqZ8n8bwjGdzBJ25X2utwnoEG>  
<https://numba.pydata.org/numba-doc/latest/user/5minguide.html>  
<https://mpi4py.readthedocs.io/en/stable/tutorial.html>

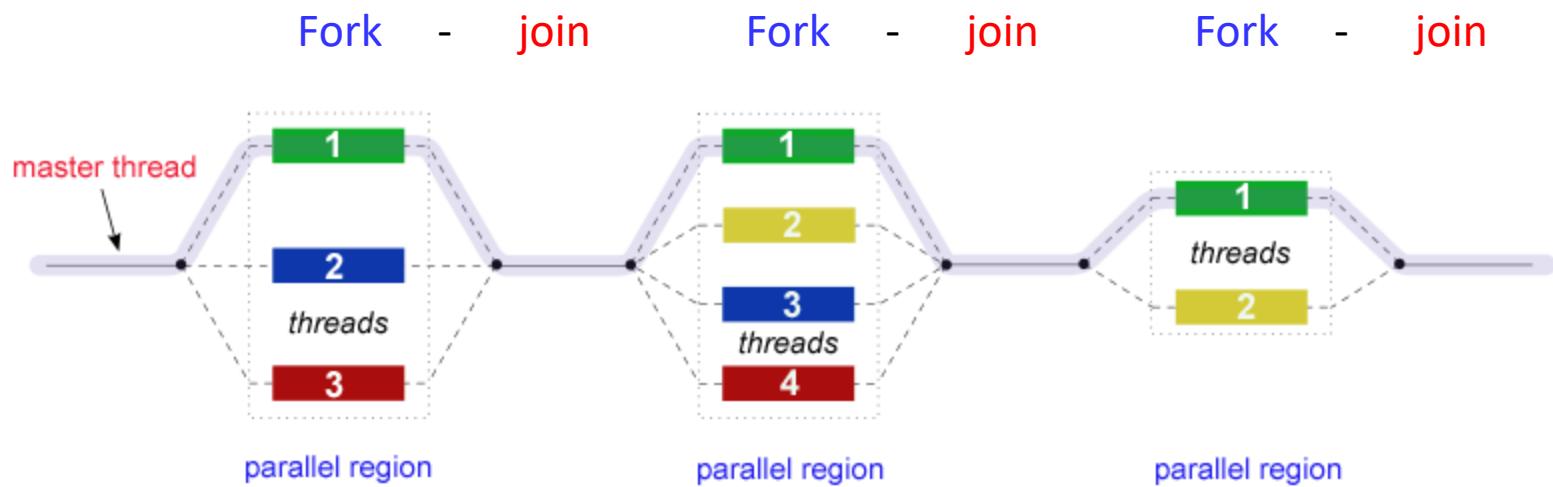
- Books:

The OpenMP Common Core: Making OpenMP Simple Again (Mattson, et al. 2019)

- In-person Courses:

Unibas Master and PhD course, <https://hpc.dmi.unibas.ch/HPC/Teaching.html>

**In particular:** Foundations of Distributed Systems (45402-01/HS): F. Ciorba, H. Schuldt, C. Tschundin  
High-Performance Computing (17164-01/FS): F. Ciorba



## FORTRAN

```
use omp_lib

 !$omp parallel
 ...
 !$omp do
 do ...
 enddo
 !$omp end do
 ...
 !$omp end parallel
```

## C/C++

```
#include <omp.h>

#pragma omp parallel
{
 ...
 #pragma omp for
 for(...){
 ...
 }
```

- Serial and parallel program share the same source code.
- Serial compiler simply overlooks parallel directives (they are comments).
- Modifying parallel directives cannot break the serial code.
- Modifying the serial code outside of parallel regions cannot break the parallelization.
- Easy to maintain.
- Compact code.

```
!$omp parallel private(id)  
id=omp_get_thread_num()  
print *, 'I am thread: ', id  
!$omp end parallel
```

Runtime library routine.

- `omp_get_num_threads()`
- `omp_set_num_threads()`
- `omp_get_wtime()`
- ...

I am thread 0  
I am thread 1  
I am thread 2  
I am thread 3

I am thread 0  
I am thread 3  
I am thread 2  
I am thread 1

I am thread 2  
I am thread 1  
I am thread 0  
I am thread 1

I am thread 0  
I am thread 1  
I am thread 2  
I am thread 3

We cannot ensure that the threads will execute and finish in order!

```
do i = 1, n  
  
    f(i) = sin(dble(i))  
  
end do
```

```
!$omp parallel  
 !$omp do  
 do i = 1, n  
  
     f(i) = sin(dble(i))  
  
 end do  
 !$omp end do  
 !$omp end parallel
```

f(i) =

sin(1)
sin(2)
sin(3)
sin(4)
sin(5)
sin(6)
sin(7)
sin(8)

f(i) =

sin(1)
sin(2)
sin(3)
sin(4)
sin(5)
sin(6)
sin(7)
sin(8)

```
do i = 1, n  
  
    f(i) = sin(dble(i))  
  
end do
```

```
!$omp parallel  
 !$omp do  
 do i = 1, n  
  
     f(i) = sin(dble(i))  
  
 end do  
 !$omp end do  
 !$omp end parallel
```

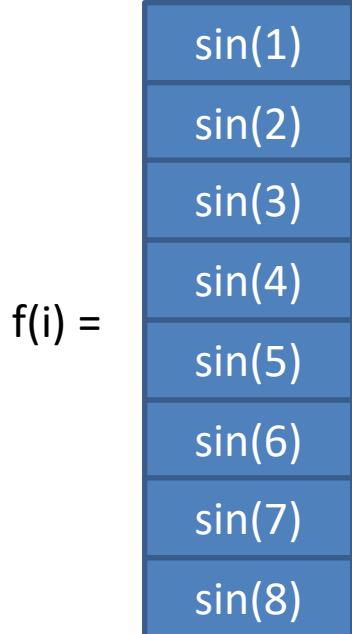
f(i) =	<table border="1"><tr><td>sin(1)</td></tr><tr><td>sin(2)</td></tr><tr><td>sin(3)</td></tr><tr><td>sin(4)</td></tr><tr><td>sin(5)</td></tr><tr><td>sin(6)</td></tr><tr><td>sin(7)</td></tr><tr><td>sin(8)</td></tr></table>	sin(1)	sin(2)	sin(3)	sin(4)	sin(5)	sin(6)	sin(7)	sin(8)
sin(1)									
sin(2)									
sin(3)									
sin(4)									
sin(5)									
sin(6)									
sin(7)									
sin(8)									

Check env. variable, f.e.:  
\$OMP\_NUM\_THREADS=4

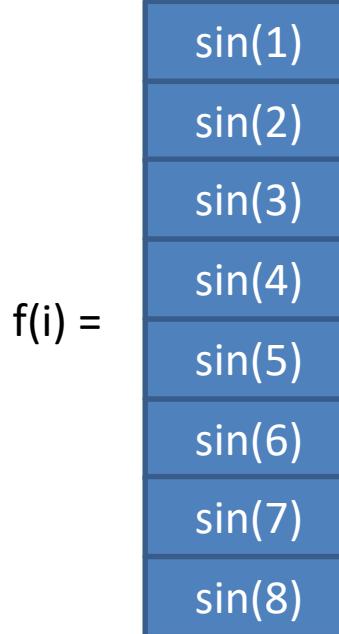
f(i) =	<table border="1"><tr><td>sin(1)</td></tr><tr><td>sin(2)</td></tr><tr><td>sin(3)</td></tr><tr><td>sin(4)</td></tr><tr><td>sin(5)</td></tr><tr><td>sin(6)</td></tr><tr><td>sin(7)</td></tr><tr><td>sin(8)</td></tr></table>	sin(1)	sin(2)	sin(3)	sin(4)	sin(5)	sin(6)	sin(7)	sin(8)
sin(1)									
sin(2)									
sin(3)									
sin(4)									
sin(5)									
sin(6)									
sin(7)									
sin(8)									

```
do i = 1, n  
  
    f(i) = sin(dble(i))  
  
end do
```

```
!$omp parallel  
 !$omp do  
 do i = 1, n  
  
     f(i) = sin(dble(i))  
  
 end do  
 !$omp end do  
 !$omp end parallel
```



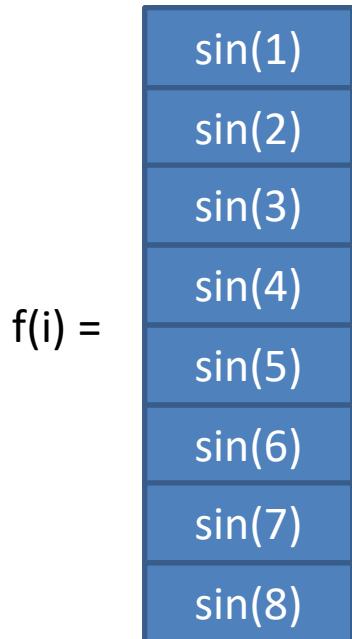
Check env. variable, f.e.:  
 \$OMP\_NUM\_THREADS=4



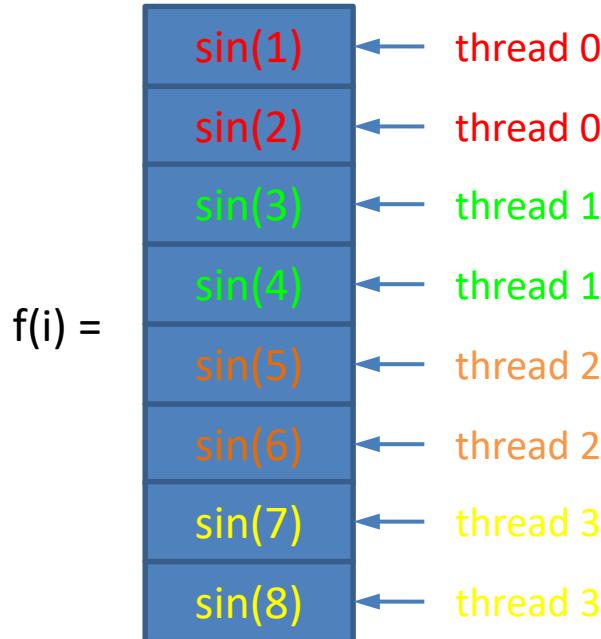
Creates 4 threads and distributes the work evenly

```
do i = 1, n  
  
    f(i) = sin(dble(i))  
  
end do
```

```
!$omp parallel  
 !$omp do  
 do i = 1, n  
  
     f(i) = sin(dble(i))  
  
 end do  
 !$omp end do  
 !$omp end parallel
```



Check env. variable, f.e.:  
 \$OMP\_NUM\_THREADS=4



x4  
Speed-up!

Creates 4 threads and distributes the work evenly

```

!$omp parallel
!$omp do
do i = 1, n

f(i) = sin(dble(i))

end do
!$omp end do
!$omp end parallel

```

These are comments!

They are understood by the compiler only when compiled with `-openmp` option (or similar). Otherwise they are ignored.

An OpenMP calculation will appear as using more than 100% of CPU

Establish parallel section  
Parallel loop evenly divided among the threads

The number of threads is given by the environment variable: `$OMP_NUM_THREADS`  
If not defined, OpenMP assumes `OMP_NUM_THREADS = # cores in the computer`.  
Can be changed with: `export OMP_NUM_THREADS = <integer>`  
(consider adding it to `.bashrc`)

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
23339	cabezon	20	0	89876	904	696	R	769.6	0.0	0:11.19	integral.out
31848	cabezon	20	0	1135m	151m	11m	S	17.8	2.0	318:52.56	chrome
21008	cabezon	20	0	3985m	2.1g	2.1g	S	5.9	27.7	10:53.79	VirtualBox
19624	cabezon	20	0	956m	118m	8828	S	2.0	1.5	241:53.73	chrome
23309	cabezon	20	0	30816	1876	1276	R	2.0	0.0	0:00.32	top
1	root	20	0	36924	588	332	S	0.0	0.0	0:01.17	init
2	root	20	0	0	0	0	S	0.0	0.0	0:00.07	kthreadd
3	root	RT	0	0	0	0	S	0.0	0.0	0:01.64	migration/0
4	root	20	0	0	0	0	S	0.0	0.0	0:05.95	ksoftirqd/0
5	root	RT	0	0	0	0	S	0.0	0.0	0:00.00	stopper/0
6	root	RT	0	0	0	0	S	0.0	0.0	0:01.97	watchdog/0
7	root	RT	0	0	0	0	S	0.0	0.0	0:00.97	migration/1
8	root	RT	0	0	0	0	S	0.0	0.0	0:00.00	stopper/1
9	root	20	0	0	0	0	S	0.0	0.0	0:04.68	ksoftirqd/1
10	root	RT	0	0	0	0	S	0.0	0.0	0:01.78	watchdog/1
11	root	RT	0	0	0	0	S	0.0	0.0	0:00.62	migration/2
12	root	RT	0	0	0	0	S	0.0	0.0	0:00.00	stopper/2

```

!$omp parallel
!$omp do schedule(static)
do i = 1, n
    f(i) = sin(dble(i))
end do
 !$omp end do
 !$omp end parallel

```

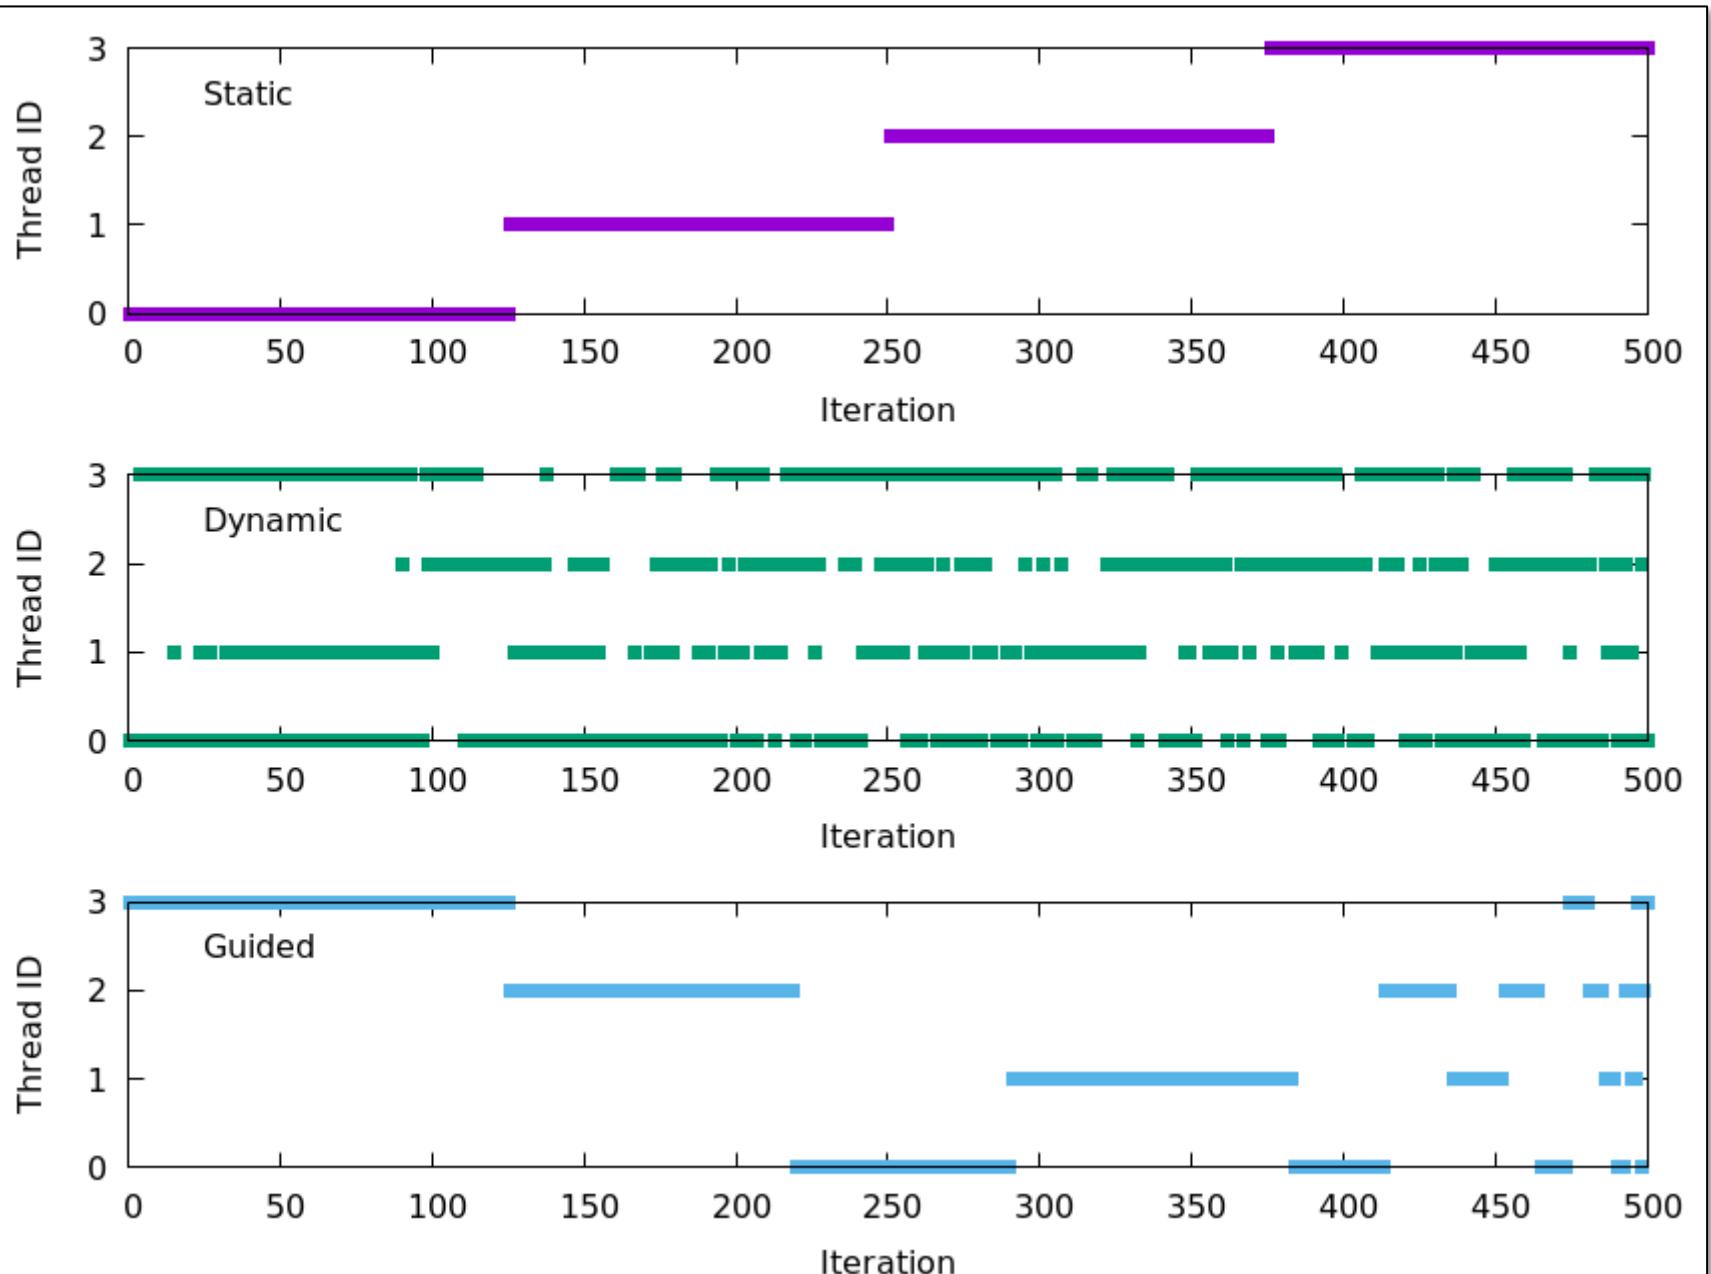
Parallel loop **evenly** divided among the threads

$$\text{chunk} = \frac{\text{loop count}}{\# \text{ threads}}$$

**static** schedule

Other options are available: **!\$omp do schedule(<kind> [,chunk\_size])**

Kind	Description
<b>static</b>	Divides the loop in equal-sized chunks (or as equal as possible).
<b>dynamic</b>	When a thread finishes, it retrieves the next chunk from the internal queue. Be aware of the extra overhead! (default chunk size = 1).
<b>guided</b>	Similar to dynamic, but starts off large and decreases to better handle imbalance. (default chunk size = same as static).
<b>auto</b>	Decision regarding scheduling is delegated to the compiler.
<b>runtime</b>	Uses OMP_SCHEDULE env. Variable to select the scheduling type.



```
factorial(1) = 1  
  
do i = 2, 6  
    factorial(i) = i * factorial(i-1)  
end do
```

factorial(i) =	1
	2
	6
	24
	120
	720
	5040

```
factorial(1) = 1  
  
 !$omp parallel  
 !$omp do schedule(static)  
 do i = 2, 6  
  
     factorial(i) = i * factorial(i-1)  
  
 end do  
 !$omp end do  
 !$omp end parallel
```

factorial(i) =

1
2
6
24
120
720
5040

There is something wrong here!

We cannot ensure this!

factorial(i) =



1
2
6
24
120
720
5040

```
factorial(1) = 1  
  
 !$omp parallel  
 !$omp do schedule(static)  
 do i = 2, 6  
  
     factorial(i) = i * factorial(i-1)  
  
 end do  
 !$omp end do  
 !$omp end parallel
```

factorial(i) =

1
2
6
24
120
720
5040

factorial(i) =

1
2

Threads access data in a disordered way

← thread 0

```
factorial(1) = 1  
  
 !$omp parallel  
 !$omp do schedule(static)  
 do i = 2, 6  
  
     factorial(i) = i * factorial(i-1)  
  
 end do  
 !$omp end do  
 !$omp end parallel
```

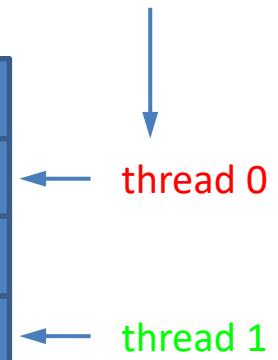
factorial(i) =

1
2
6
24
120
720
5040

factorial(i) =

1
2
0

Threads access data in a disordered way



```
factorial(1) = 1  
  
 !$omp parallel  
 !$omp do schedule(static)  
 do i = 2, 6  
  
     factorial(i) = i * factorial(i-1)  
  
 end do  
 !$omp end do  
 !$omp end parallel
```

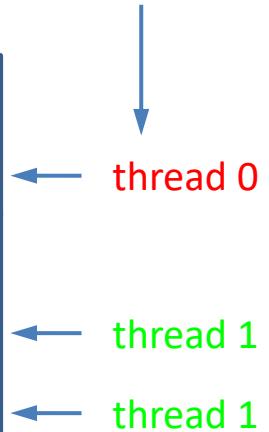
factorial(i) =

1
2
6
24
120
720
5040

factorial(i) =

1
2
0
0

Threads access data in a disordered way



```
factorial(1) = 1  
  
 !$omp parallel  
 !$omp do schedule(static)  
 do i = 2, 6  
  
     factorial(i) = i * factorial(i-1)  
  
 end do  
 !$omp end do  
 !$omp end parallel
```

factorial(i) =

1
2
6
24
120
720
5040

factorial(i) =

1
2
6
0
0

Threads access data in a disordered way



- ← thread 0
- ← thread 0
- ← thread 1
- ← thread 1

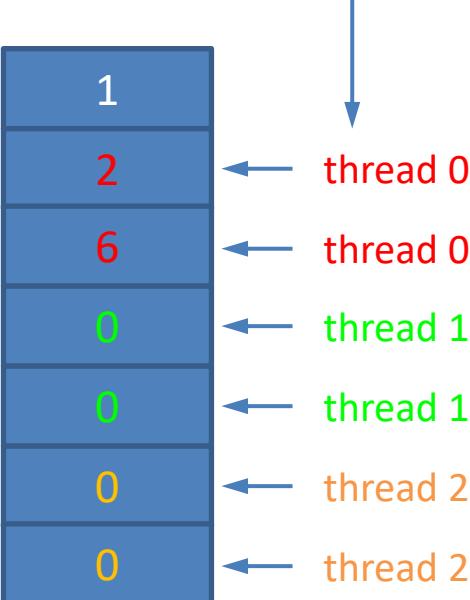
```
factorial(1) = 1  
  
 !$omp parallel  
 !$omp do schedule(static)  
 do i = 2, 6  
  
     factorial(i) = i * factorial(i-1)  
  
 end do  
 !$omp end do  
 !$omp end parallel
```

factorial(i) =

1
2
6
24
120
720
5040

factorial(i) =

Threads access data in a disordered way



The diagram illustrates the state of the factorial array after three threads have executed their iterations. Thread 0 (red) has updated elements at indices 2 and 3. Thread 1 (green) has updated elements at indices 4 and 5. Thread 2 (orange) has updated elements at indices 6 and 7. The resulting values in the array are 1, 2, 6, 0, 0, 0, 0, 0.

1
2
6
0
0
0
0
0

```

factorial(1) = 1

 !$omp parallel
 !$omp do schedule(static)
 do i = 2, 6

     factorial(i) = i * factorial(i-1)

 end do
 !$omp end do
 !$omp end parallel

```

factorial(i) =

1
2
6
24
120
720
5040

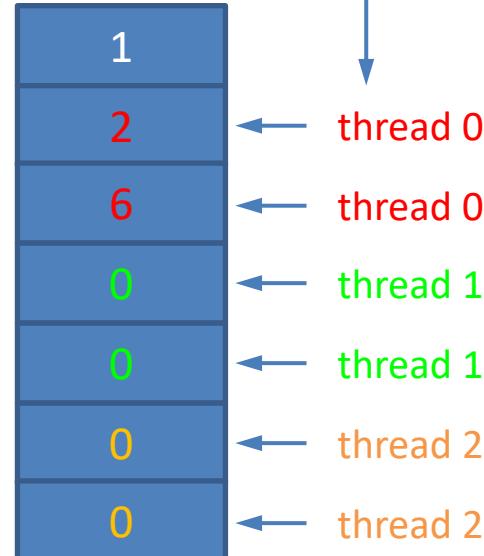
This is called **Race Condition**

It can even be worse if factorial was not properly initialized!

Difficult to detect because for a given case, system or run, the threads may win the race in an order that happens to make the program run correctly.

Threads access data in a disordered way

factorial(i) =



The diagram illustrates a race condition in a shared memory space. A vertical stack of eight cells represents memory. Thread 0 (blue arrow) writes values 2, 6, and 0 to the first three cells. Thread 1 (green arrow) writes values 0 and 0 to the fourth and fifth cells. Thread 2 (orange arrow) writes values 0 and 0 to the sixth and eighth cells. The final state shows the array containing [1, 2, 6, 0, 0, 0, 0, 0], which is incorrect for  $i=6$ .

1
2
6
0
0
0
0
0

```
!$omp parallel
!$omp do schedule(static)
do i = 1, 6

    x = i
    y(i) = x * 2

end do
!$omp end do
!$omp end parallel
```

$y(i) =$

2
4
6
8
10
12

$y(i) =$

x	$y(i)$
1	

← thread 0

## Another example of **Race Condition**

```
!$omp parallel
!$omp do schedule(static)
do i = 1, 6

    x = i
    y(i) = x * 2

end do
!$omp end do
!$omp end parallel
```

$y(i) =$	2
	4
	6
	8
	10
	12

$y(i) =$

x	$y(i)$
1	2

← thread 0

## Another example of **Race Condition**

```
!$omp parallel
 !$omp do schedule(static)
 do i = 1, 6

   x = i
   y(i) = x * 2

 end do
 !$omp end do
 !$omp end parallel
```

$y(i) =$	<table border="1"><tr><td>2</td></tr><tr><td>4</td></tr><tr><td>6</td></tr><tr><td>8</td></tr><tr><td>10</td></tr><tr><td>12</td></tr></table>	2	4	6	8	10	12
2							
4							
6							
8							
10							
12							

## Another example of **Race Condition**

x	y(i)
1	2
2	

← thread 0  
← thread 0

```

!$omp parallel
 !$omp do schedule(static)
 do i = 1, 6

   x = i
   y(i) = x * 2

 end do
 !$omp end do
 !$omp end parallel

```

$y(i) =$	2
	4
	6
	8
	10
	12

$y(i) =$

x	y(i)
1	2
2	
5	

← thread 0  
← thread 0  
← thread 2

Another example of **Race Condition**

```

!$omp parallel
 !$omp do schedule(static)
 do i = 1, 6

   x = i
   y(i) = x * 2

 end do
 !$omp end do
 !$omp end parallel

```

$y(i) =$	2
	4
	6
	8
	10
	12

$y(i) =$

x	$y(i)$
1	2
2	10
5	

← thread 0  
← thread 0  
← thread 2

Another example of **Race Condition**

```

!$omp parallel
 !$omp do schedule(static)
 do i = 1, 6

   x = i
   y(i) = x * 2

 end do
 !$omp end do
 !$omp end parallel

```

$y(i) =$	<table border="1"><tr><td>2</td></tr><tr><td>4</td></tr><tr><td>6</td></tr><tr><td>8</td></tr><tr><td>10</td></tr><tr><td>12</td></tr></table>	2	4	6	8	10	12
2							
4							
6							
8							
10							
12							

$y(i) =$

x	y(i)	
1	2	← thread 0
2	10	← thread 0
3		← thread 1
5		← thread 2

Another example of **Race Condition**

```

!$omp parallel
 !$omp do schedule(static)
 do i = 1, 6

   x = i
   y(i) = x * 2

 end do
 !$omp end do
 !$omp end parallel

```

$y(i) =$	<table border="1"><tr><td>2</td></tr><tr><td>4</td></tr><tr><td>6</td></tr><tr><td>8</td></tr><tr><td>10</td></tr><tr><td>12</td></tr></table>	2	4	6	8	10	12
2							
4							
6							
8							
10							
12							

## Another example of **Race Condition**

x	y(i)	
1	2	← thread 0
2	10	← thread 0
3		← thread 1
5	6	← thread 2

```

!$omp parallel
 !$omp do schedule(static)
 do i = 1, 6

   x = i
   y(i) = x * 2

 end do
 !$omp end do
 !$omp end parallel

```

$y(i) =$	<table border="1"><tr><td>2</td></tr><tr><td>4</td></tr><tr><td>6</td></tr><tr><td>8</td></tr><tr><td>10</td></tr><tr><td>12</td></tr></table>	2	4	6	8	10	12
2							
4							
6							
8							
10							
12							

$y(i) =$

x	y(i)	
1	2	← thread 0
2	10	← thread 0
3	6	← thread 1
5	6	← thread 2

Another example of **Race Condition**

```

!$omp parallel
 !$omp do schedule(static)
 do i = 1, 6

   x = i
   y(i) = x * 2

 end do
 !$omp end do
 !$omp end parallel

```

$y(i) =$	<table border="1"><tr><td>2</td></tr><tr><td>4</td></tr><tr><td>6</td></tr><tr><td>8</td></tr><tr><td>10</td></tr><tr><td>12</td></tr></table>	2	4	6	8	10	12
2							
4							
6							
8							
10							
12							

$y(i) =$

x	y(i)	
1	2	← thread 0
2	10	← thread 0
3	6	← thread 1
4		← thread 1
5	6	← thread 2

Another example of **Race Condition**

```

!$omp parallel
 !$omp do schedule(static)
 do i = 1, 6

   x = i
   y(i) = x * 2

 end do
 !$omp end do
 !$omp end parallel

```

$y(i) =$	2
	4
	6
	8
	10
	12

$y(i) =$

x	y(i)	
1	2	← thread 0
2	10	← thread 0
3	6	← thread 1
4		← thread 1
5	6	← thread 2
6		← thread 2

Another example of **Race Condition**

```

!$omp parallel
 !$omp do schedule(static)
 do i = 1, 6

   x = i
   y(i) = x * 2

 end do
 !$omp end do
 !$omp end parallel

```

$y(i) =$	<table border="1"><tr><td>2</td></tr><tr><td>4</td></tr><tr><td>6</td></tr><tr><td>8</td></tr><tr><td>10</td></tr><tr><td>12</td></tr></table>	2	4	6	8	10	12
2							
4							
6							
8							
10							
12							

$y(i) =$

x	y(i)	
1	2	← thread 0
2	10	← thread 0
3	6	← thread 1
4	12	← thread 1
5	6	← thread 2
6	12	← thread 2

Another example of **Race Condition**

To solve this we have the **private** clause.

```

!$omp parallel private (x)
 !$omp do schedule(static)
 do i = 1, 6

   x = i
   y(i) = x * 2

 end do
 !$omp end do
 !$omp end parallel

```

	2
	4
	6
	8
	10
	12

y(i) =

x <sub>0</sub>	x <sub>1</sub>	x <sub>2</sub>	y(i)
1			2
2			4
	3		6
	4		8
	5	10	
	6	12	

y(i) =

← thread 0  
← thread 0  
← thread 1  
← thread 1  
← thread 2  
← thread 2

## Another example of **Race Condition**

To solve this we have the **private** clause.

With this, each thread has its private copy of x and the problem is solved.

This comes at the cost of an overhead in memory.

Can you declare all private variables of the following code?

**NOTE:** jumpx and w1d are outputs of the corresponding subroutine calls.

```
!$omp parallel private (???)  
!$omp do schedule(static)  
do i = ini, end  
    ii = 1+dim*(i-1)  
    do k = 1, nvi(i)  
        j = neighbors(i,k)  
        call apply_PBC(i,k,0,jumpx)  
        d1 = a(ii) - a(jj) - jumpx  
        d05 = sqrt(d1)  
        v1 = d05/h(i)  
        call Wkernel(v1,w1d)  
        dter = pk(i)*w1d  
        sumwh(i) = sumwh(i) + xmass(j) * dter  
        if (equ) then  
            dlw1d = log10(w1d)/indice(i)  
        end if  
    end do  
    do it = 1, nut  
        do ie = 1, ne  
            f(ie, it) = kfactor * alpha( ie, it, i)  
            ftot(ie, it, i) = f(ie, it) / dens(i)  
        end do  
    end do  
end do  
!$omp end do  
!$omp end parallel
```

## Automatized help: Intel Inspector

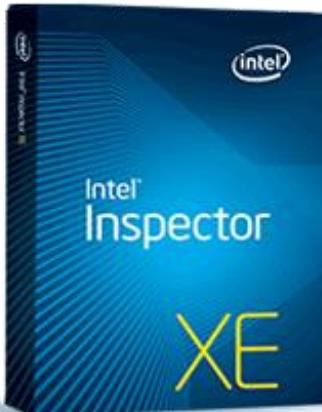
Commercial tool, but free licenses for academia and open source projects.

<https://www.intel.com/content/www/us/en/developer/articles/tool/oneapi-standalone-components.html#inspector>

Available for Linux and Windows. Supports C, C++, and Fortran. GUI + command line.

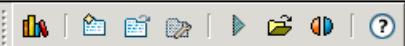
10 min intro: <https://www.intel.com/content/www/us/en/developer/videos/introduction-to-intel-inspector.html>

Detects Race conditions and Deadlocks.



```
> ifort -g -openmp myprogram.f90 -o myprogram  
  
> export NUM_THREADS_OMP=2  
> inspxe-gui
```

File View Help



Project Navigator



Welcome 

 [Getting Started](#)

# Welcome to Intel Inspector 2020

## *Memory and Thread Debugging*

 [New Project...](#)



 [Open Project...](#)

 [Open Result](#)

Recent Projects:

Recent Results:

## Target

## Suppressions

## Binary/Symbol Search

## Source Search

Specify and configure your analysis target: an application or a script to execute. Press F1 for more details.

Application: /scicore/home/scicore/cabezon/openmp\_course/a.out

Application parameters:

Use application directory as working directory

Working directory: /scicore/home/scicore/cabezon/openmp\_course

User-defined environment variables:

Store result in the project directory: /scicore/home/scicore/cabezon/intel/inspxe/projects/test\_race\_condition

Store result in (and create link file to) another directory

/scicore/home/scicore/cabezon/intel/inspxe/projects/test\_race\_condition

Result location:

/scicore/home/scicore/cabezon/intel/inspxe/projects/test\_race\_condition/r@@@{at}

Advanced

OK

Cancel

Search for the executable

test\_race\_condition - Project Properties@login20.cluster.bc2.ch

Target    Suppressions    Binary/Symbol Search    Source Search

Specify and configure your analysis target: an application or a script to execute. Press F1 for more details.

Application: /scicore/home/scicore/cabezon/openmp\_course/a.out    Browse...

Application parameters:    Modify...

Use application directory as working directory

Working directory: /scicore/home/scicore/cabezon/openmp\_course    Browse...

User-defined environment variables:    Modify...

Store result in the project directory: /scicore/home/scicore/cabezon/intel/inspxe/projects/test\_race\_condition

Store result in (and create link file to) another directory  
/scicore/home/scicore/cabezon/intel/inspxe/projects/test\_race\_condition    Browse...

Result location:  
/scicore/home/scicore/cabezon/intel/inspxe/projects/test\_race\_condition/r@@@{at}

Advanced

OK    Cancel



② [Getting Started](#)

# Welcome to Intel Inspector 2020

## *Memory and Thread Debugging*

**Current project: test\_race\_condition**

- ▶ [Memory Error Analysis / Locate Memory Problems](#)
- ▶ [Threading Error Analysis / Locate Deadlocks and Data Races](#)
- ▶ [Memory Error Analysis / Detect Leaks](#)
- ▶ [New Analysis...](#)

 [New Project...](#)

 [Open Project...](#)

 [Open Result](#)

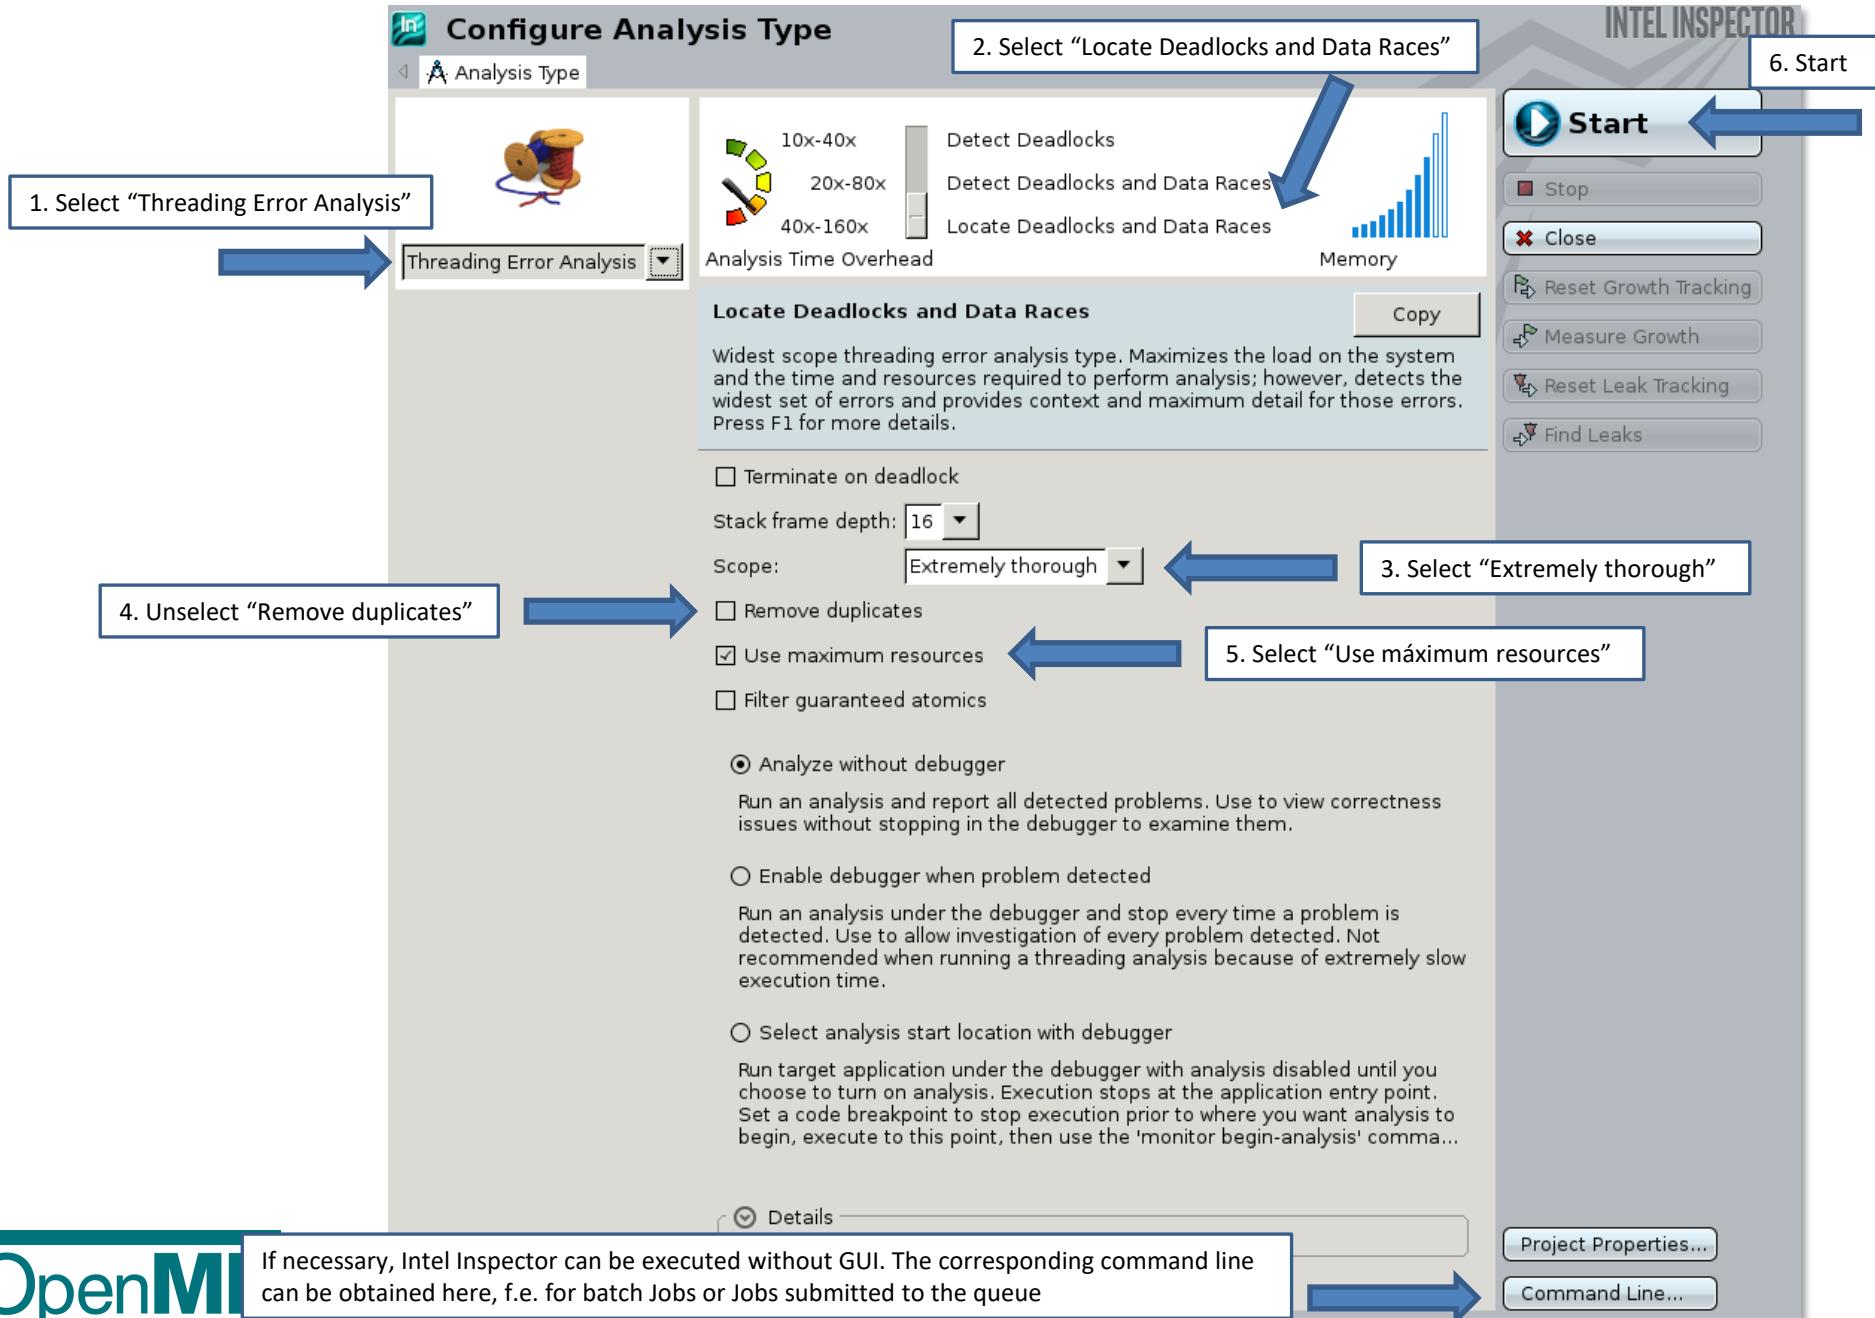
Recent Projects:

Recent Results:



**Configure Analysis Type**

1. Select "Threading Error Analysis"



2. Select "Locate Deadlocks and Data Races"

3. Select "Extremely thorough"

4. Unselect "Remove duplicates"

5. Select "Use maximum resources"

6. Start

Threading Error Analysis

Analysis Time Overhead

10x-40x  
20x-80x  
40x-160x

Detect Deadlocks  
Detect Deadlocks and Data Races  
Locate Deadlocks and Data Races

Memory

Locate Deadlocks and Data Races

Widest scope threading error analysis type. Maximizes the load on the system and the time and resources required to perform analysis; however, detects the widest set of errors and provides context and maximum detail for those errors. Press F1 for more details.

Terminate on deadlock

Stack frame depth: 16

Scope: Extremely thorough

Remove duplicates

Use maximum resources

Filter guaranteed atomics

Analyze without debugger

Run an analysis and report all detected problems. Use to view correctness issues without stopping in the debugger to examine them.

Enable debugger when problem detected

Run an analysis under the debugger and stop every time a problem is detected. Use to allow investigation of every problem detected. Not recommended when running a threading analysis because of extremely slow execution time.

Select analysis start location with debugger

Run target application under the debugger with analysis disabled until you choose to turn on analysis. Execution stops at the application entry point. Set a code breakpoint to stop execution prior to where you want analysis to begin, execute to this point, then use the 'monitor begin-analysis' comma...

Details

If necessary, Intel Inspector can be executed without GUI. The corresponding command line can be obtained here, f.e. for batch Jobs or Jobs submitted to the queue

Project Properties...

Command Line...

**Locate Deadlocks and Data Races**

Target Analysis Type Collection Log Summary

**Problems**

ID	Type	Sources	Modules	State
P1	Data race	pi.f90	a.out	New

List of errors

Location in source code

**Filters**

- Severity: Error (1 item(s))
- Type: Data race (1 item(s))
- Source: pi.f90 (1 item(s))
- Module: a.out (1 item(s))
- State: New (1 item(s))
- Suppressed: Not suppressed (1 item(s))
- Investigated: Not investigated (1 item(s))

**Code Locations: Data race**

Description	Source	Function	Module	Variable
Write	pi.f90:17	MAIN__\$omp\$parallel@14	a.out	x
	15 !\$omp do reduction(+:sum)		a.out!MAIN__\$omp\$parallel@14 - pi.f90	
	16 do i=1,steps			
	17 x=(dble(i)-.5d0)*dx			
	18 sum=sum+4.d0/(1.d0+x*x)			
	19 enddo			

Description	Source	Function	Module	Variable
Write	pi.f90:17	MAIN__\$omp\$parallel@14	a.out	x
	15 !\$omp do reduction(+:sum)		a.out!MAIN__\$omp\$parallel@14 - pi.f90	
	16 do i=1,steps			
	17 x=(dble(i)-.5d0)*dx			
	18 sum=sum+4.d0/(1.d0+x*x)			
	19 enddo			

**Timeline**

- OMP Worker Thread #8 (18303)
- OMP Worker Thread #15 (18311)

A note regarding private variables: They are not initialized!

```
result = 20  
  
!$omp parallel private(result)  
result = result + 10  
print *, 'Thread ',omp_get_thread_num(), result  
!$omp end parallel  
  
print *, result
```

To control the initialization of private variables we have **firstprivate**

```
Thread 3 10  
Thread 0 10  
Thread 1 71  
Thread 2 10
```

```
20
```

If you are lucky, result may access a section of the memory that is empty. Otherwise, any value can be used to initialize result.

A note regarding private variables: **They are not initialized!**

```
result = 20  
  
!$omp parallel firstprivate(result)  
result = result + 10  
print *, 'Thread ',omp_get_thread_num(), result  
!$omp end parallel  
  
print *, result
```

To control the initialization of private variables we have **firstprivate**

```
Thread 3 30  
Thread 0 30  
Thread 1 30  
Thread 2 30  
  
20
```

Note that this output has no relation with the OpenMP section! Private variables are destroyed once the parallel section ends.

Consider now this code. Is it correct?

```
sum = 0

 !$omp parallel
 !$omp do
 do i = 1, 6

 sum = sum + i

end do
 !$omp end do
 !$omp end parallel
```

The variable sum should be private to give the correct answer, but it should be also shared to be accessed by all threads!

Consider now this code. Is it correct?

```
sum = 0

 !$omp parallel
 !$omp do reduction(+: sum)
 do i = 1, 6

     sum = sum + i

 end do
 !$omp end do
 !$omp end parallel
```

The variable sum should be private to give the correct answer, but it should be also shared to be accessed by all threads!

We can solve this with **reduction**

This allows each thread to have a private copy of sum where they store their partial calculation and then, when the threads exit, it groups all partial values from all threads using the defined operation (+ in this case) in one global variable.

Reduction variables have to meet the following requirements:

- They can only be listed in one reduction
- Cannot be declared constant
- Cannot be declared private in the parallel construct

Consider now this code. Is it correct?

```
sum = 0

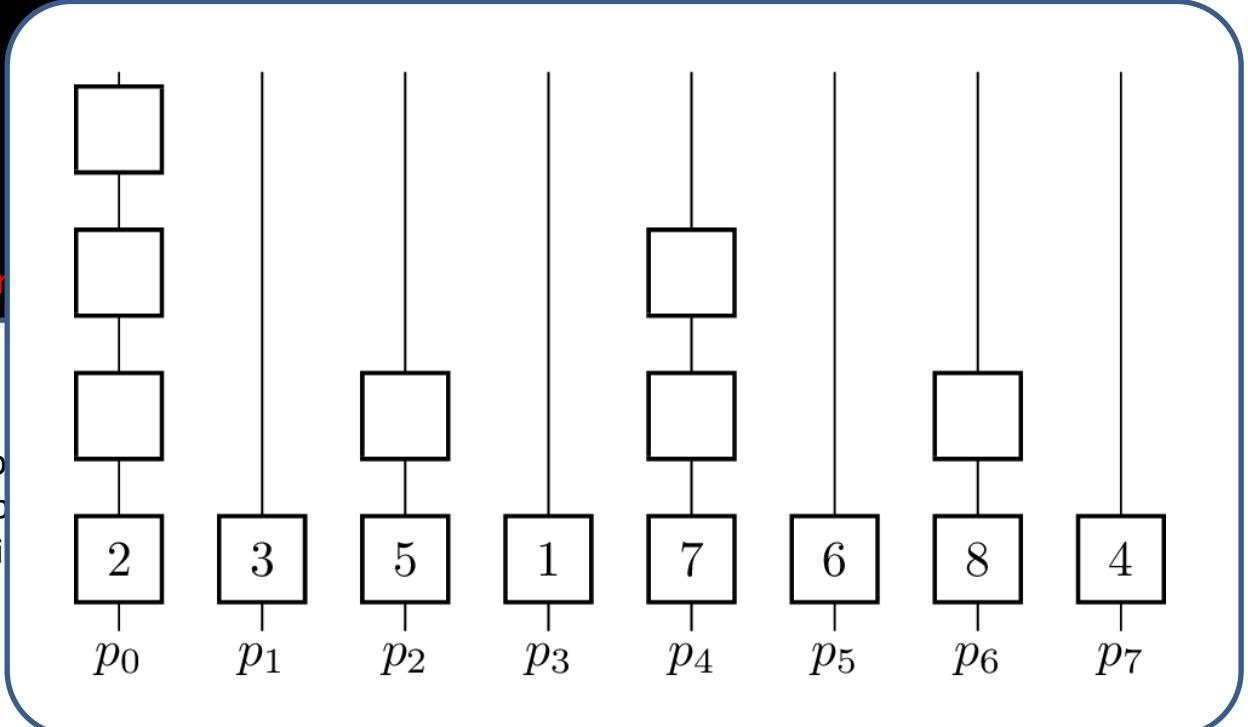
 !$omp parallel
 !$omp do reduction(+: sum)
 do i = 1, 6

     sum = sum

 end do
 !$omp end do
 !$omp end par
```

This allows each thread to store their partial sum. When they exit, it groups all partial sums and performs a reduction operation (+ in this case).

The variable sum should be private to give the correct answer, but it should be also shared to be accessed by all threads.



Source: wikipedia

- They can only be listed in one reduction
- Cannot be declared constant
- Cannot be declared private in the parallel construct

# Getting the course exercises

Clone the repository of the exercises:

```
git clone https://git.scicore.unibas.ch/scicore/openmp_course.git
```

Or download them directly and extract:

```
https://git.scicore.unibas.ch/scicore/openmp\_course/-/archive/master/openmp\_course-master.tar.gz
```

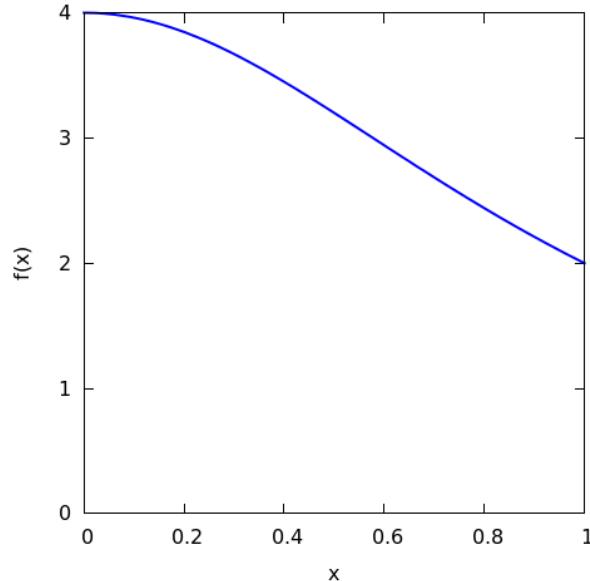
Enter in openmp\_course/:

```
cd openmp_course
```

## Exercise: Numerical integration

We know that:

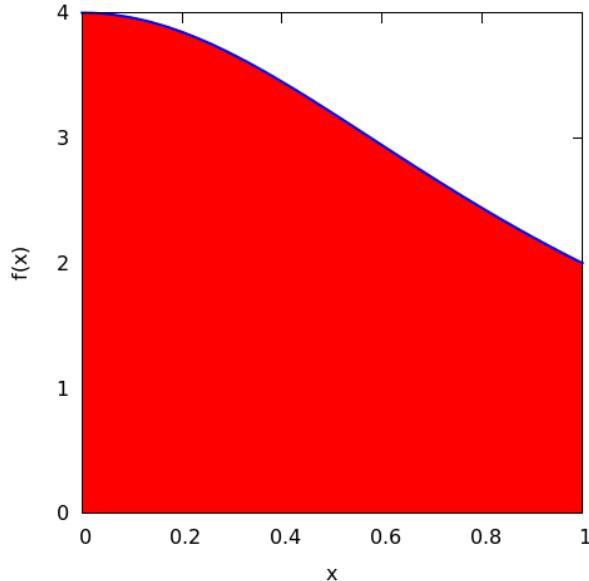
$$\int_0^1 \frac{4}{(1+x^2)} dx = \pi$$



## Exercise: Numerical integration

We know that:

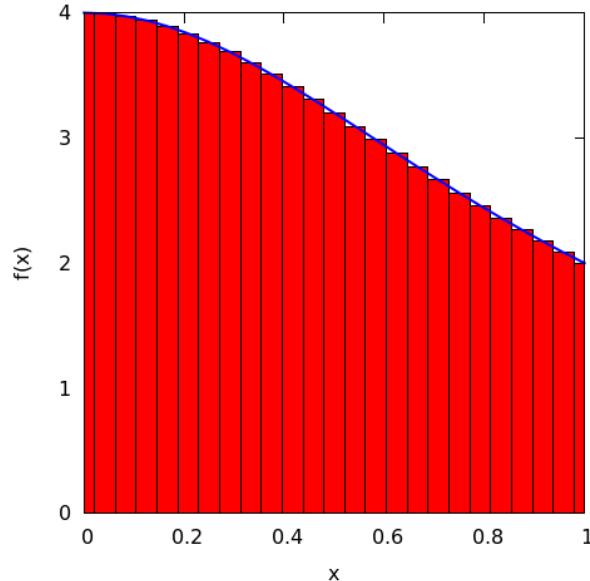
$$\int_0^1 \frac{4}{(1+x^2)} dx = \pi$$



## Exercise: Numerical integration

We know that:

$$\int_0^1 \frac{4}{(1+x^2)} dx = \pi$$



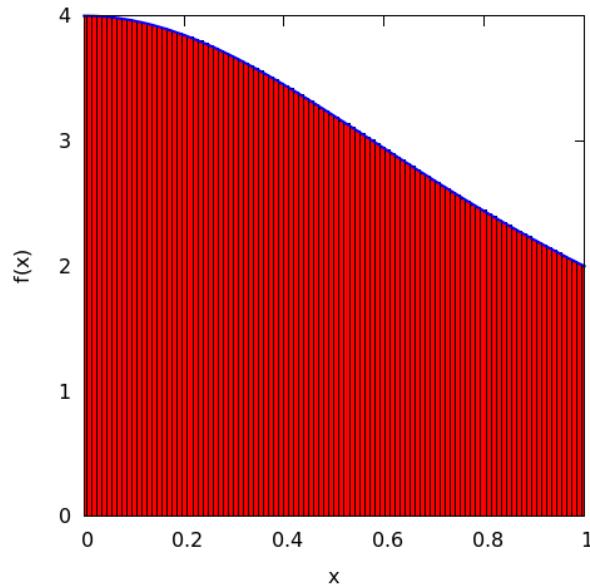
Using the composite formula we can approximate the integral as:

$$\int_a^b f(x) dx \approx \left( \sum_{i=0}^n f(x_i) \right) \frac{b-a}{n}$$

## Exercise: Numerical integration

We know that:

$$\int_0^1 \frac{4}{(1+x^2)} dx = \pi$$



Using the composite formula we can approximate the integral as:

$$\int_a^b f(x) dx \approx \left( \sum_{i=0}^n f(x_i) \right) \frac{b-a}{n}$$

## Exercise: Numerical integration

Parallelize the following program with openMP

```
steps = 1000000000
dx = 1./dble(steps)
sum = 0.
do i = 1, steps
    x = (dble(i)-0.5)*dx
    sum = sum + 4./(1. + x*x)
end do
sum = sum * dx
```

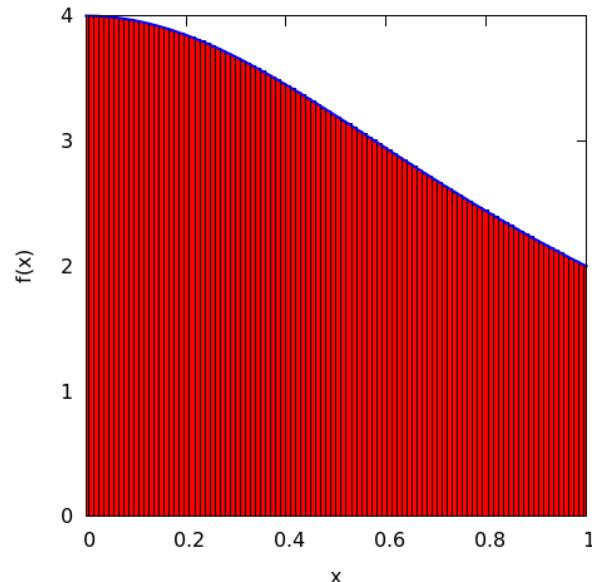
Compile with: **gfortran -fopenmp pi.f90**

Fix the number of threads: **export OMP\_NUM\_THREADS=4**

Execute with: **time ./a.out**

We know that:

$$\int_0^1 \frac{4}{(1+x^2)} dx = \pi$$



Using the composite formula we can approximate the integral as:

$$\int_a^b f(x) dx \approx \left( \sum_{i=0}^n f(x_i) \right) \frac{b-a}{n}$$

What happens if we have nested loops?

```
do i = 1, ni
  do j = 1, nj
    do k = 1, nk
      <calculations>
    end do
  end do
end do
```

## What happens if we have nested loops?

Not much... we simply parallelize the outer loop.

```
!$omp parallel private(i,j,k)
 !$omp do
 do i = 1, ni
   do j = 1, nj
     do k = 1, nk
       <calculations>
     end do
   end do
 end do
 !$omp end do
 !$omp end parallel
```

But this depends on the number of iterations for each loop.

## What happens if we have nested loops?

Not much... we simply parallelize the outer loop.

```
!$omp parallel private(i,j,k)
 !$omp do
 do i = 1, 4 ←
   do j = 1, 20
     do k = 1, 10000
       <calculations>
     end do
   end do
 end do
 !$omp end do
 !$omp end parallel
```

But this depends on the number of iterations for each loop.

Having such small count of iterations in the outer loop is not efficient. We might even have less iterations than threads!

## What happens if we have nested loops?

Not much... we simply parallelize the outer loop.

```
!$omp parallel private(i,j,k)
 !$omp do
 do k = 1, 10000
   do j = 1, 20
     do i = 1, 4
       <calculations>
     end do
   end do
 end do
 !$omp end do
 !$omp end parallel
```

But this depends on the number of iterations for each loop.

Having such small count of iterations in the outer loop is not efficient. We might even have less iterations than threads!

Re-arranging the nested loops will solve it!

But this is not always possible, or maybe all loops have relatively small counts!

## What happens if we have nested loops?

Not much... we simply parallelize the outer loop.

```
!$omp parallel private(i,j,k)
 !$omp do
 do i = 1, 10
   do j = 1, 10
     do k = 1, 10
       <calculations>
     end do
   end do
 end do
 !$omp end do
 !$omp end parallel
```

But this depends on the number of iterations for each loop.

Having such small count of iterations in the outer loop is not efficient. We might even have less iterations than threads!

Re-arranging the nested loops will solve it!

But this is not always possible, or maybe all loops have relatively small counts!

## What happens if we have nested loops?

Not much... we simply parallelize the outer loop.

```
!$omp parallel private(i,j,k)
!$omp do collapse(3)
do i = 1, 10
  do j = 1, 10
    do k = 1, 10
      <calculations>
    end do
  end do
end do
 !$omp end do
 !$omp end parallel
```

To solve this we have the clause **collapse(n)**

It specifies how many nested loops will be collapsed in a single loop.

But this depends on the number of iterations for each loop.

Having such small count of iterations in the outer loop is not efficient. We might even have less iterations than threads!

Re-arranging the nested loops will solve it!

But this is not always possible, or maybe all loops have relatively small counts!

```
!$omp parallel private(ii,i,j,k)
!$omp do
do ii = 1, 1000
  i = mod(ii/100, 10)+1
  j = mod(ii/10, 10)+1
  k = mod(ii, 10)+1
  <calculations>
end do
 !$omp end do
 !$omp end parallel
```

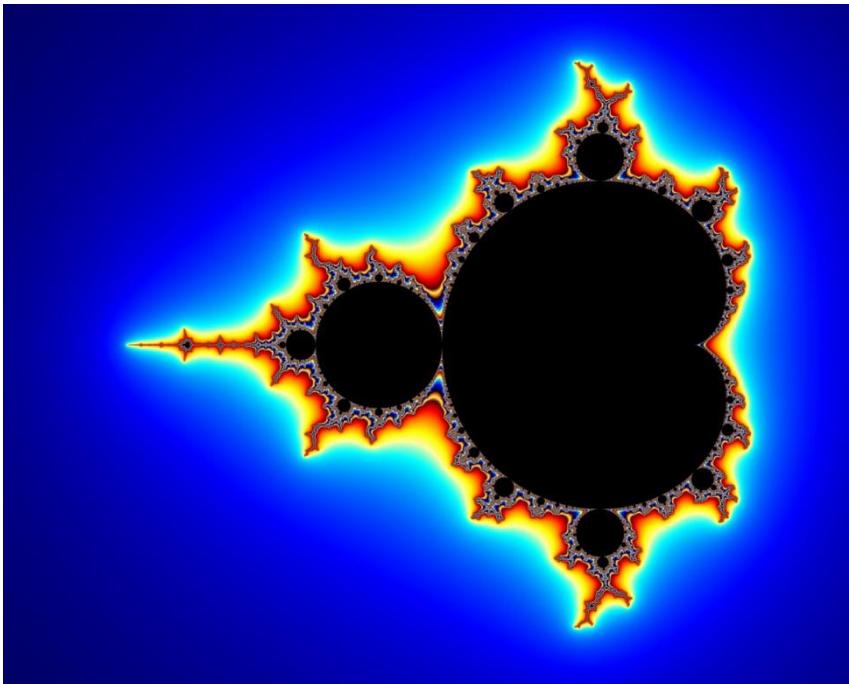
What about subroutines? We can use **orphaned directives**

```
!$omp parallel private(i,j,k)  
  
call calculate_manythings()  
  
!$omp end parallel
```



```
subroutine calculate_manythings()  
  
  !$omp do schedule(static)  
    do i=1,n  
      <calculations>  
    enddo  
  !$omp end do  
  
  return  
end subroutine
```

## Mandelbrot set area calculation



Mathematically it can be calculated exactly as:

$$A = \pi \left( 1 - \sum_{n=1}^{\infty} n b_n^2 \right)$$

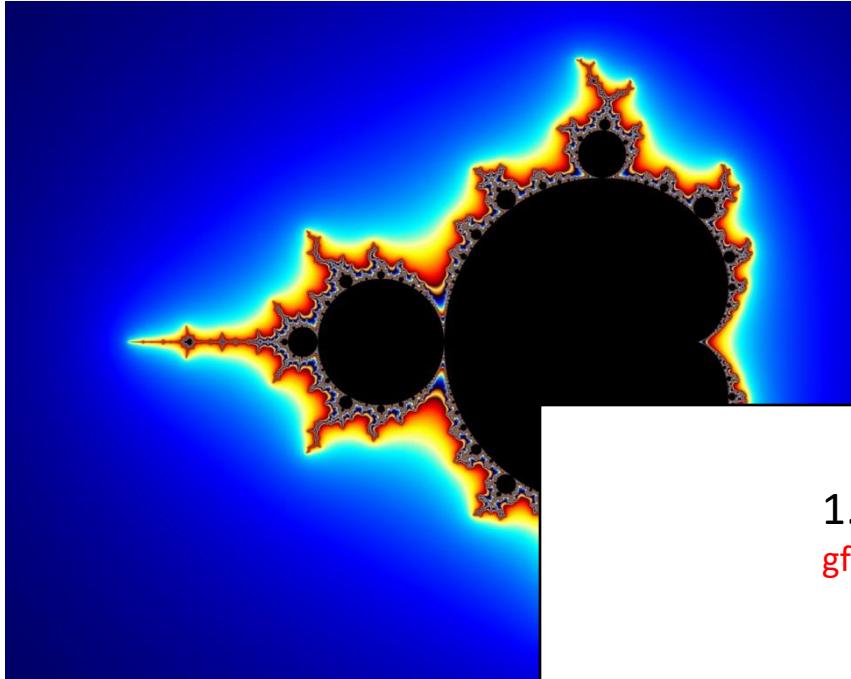
But this series converges VERY slowly. It needs  $10^{118}$  terms to get the first two digits, and  $10^{1181}$  to get the third!

The area obtained by pixel counting is:

**1.50659177  $\pm$  0.00000008**

(Munafo et al. OEIS A098403, 2000)

## Mandelbrot set area calculation



Mathematically it can be calculated exactly as:

$$A = \pi \left( 1 - \sum_{n=1}^{\infty} n b_n^2 \right)$$

But this series converges VERY slowly. It needs  $10^{118}$  terms to get the first two digits, and  $10^{1181}$  to get the third!

The area obtained by pixel counting is:

**1.50659177  $\pm$  0.00000008**

(Munafo et al. OEIS A098403, 2000)

1. Compile and execute the serial code

```
gfortran mandelbrot.f90 -o mandelbrot_serial  
time ./mandelbrot_serial
```

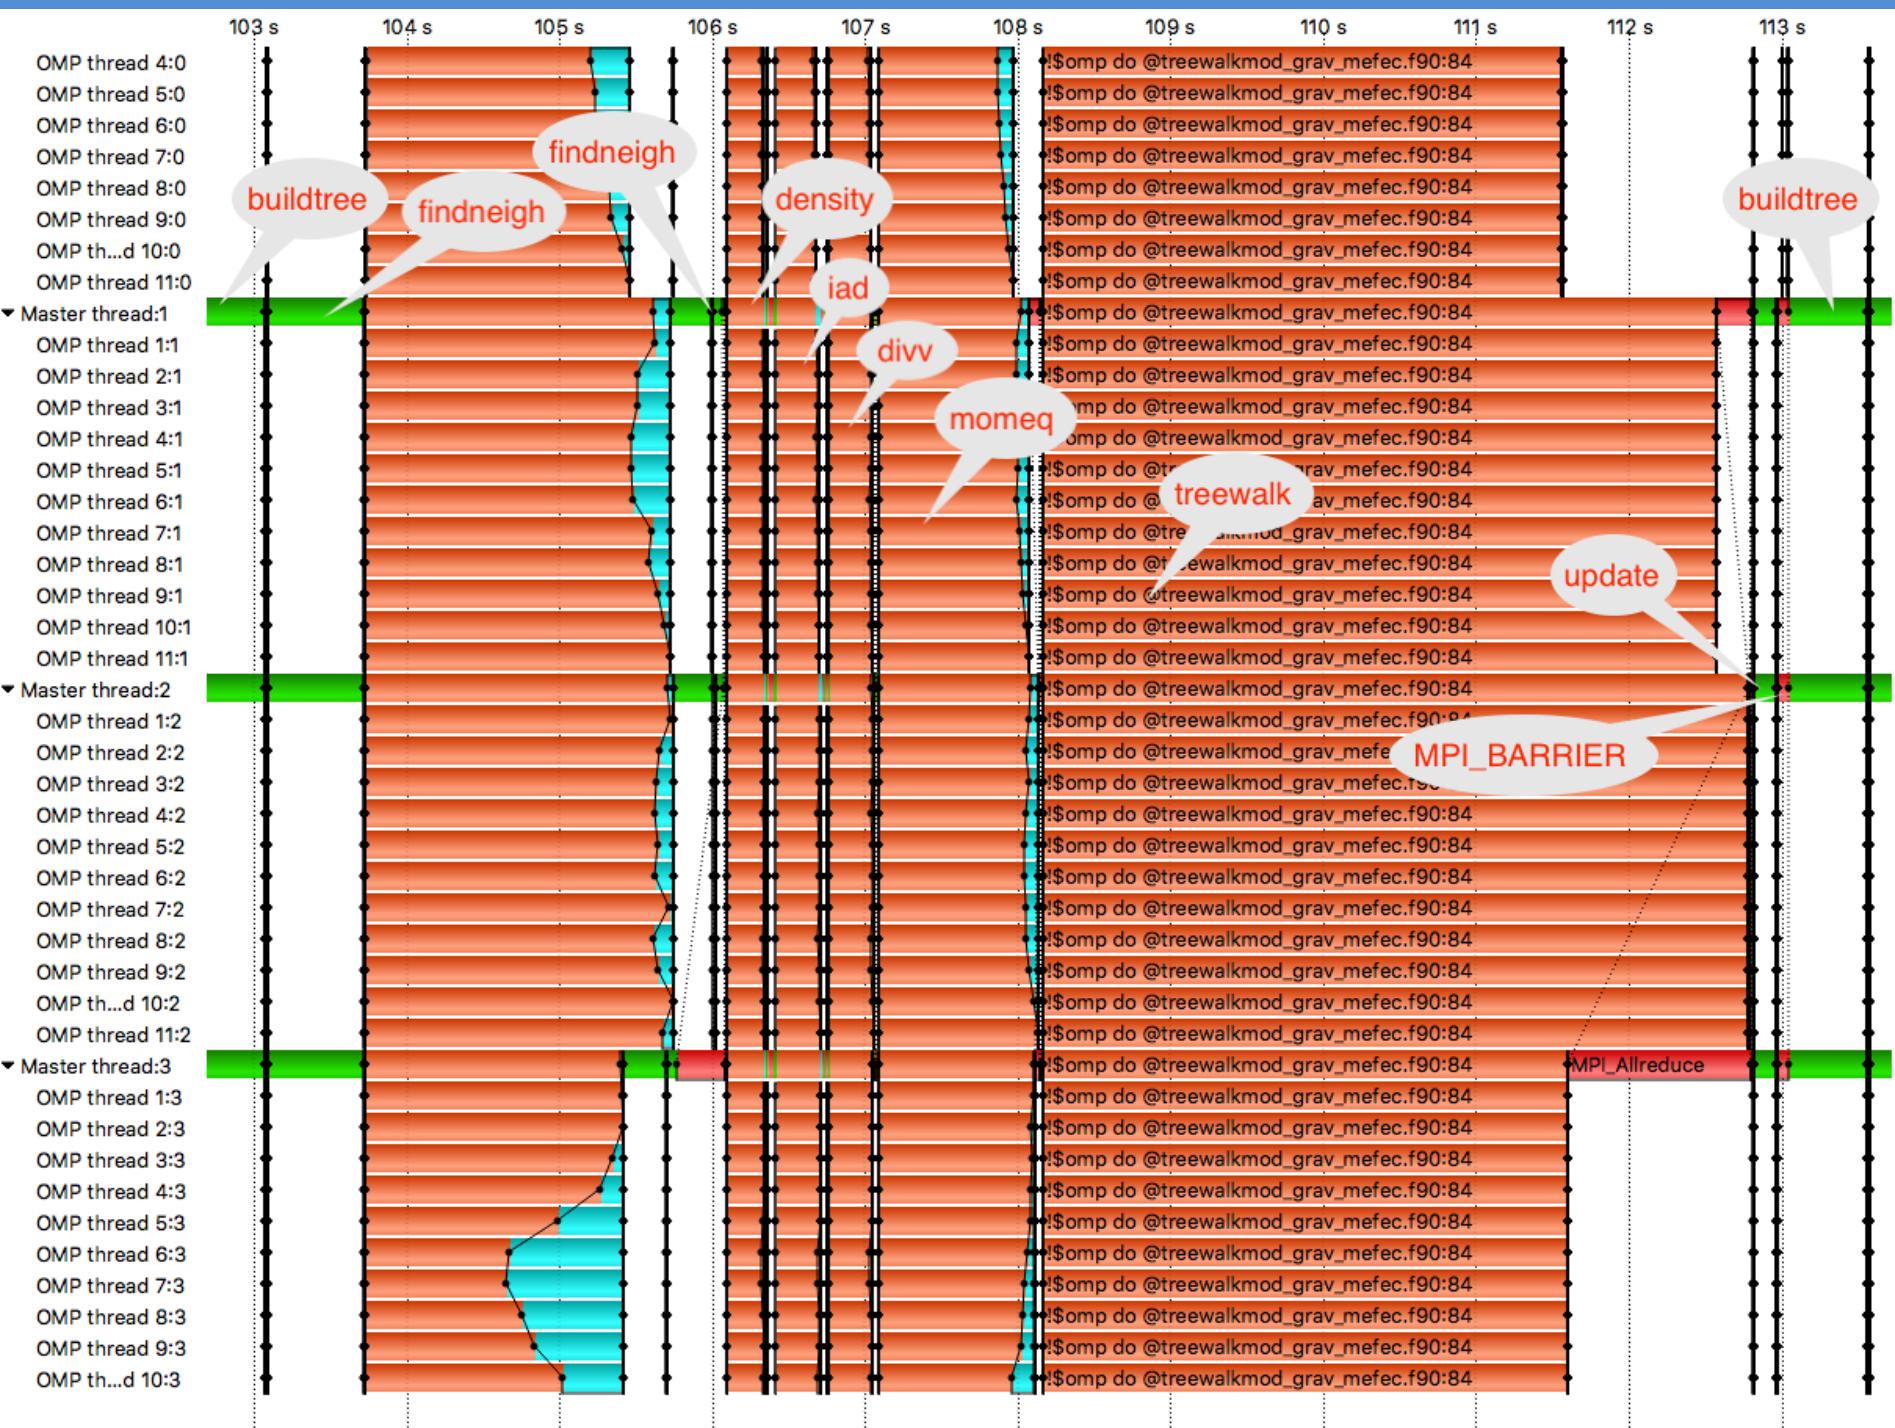
2. Parallelize the code with OpenMP

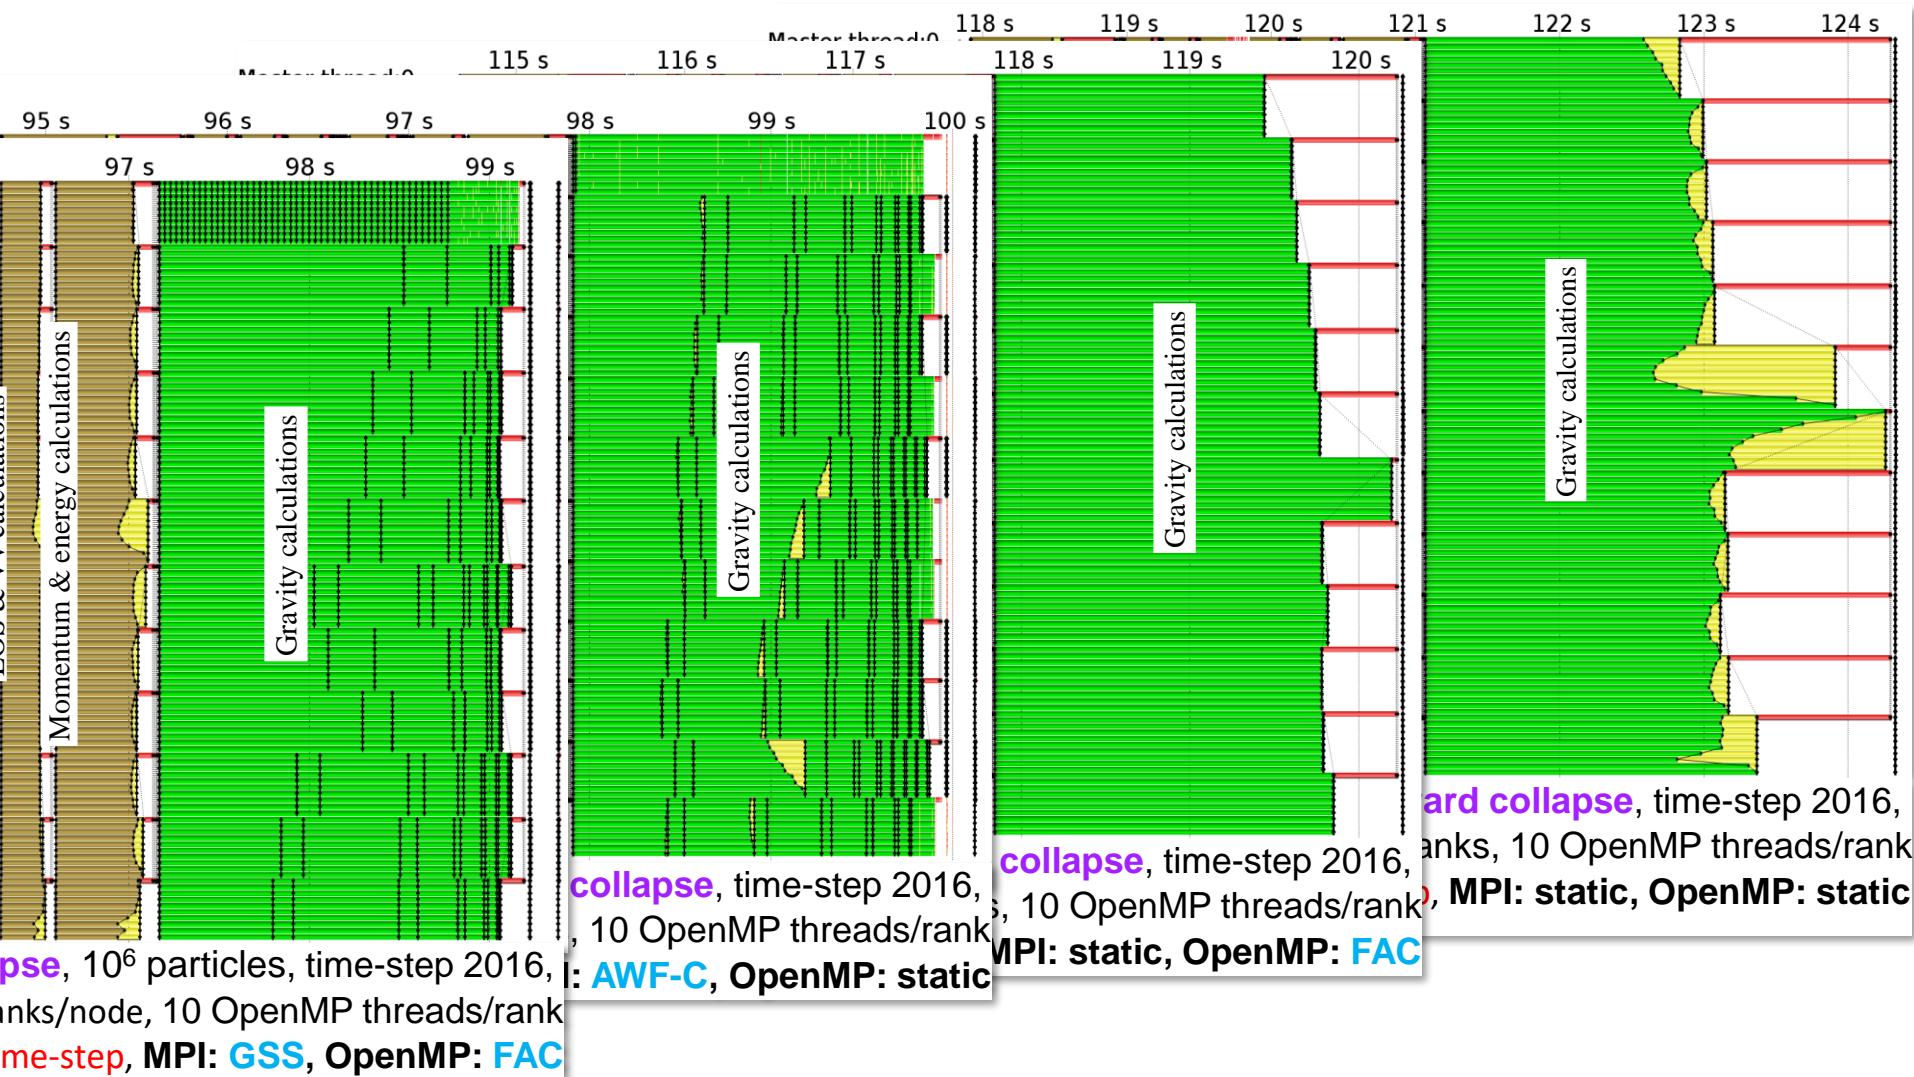
3. Compile and execute the parallel code

```
gfortran -fopenmp mandelbrot.f90 -o mandelbrot_parallel  
time ./mandelbrot_parallel
```

4. Change the number of threads (2-4) and check if it scales

```
export OMP_NUM_THREADS = <# of threads>
```





Two-level dynamic load balancing for high performance scientific applications  
 Mohammed, A. et al (Siam PP20)

Can we do any of this in Python? **Kind of...**

```
from multiprocessing import Process

def func1():
    j=0
    print ('func1: starting')
    for i in range(10000000000):
        j=j+1
    print ('func1: finishing',j)

def func2():
    j=0
    print ('func2: starting')
    for i in range(10000000000):
        j=j+1
    print ('func2: finishing',j)

def func3():
    j=0
    print ('func3: starting')
    for i in range(10000000000):
        j=j+1
    print ('func3: finishing',j)

if __name__ == '__main__':
    p1 = Process(target=func1)
    p1.start()
    p2 = Process(target=func2)
    p2.start()
    p3 = Process(target=func3)
    p3.start()
    p1.join()
    p2.join()
    p3.join()
```

```
top - 11:28:00 up 2 days, 2:02, 1 user, load average: 0,57, 0,56, 0,48
Tasks: 306 total, 4 running, 301 sleeping, 0 stopped, 1 zombie
%Cpu(s): 38,3 us, 0,4 sy, 0,0 ni, 61,1 id, 0,1 wa, 0,0 hi, 0,1 si, 0,0 st
KiB Mem : 8079052 total, 294168 free, 5072908 used, 2711976 buff/cache
KiB Swap: 19529724 total, 19506740 free, 22984 used. 1438376 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
26027	ruben	20	0	50896	8488	3072	R	100,0	0,1	0:09.43	python
26028	ruben	20	0	50896	8488	3072	R	100,0	0,1	0:09.43	python
26026	ruben	20	0	50896	8488	3072	R	99,0	0,1	0:09.43	python
24587	ruben	20	0	4736720	2,196g	2,119g	S	4,0	28,5	7:45.81	VirtualBox
1128	root	20	0	756332	231528	197116	S	1,0	2,9	11:15.62	Xorg
7796	ruben	20	0	1325324	184588	52036	S	1,0	2,3	21:43.70	skypeforli+
26032	ruben	20	0	437152	21932	18568	S	1,0	0,3	0:00.25	gnome-scre+
1	root	20	0	185332	4984	3008	S	0,0	0,1	0:01.58	svstemd

```
ruben@jarvis:~/test/parallelpython$ time python test.py
func1: starting
func2: starting
func3: starting
func3: finishing 10000000000
func2: finishing 10000000000
func1: finishing 10000000000
[1]+ Done                  emacs test.py

real    0m36.208s
user    1m50.284s
sys     0m0.544s
```

Can we do any of this in Python? **Kind of...**

```
from multiprocessing import Process

def func1():
    j=0
    print ('func1: starting')
    for i in range(1000000000):
        j=j+1
    print ('func1: finishing',j)

def func2():
    j=0
    print ('func2: starting')
    for i in range(1000000000):
        j=j+1
    print ('func2: finishing',j)

def func3():
    j=0
    print ('func3: starting')
    for i in range(1000000000):
        j=j+1
    print ('func3: finishing',j)

def runInParallel(*fns):
    proc = []
    for fn in fns:
        p = Process(target=fn)
        p.start()
        proc.append(p)
    for p in proc:
        p.join()

if __name__ == '__main__':
    runInParallel(func1, func2, func3)
```

```
top - 11:28:00 up 2 days, 2:02, 1 user, load average: 0,57, 0,56, 0,48
Tasks: 306 total, 4 running, 301 sleeping, 0 stopped, 1 zombie
%Cpu(s): 38,3 us, 0,4 sy, 0,0 ni, 61,1 id, 0,1 wa, 0,0 hi, 0,1 si, 0,0 st
KiB Mem : 8079052 total, 294168 free, 5072908 used, 2711976 buff/cache
KiB Swap: 19529724 total, 19506740 free, 22984 used. 1438376 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
26027	ruben	20	0	50896	8488	3072	R	100,0	0,1	0:09.43	python
26028	ruben	20	0	50896	8488	3072	R	100,0	0,1	0:09.43	python
26026	ruben	20	0	50896	8488	3072	R	99,0	0,1	0:09.43	python
24587	ruben	20	0	4736720	2,196g	2,119g	S	4,0	28,5	7:45.81	VirtualBox
1128	root	20	0	756332	231528	197116	S	1,0	2,9	11:15.62	Xorg
7796	ruben	20	0	1325324	184588	52036	S	1,0	2,3	21:43.70	skypeforli+
26032	ruben	20	0	437152	21932	18568	S	1,0	0,3	0:00.25	gnome-scre+
1	root	20	0	185332	4984	3008	S	0,0	0,1	0:01.58	svstemd

```
ruben@jarvis:~/test/parallelpython$ time python test.py
func1: starting
func2: starting
func3: starting
func3: finishing 1000000000
func2: finishing 1000000000
func1: finishing 1000000000
[1]+ Done                  emacs test.py

real    0m36.208s
user    1m50.284s
sys     0m0.544s
```

<http://stackabuse.com/parallel-processing-in-python/>

We want to sum up the integer half-value of the first 3e8 integers.

```
from multiprocessing import Process,Manager

def func1(id,results):
    j=0
    print('func1: starting')
    for i in range(100000000):
        j=j+i/2
    print('func1: finishing')
    results[id]=j

def func2(id,results):
    j=0
    print('func2: starting')
    for i in range(100000001,200000000):
        j=j+i/2
    print('func2: finishing')
    results[id]=j

def func3(id,results):
    j=0
    print('func3: starting')
    for i in range(200000001,300000000):
        j=j+i/2
    print('func3: finishing')
    results[id]=j

def runInParallel(*fns):
    proc=[]
    for i,fn in enumerate(fns):
        p = Process(target=fn,args=(i,results))
        p.start()
        proc.append(p)
    for p in proc:
        p.join()
    print(results)
    print(sum(results.values()))

if __name__ == '__main__':
    results=Manager().dict()
    runInParallel (func1,func2,func3)
```

```
[cabezón@login10 openmp]$ time python python_multiproc_2.py
func1: starting
func2: starting
func3: starting
func3: finishing
func2: finishing
func1: finishing
[0: 2499999950000000, 1: 7499999900000000, 2: 1249999985000000]
2249999700000000
real    0m12.819s
user    0m27.136s
sys     0m8.865s
[cabezón@login10 openmp]$
```

Note that threads finish in an unstructured way!

We store the results in a dictionary and sum them up.  
(mimicking a 'reduce' in OpenMP)

More elegant way to do this is using 'numba'. <https://numba.pydata.org/>



Adding a decorator numba compiles the code on-the-fly creating an optimized machine code

```
from numba import jit

@jit() ←
def func1():
    j=0
    print('func1: starting')
    for i in range(300000000):
        j=j+i/2
    print('func1: finishing',j)

if __name__ == '__main__':
    func1()
```

```
[cabezón@login10 openmp]$ time python python_multiproc_2.py
func1: starting
func2: starting
func3: starting
func3: finishing
func2: finishing
func1: finishing
[0: 2499999950000000, 1: 7499999900000000, 2: 1249999985000000]
2249999700000000

real 0m12.819s
user 0m27.130s
sys 0m8.865s
[cabezón@login10 openmp]$
```

```
[cabezón@login10 openmp]$ source numbatest/bin/activate
(numbatest) [cabezón@login10 openmp]$ time python python_numba.py
func1: starting
func1: finishing 2.24999988355443e+16

real 0m1.142s
user 0m1.205s
sys 0m1.294s
(numbatest) [cabezón@login10 openmp]$
```

More elegant way to do this is using 'numba'. <https://numba.pydata.org/>



Adding a decorator numba compiles the code on-the-fly creating an optimized machine code

```
from numba import jit

@jit() ←
def func1():
    j=0
    print('func1: starting')
    for i in range(30000000
                    j=j+i/2
    print('func1: finishing')

if __name__ == '__main__':
    func1()
```

```
[cabezon@login10 openmp]$ time python python_multiproc_2.py
func1: starting
func2: starting
func3: starting
func3: finishing
func2: finishing
func1: finishing
[0: 2499999950000000, 1: 7499999900000000, 2: 1249999985000000]
2249999700000000
real 0m12.819s
user 0m27.130s
sys 0m8.865s
```

We need to install numba. For that we will install a virtual environment of Python:

```
virtualenv testnumba
source testnumba/bin/activate
```

2. Install numba  

```
pip install numba
```

3. Execute the serial version with the numba decorator  
(don't forget to import jit at the beginning of the Python script)  

```
time python python_serial.py
```

More elegant way to do this is using 'numba'. <https://numba.pydata.org/>



We are using just 1 core with numba, but we can also use all cores in the machine, just like with OpenMP, with prange:

```
from numba import jit,prange
@jit(parallel=True)
def func1():
    j=0
    print('func: starting (from 1 to 300000000)')
    for i in prange(1,30000000001):
        j=j+i/2
    print('func1: finishing',j)

if __name__ == '__main__':
    func1()
```

```
cabezón@sh101:~/openmp_course $time python python_serial.py
func: starting (from 1 to 30000000000)
func1: finishing 2.2500000000007956e+20

real 0m31.206s
user 0m32.340s
sys 0m2.832s
cabezón@sh101:~/openmp_course $
```

```
cabezón@sh101:~/openmp_course $time python python_serial.py
func: starting (from 1 to 30000000000)
func1: finishing 2.2500000000005246e+20

real 0m1.781s
user 0m41.160s
sys 0m3.124s
cabezón@sh101:~/openmp_course $
```

**NOTE:** we are doing x100 more iterations in this example than in the previous one so that we can actually see differences in the times between 1 and many cores.

```
saktho00@worker04:~/temp/numbatest
File Edit View Search Terminal Help
genotypes[cellID,2+targetoffset]=maxsnps#genotypes[c
ellID,3+targetoffset] + 1
#genotypes[cellThreadID,1+targetoffset]=gen
#genotypes[cellID,2+targetoffset]=genotypes[cellID,2+targetoffse
t] + 1
    cuda.syncthreads()
    #cuda.atomic.add(syncObj,0,1)
    #while(syncObj[0]<N*(numbIt+1+gen)):
    #    1+1
@jit
def pairwisedist(genotypes):
    difflist = []
    N=int(genotypes.shape[0])
    L=int(genotypes.shape[1]/2)
    for i in range(N-1):
        for j in range(i+1,N):
            delta=0
            diff = 0
            for k in range(L):
                delta = genotypes[i,k] ^ genotypes[j,k]
                for m in range(64):
                    if int64(delta) & 1 == 1:
                        diff = diff + 1
                delta = math.floor(delta / 2)
                if delta == 0:
                    break
            difflist.append(diff/(L*64))
    return(difflist)

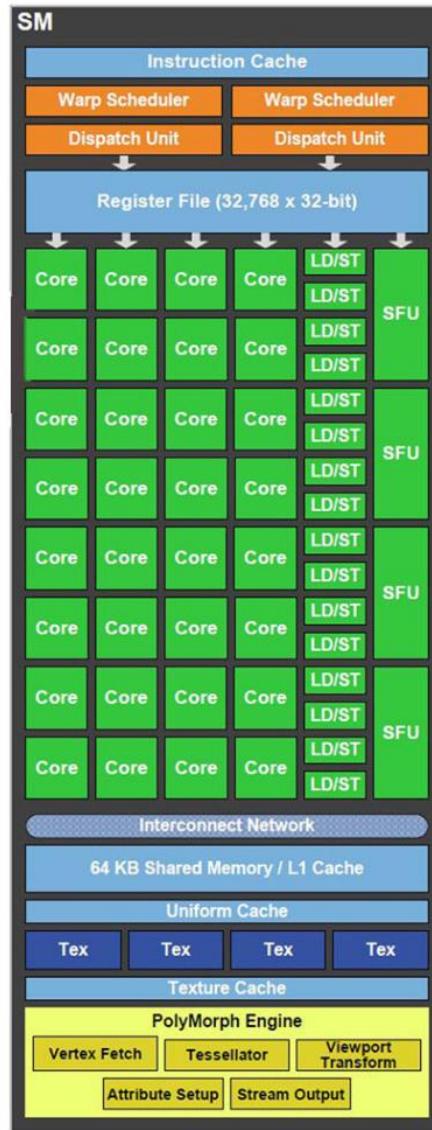
@jit
def closedist(genotypes):
    "recombSim.py" 329L, 11141C written
```

Thomas Sakoparnig

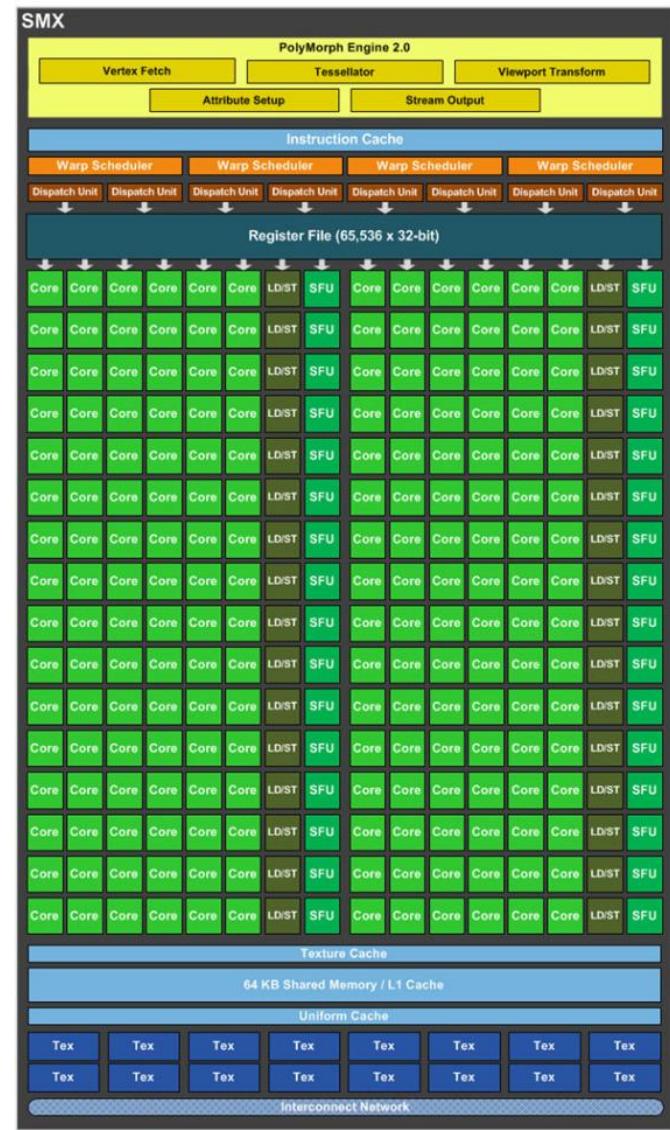
```
saktho00@worker04:~/temp/numbatest
File Edit View Search Terminal Help
[saktho00@sgi27 numbatest]$ vi recombSim.py
[saktho00@sgi27 numbatest]$ python recombSim.py
ttime: 3.4206590335816145 ←
0.0170744024235
[saktho00@sgi27 numbatest]$ vi recombSim.py
[saktho00@sgi27 numbatest]$ python recombSim.py
ttime: 2637.8457502331585 ←
0.0178537920918
[saktho00@sgi27 numbatest]$
```

# What about GPU?

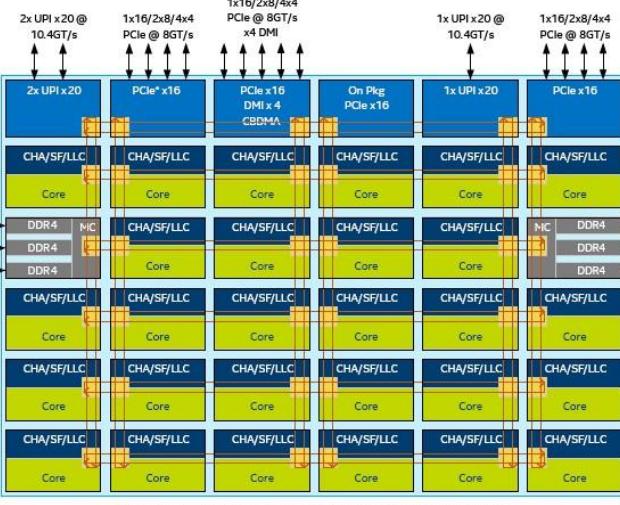
Tesla SM unit (2007)



Fermi SMX unit (2011)



Intel Skylake (2017)

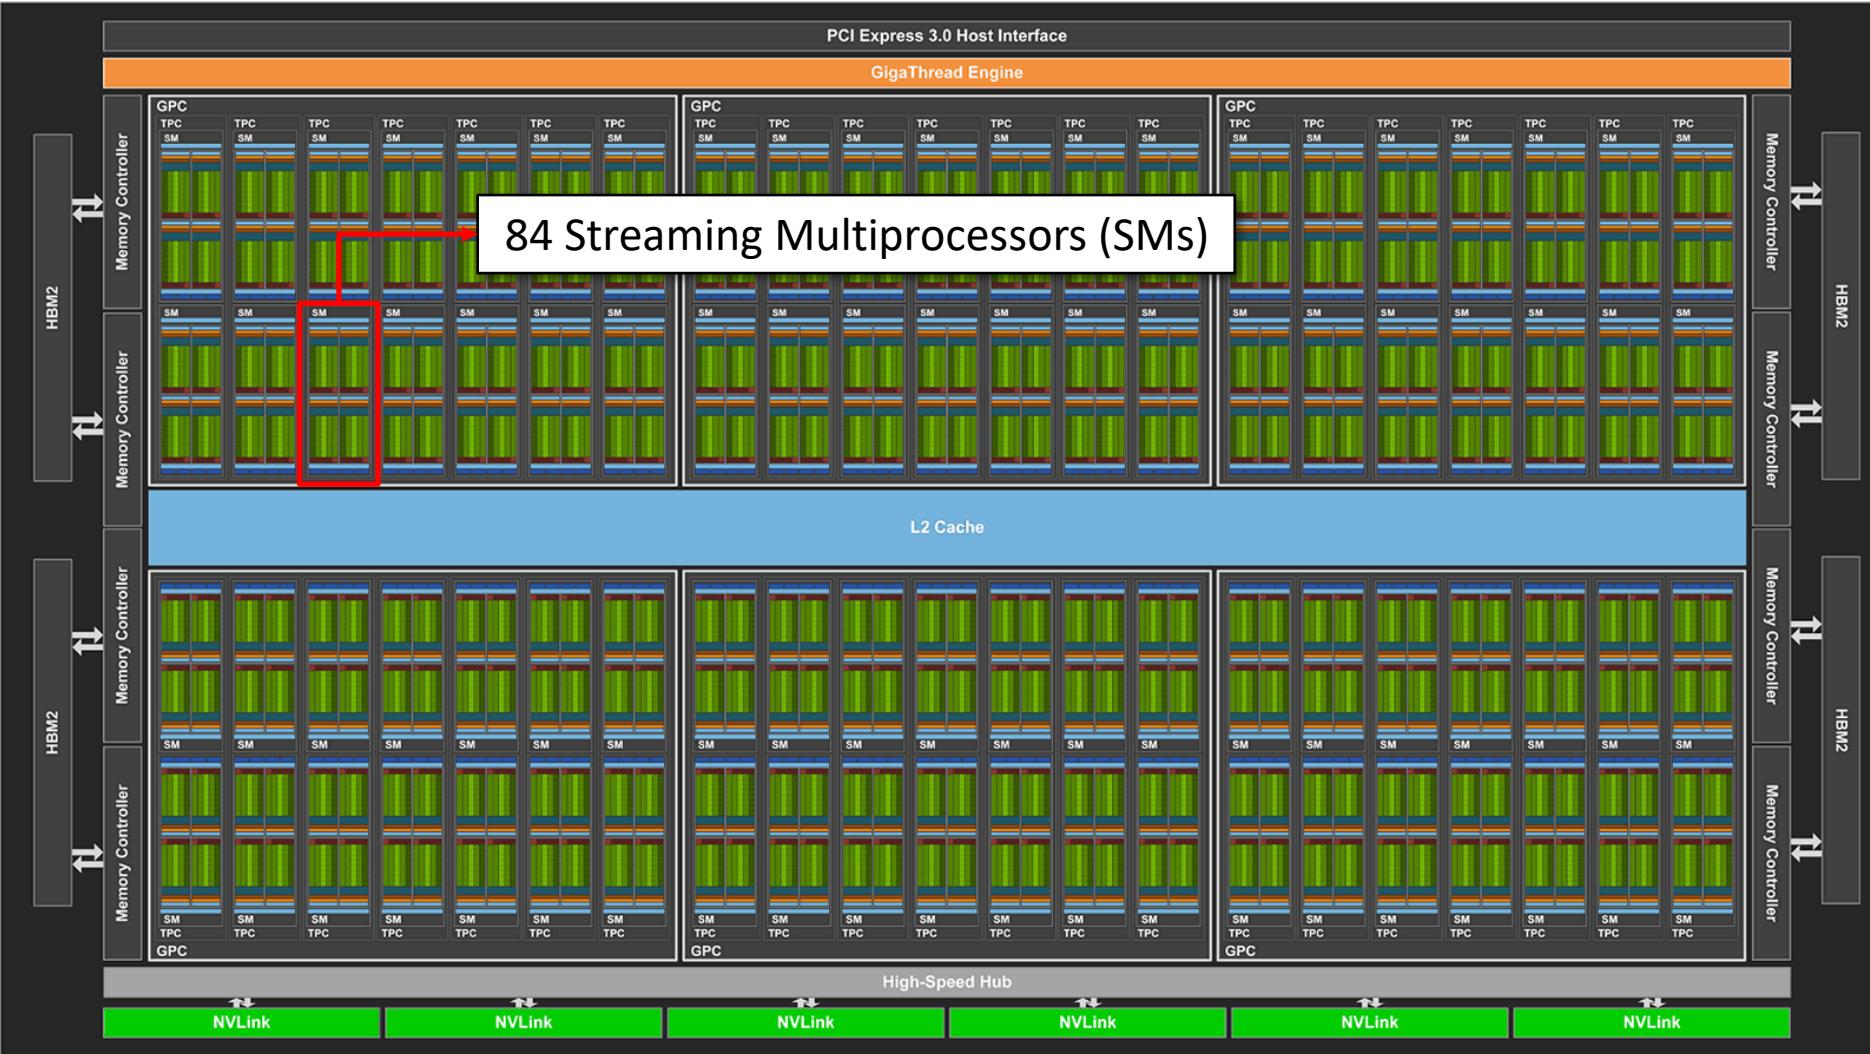


CHA – Caching and Home Agent ; SF – Snoop Filter; LLC – Last Level Cache;  
Core – Skylake-SP Core; UPI – Intel® UltraPath Interconnect

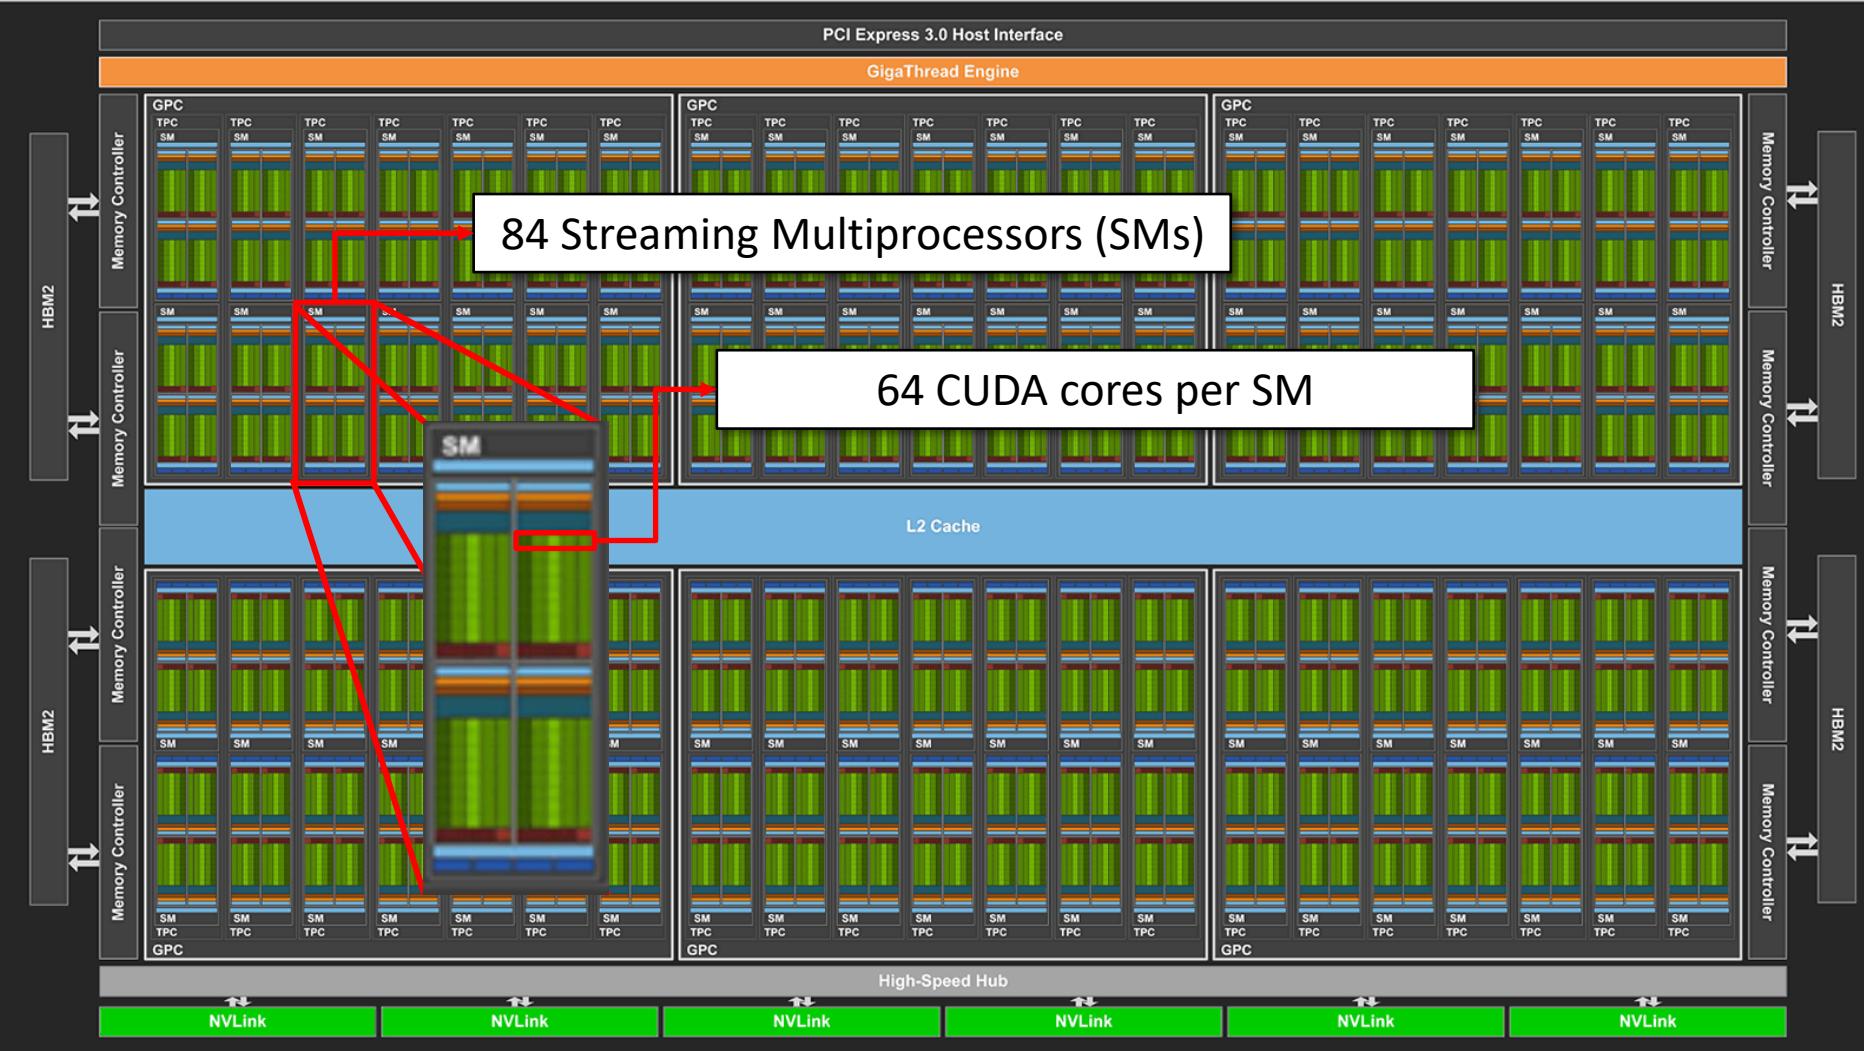
Volta (2017)



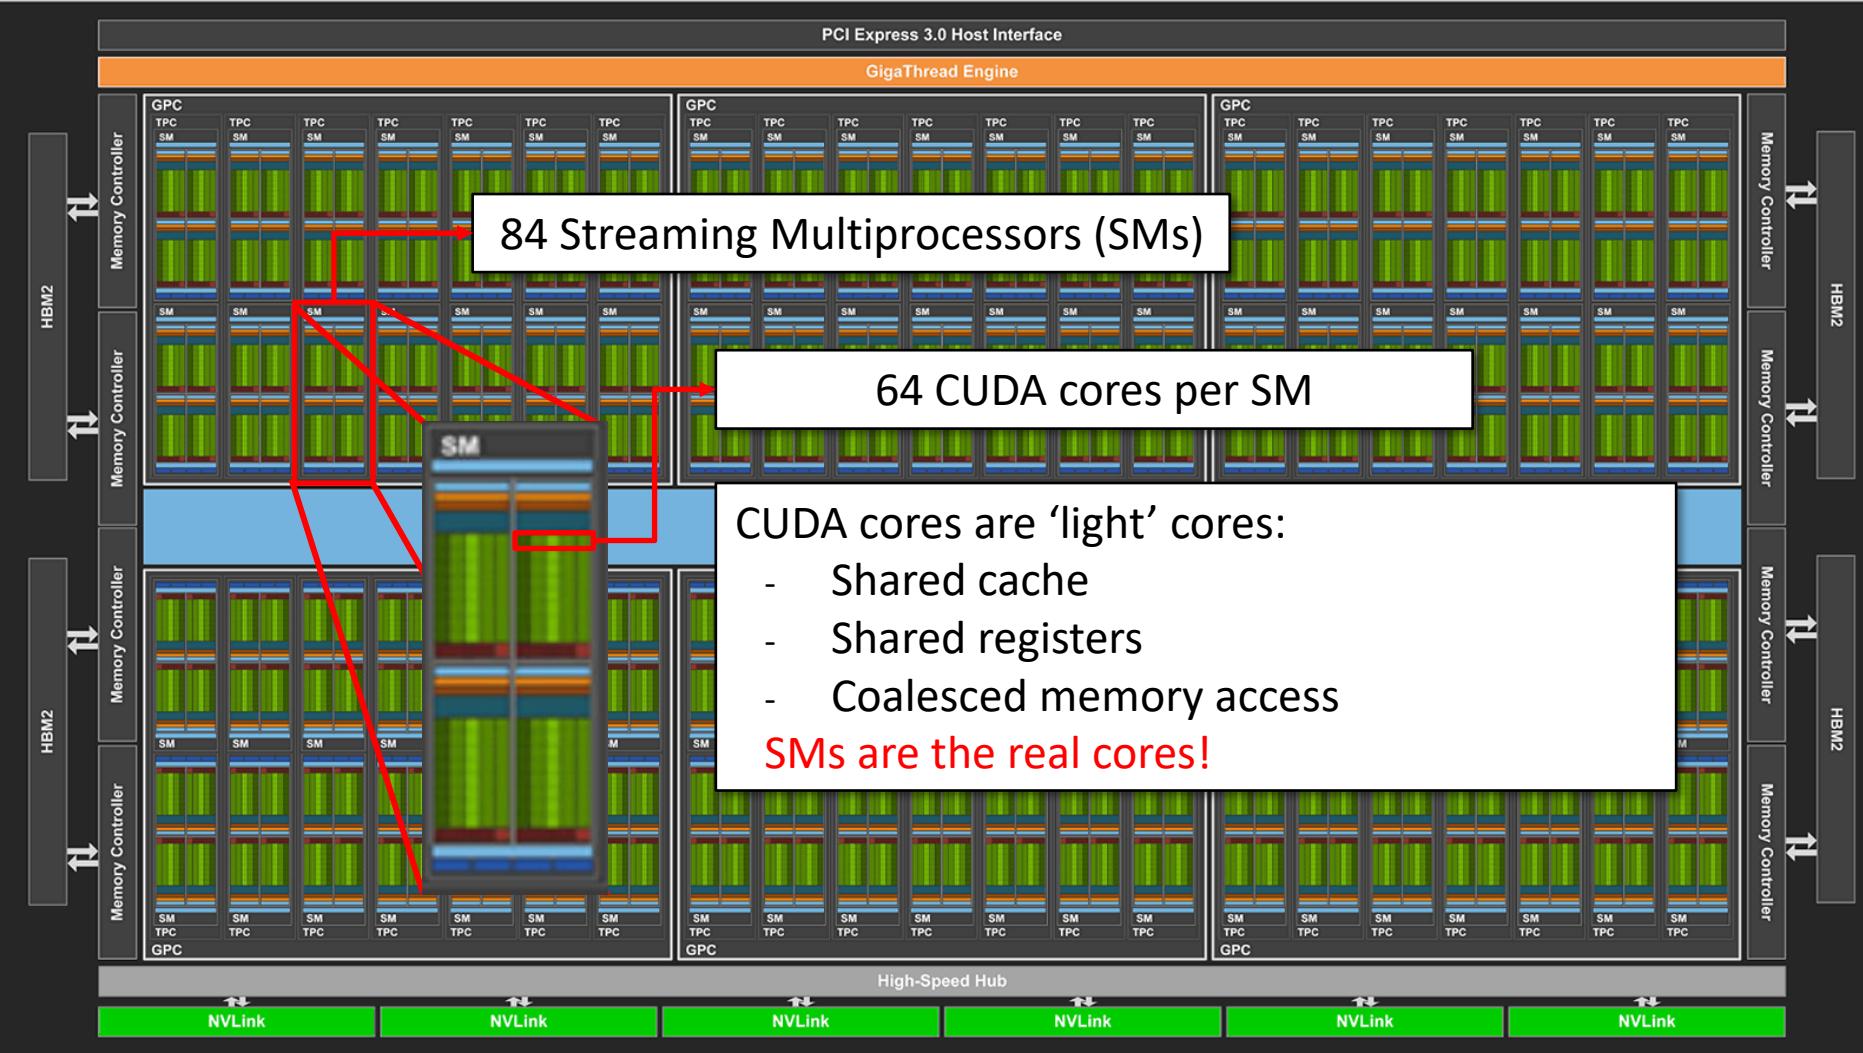
## Volta (2017)



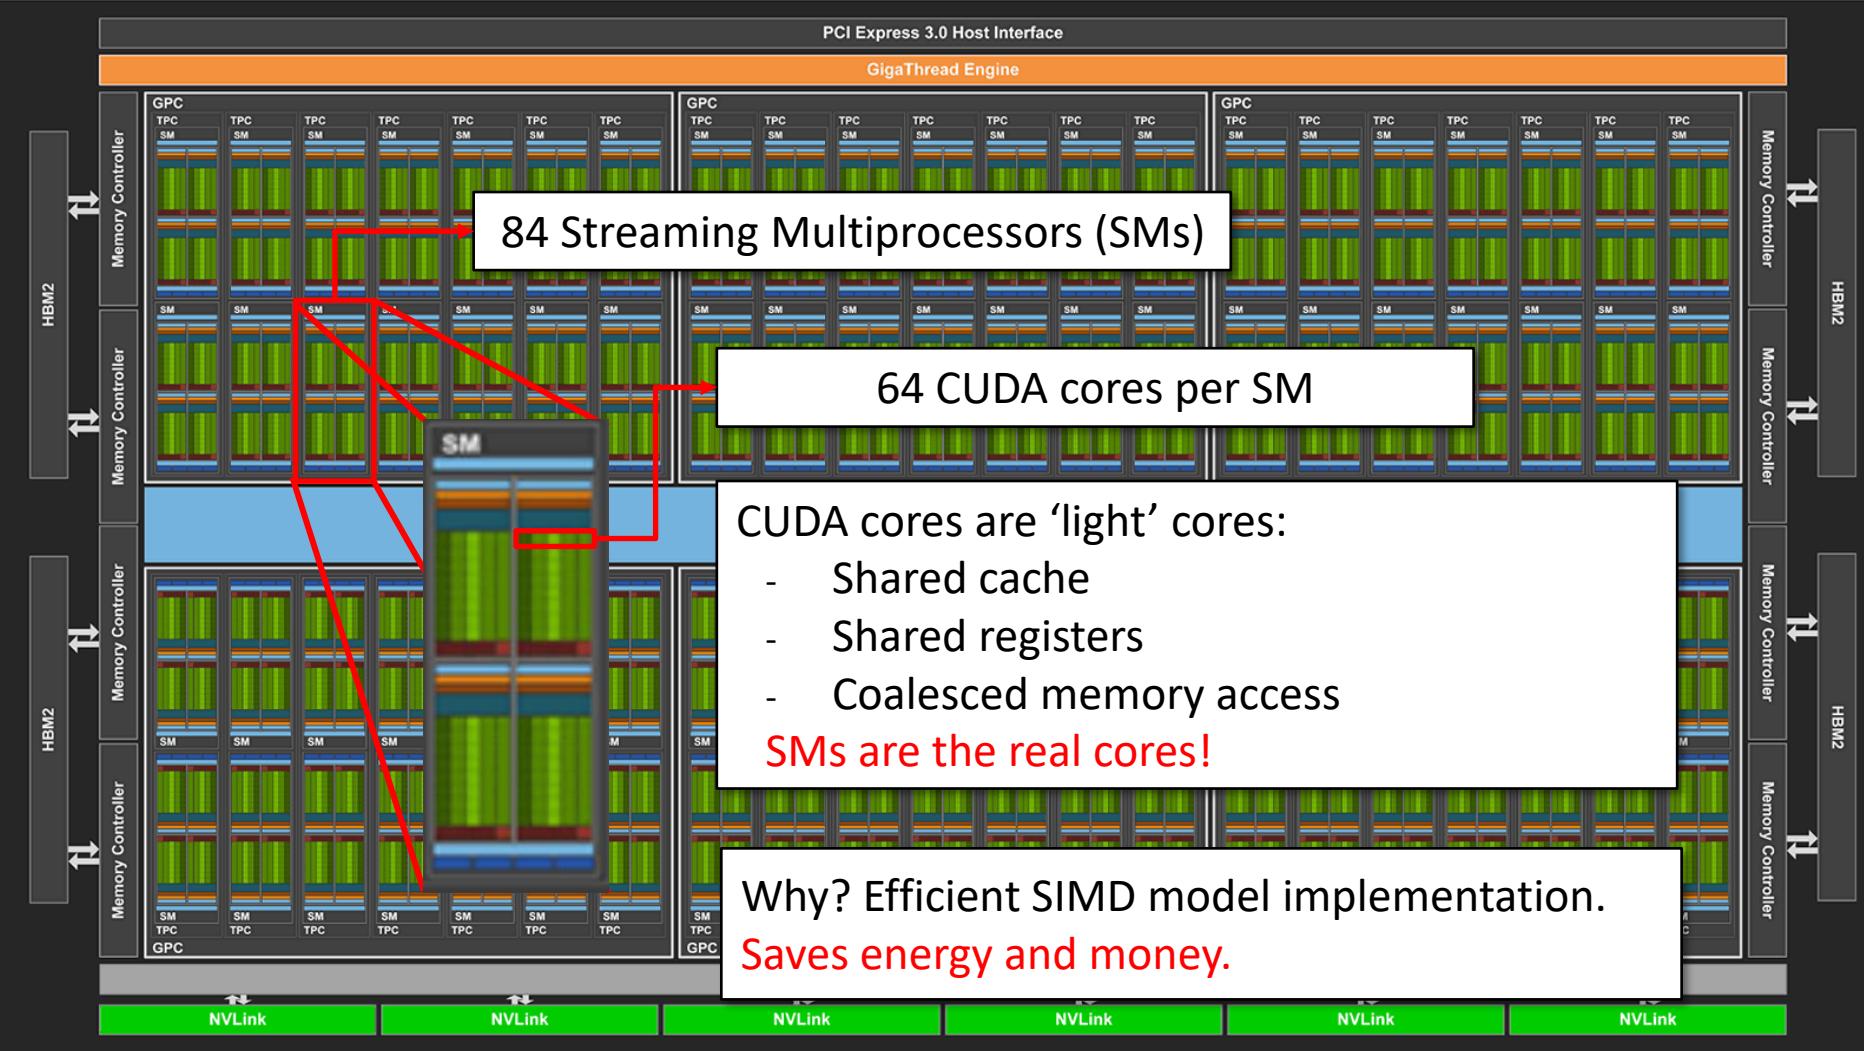
## Volta (2017)



Volta (2017)



## Volta (2017)



Since version 4.0, openMP can do GPU offloading.

```
int A[1] = {-1};  
#pragma omp target  
{  
    A[0] = omp_is_initial_device();  
}  
if (!A[0]) {  
    printf("Able to use offloading!\n");  
}
```

You can control the transfer of data

omp\_is\_initial\_device() returns 0 when called from the accelerator

```
#pragma omp target data map(to: b[:n])  
{  
    // 'b' copied onto device  
  
    // This region is offloaded to device  
    #pragma omp target map(tofrom: a[:n])  
{  
        // 'a' copied onto device  
        #pragma omp teams distribute  
        for(int ii = 0; ii < n; ++ii)  
        {  
            a[ii] = a[ii] + alpha * b[ii];  
        }  
        // 'a' copied back  
  
        // ... potentially more target regions  
    }  
    // 'b' is discarded by device
```

Since version 4.0, openMP can do **GPU offloading**.

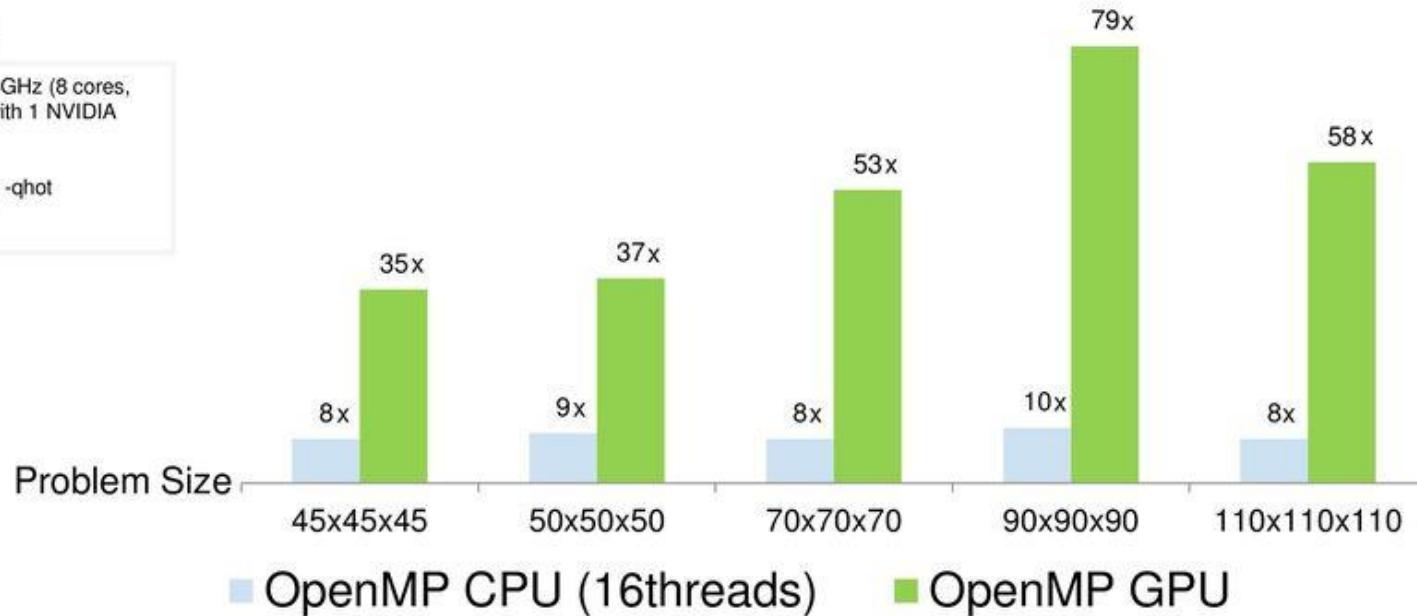
## Performance Results – End to End

**LULESH – Speedup Over Serial  
(Higher is Better)**

### Test Specs

2 Power8 sockets @ 4GHz (8 cores, with 8 threads each) with 1 NVIDIA Pascal P100 GPU.

Compiler Options: -O3 -qhot  
-qsmp=omp -qoffload\*  
\* Where applicable

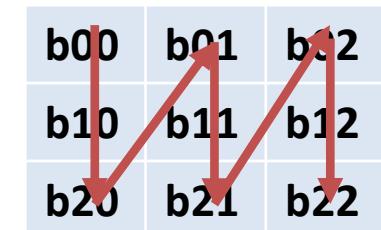


## Full example: matrix multiplication

Sub-optimal memory access will kill the performance...

The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```

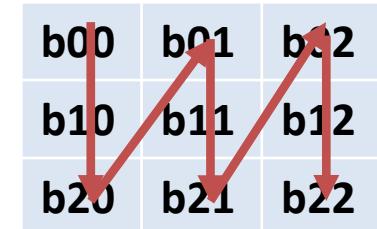


## Full example: matrix multiplication

Sub-optimal memory access will kill the performance...

The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



- Processors **load data into their cache** for fast reuse
- They always load a bunch of data at the same time
- They do **branch-prediction** and **pre-fetching**!

However here we are not helping him...

Processor cache			Cache Misses: 0					
a00	a01	a02	a10	a11	a12	a20	a21	a22
b00	b01	b02	b10	b11	b12	b20	b21	b22
c00	c01	c02	c10	c11	c12	c20	c21	c22

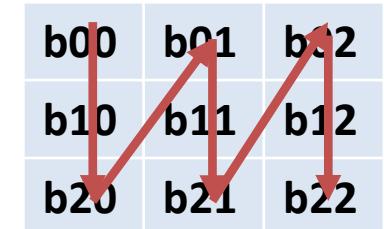
i	j	k
0	0	0

## Full example: matrix multiplication

Sub-optimal memory access will kill the performance...

The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



- Processors **load data into their cache** for fast reuse
- They always load a bunch of data at the same time
- They do **branch-prediction** and **pre-fetching**!

However here we are not helping him...

Processor cache			Cache Misses: 1					
a <sub>00</sub>	a <sub>01</sub>	a <sub>02</sub>	a <sub>10</sub>	a <sub>11</sub>	a <sub>12</sub>	a <sub>20</sub>	a <sub>21</sub>	a <sub>22</sub>
b <sub>00</sub>	b <sub>01</sub>	b <sub>02</sub>	b <sub>10</sub>	b <sub>11</sub>	b <sub>12</sub>	b <sub>20</sub>	b <sub>21</sub>	b <sub>22</sub>
c <sub>00</sub>	c <sub>01</sub>	c <sub>02</sub>	c <sub>10</sub>	c <sub>11</sub>	c <sub>12</sub>	c <sub>20</sub>	c <sub>21</sub>	c <sub>22</sub>

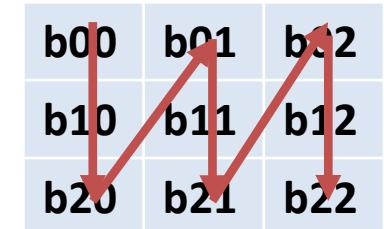
i	j	k
0	0	1

## Full example: matrix multiplication

Sub-optimal memory access will kill the performance...

The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



- Processors **load data into their cache** for fast reuse
- They always load a bunch of data at the same time
- They do **branch-prediction** and **pre-fetching!**

However here we are not helping him...

Processor cache			Cache Misses: 2					
a00	a01	a02	a10	a11	a12	a20	a21	a22
b00	b01	b02	b10	b11	b12	b20	b21	b22
c00	c01	c02	c10	c11	c12	c20	c21	c22

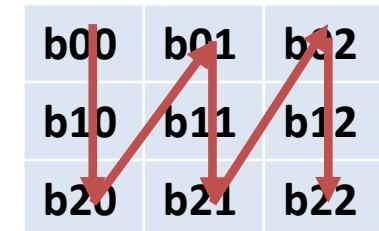
i	j	k
0	0	2

## Full example: matrix multiplication

Sub-optimal memory access will kill the performance...

The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



- Processors **load data into their cache** for fast reuse
- They always load a bunch of data at the same time
- They do **branch-prediction** and **pre-fetching**!

However here we are not helping him...

Processor cache										Cache Misses: 3								
a00	a01	a02	a10	a11	a12	a20	a21	a22		b00	b01	b02	b10	b11	b12	b20	b21	b22
c00	c01	c02	c10	c11	c12	c20	c21	c22										

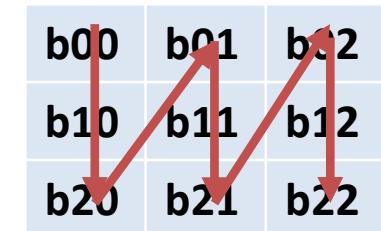
i	j	k
0	1	0

## Full example: matrix multiplication

Sub-optimal memory access will kill the performance...

The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



- Processors **load data into their cache** for fast reuse
- They always load a bunch of data at the same time
- They do **branch-prediction** and **pre-fetching**!

However here we are not helping him...

Processor cache			Cache Misses: 4					
a00	a01	a02	a10	a11	a12	a20	a21	a22
b00	b01	b02	b10	b11	b12	b20	b21	b22
c00	c01	c02	c10	c11	c12	c20	c21	c22

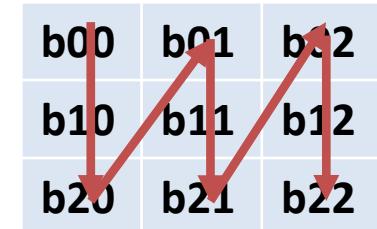
i	j	k
0	1	1

## Full example: matrix multiplication

Sub-optimal memory access will kill the performance...

The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



- Processors **load data into their cache** for fast reuse
- They always load a bunch of data at the same time
- They do **branch-prediction** and **pre-fetching**!

However here we are not helping him...

Processor cache			Cache Misses: 5					
a00	a01	a02	a10	a11	a12	a20	a21	a22
b00	b01	b02	b10	b11	b12	b20	b21	b22
c00	c01	c02	c10	c11	c12	c20	c21	c22

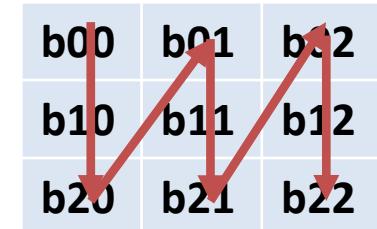
i	j	k
0	1	2

## Full example: matrix multiplication

Sub-optimal memory access will kill the performance...

The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



- Processors **load data into their cache** for fast reuse
- They always load a bunch of data at the same time
- They do **branch-prediction** and **pre-fetching**!

However here we are not helping him...

Processor cache			Cache Misses: 6					
a00	a01	a02	a10	a11	a12	a20	a21	a22
b00	b01	b02	b10	b11	b12	b20	b21	b22
c00	c01	c02	c10	c11	c12	c20	c21	c22

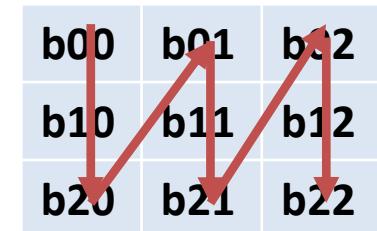
i	j	k
0	2	0

## Full example: matrix multiplication

Sub-optimal memory access will kill the performance...

The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



- Processors **load data into their cache** for fast reuse
- They always load a bunch of data at the same time
- They do **branch-prediction** and **pre-fetching**!

However here we are not helping him...

Processor cache			Cache Misses: 7					
a00	a01	a02	a10	a11	a12	a20	a21	a22
b00	b01	b02	b10	b11	b12	b20	b21	b22
c00	c01	c02	c10	c11	c12	c20	c21	c22

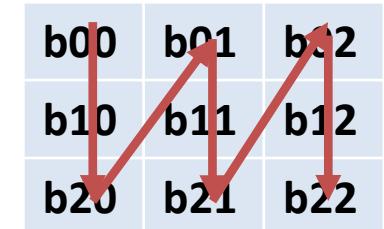
i	j	k
0	2	1

## Full example: matrix multiplication

Sub-optimal memory access will kill the performance...

The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



- Processors **load data into their cache** for fast reuse
- They always load a bunch of data at the same time
- They do **branch-prediction** and **pre-fetching**!

However here we are not helping him...

Processor cache			Cache Misses: 8					
a00	a01	a02	a10	a11	a12	a20	a21	a22
b00	b01	b02	b10	b11	b12	b20	b21	b22
c00	c01	c02	c10	c11	c12	c20	c21	c22

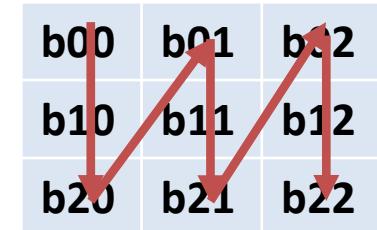
i	j	k
0	2	2

## Full example: matrix multiplication

Sub-optimal memory access will kill the performance...

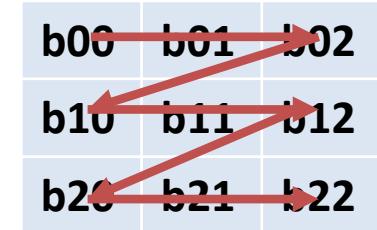
The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



The better way (row-major):

```
for(int i=0; i<n; ++i)
    for(int k=0; k<n; ++k)
        for(int j=0; j<n; ++j)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



Let's go...

a00	a01	a02	a10	a11	a12	a20	a21	a22
b00	b01	b02	b10	b11	b12	b20	b21	b22
c00	c01	c02	c10	c11	c12	c20	c21	c22

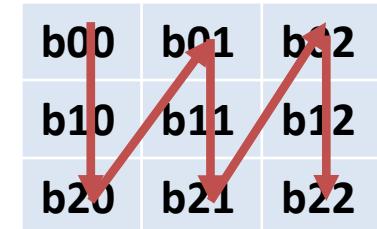
i	k	j
0	0	0

## Full example: matrix multiplication

Sub-optimal memory access will kill the performance...

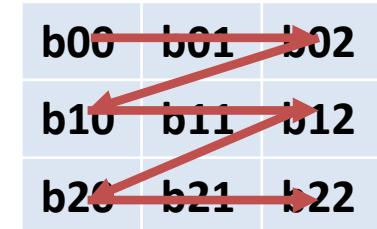
The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



The better way (row-major):

```
for(int i=0; i<n; ++i)
    for(int k=0; k<n; ++k)
        for(int j=0; j<n; ++j)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



Cache Misses: 0								
a00	a01	a02	a10	a11	a12	a20	a21	a22
b00	b01	b02	b10	b11	b12	b20	b21	b22
c00	c01	c02	c10	c11	c12	c20	c21	c22

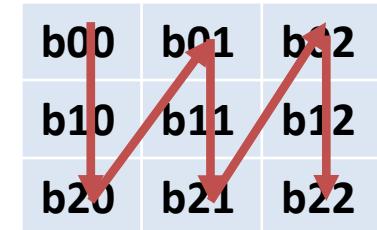
i	k	j
0	0	1

## Full example: matrix multiplication

Sub-optimal memory access will kill the performance...

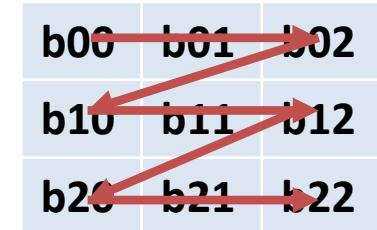
The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



The better way (row-major):

```
for(int i=0; i<n; ++i)
    for(int k=0; k<n; ++k)
        for(int j=0; j<n; ++j)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



Cache Misses: 0

a00	a01	a02	a10	a11	a12	a20	a21	a22
b00	b01	b02	b10	b11	b12	b20	b21	b22
c00	c01	c02	c10	c11	c12	c20	c21	c22

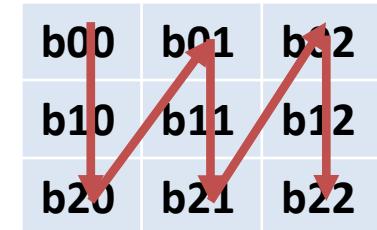
i	k	j
0	0	2

## Full example: matrix multiplication

Sub-optimal memory access will kill the performance...

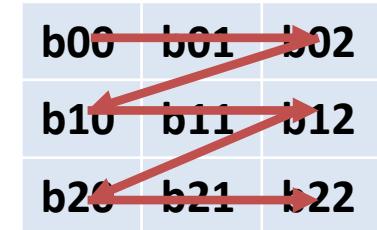
The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



The better way (row-major):

```
for(int i=0; i<n; ++i)
    for(int k=0; k<n; ++k)
        for(int j=0; j<n; ++j)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



Cache Misses: 1

a00	a01	a02	a10	a11	a12	a20	a21	a22
b00	b01	b02	b10	b11	b12	b20	b21	b22
c00	c01	c02	c10	c11	c12	c20	c21	c22

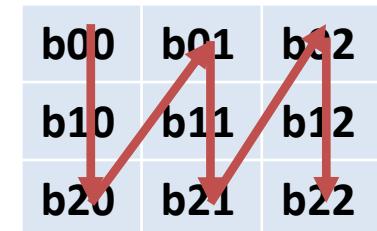
i	k	j
0	1	0

## Full example: matrix multiplication

Sub-optimal memory access will kill the performance...

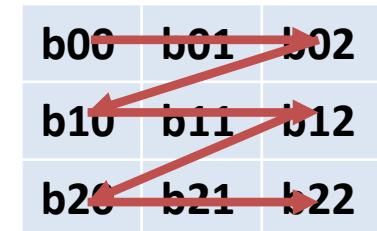
The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



The better way (row-major):

```
for(int i=0; i<n; ++i)
    for(int k=0; k<n; ++k)
        for(int j=0; j<n; ++j)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



Cache Misses: 1

a00	a01	a02	a10	a11	a12	a20	a21	a22
b00	b01	b02	b10	b11	b12	b20	b21	b22
c00	c01	c02	c10	c11	c12	c20	c21	c22

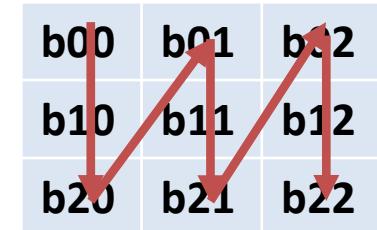
i	k	j
0	1	1

## Full example: matrix multiplication

Sub-optimal memory access will kill the performance...

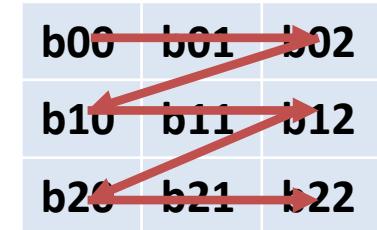
The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



The better way (row-major):

```
for(int i=0; i<n; ++i)
    for(int k=0; k<n; ++k)
        for(int j=0; j<n; ++j)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



Cache Misses: 1

a00	a01	a02	a10	a11	a12	a20	a21	a22
b00	b01	b02	b10	b11	b12	b20	b21	b22
c00	c01	c02	c10	c11	c12	c20	c21	c22

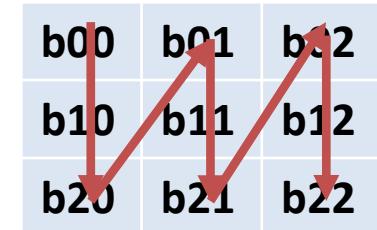
i	k	j
0	1	2

## Full example: matrix multiplication

Sub-optimal memory access will kill the performance...

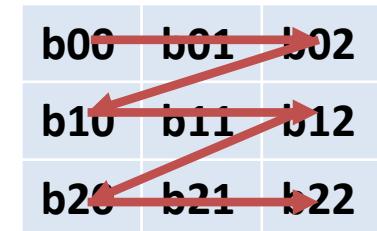
The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



The better way (row-major):

```
for(int i=0; i<n; ++i)
    for(int k=0; k<n; ++k)
        for(int j=0; j<n; ++j)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



Cache Misses: 2

a00	a01	a02	a10	a11	a12	a20	a21	a22
b00	b01	b02	b10	b11	b12	b20	b21	b22
c00	c01	c02	c10	c11	c12	c20	c21	c22

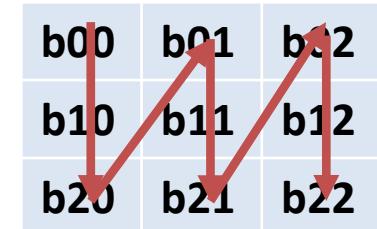
i	k	j
0	2	0

## Full example: matrix multiplication

Sub-optimal memory access will kill the performance...

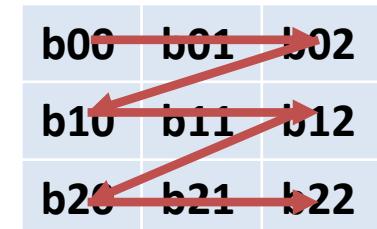
The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```

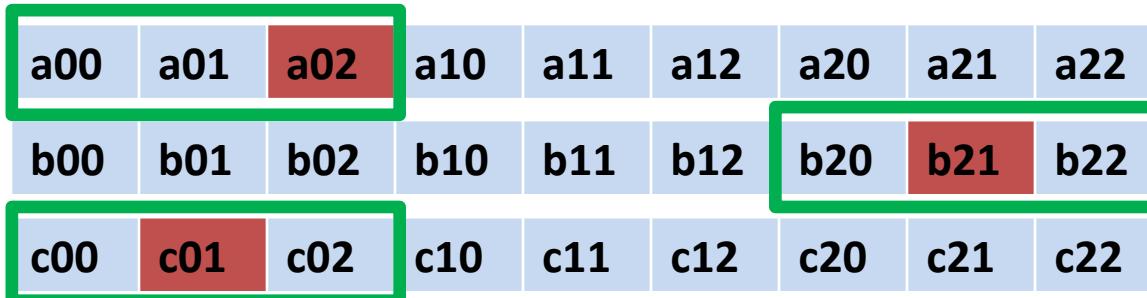


The better way (row-major):

```
for(int i=0; i<n; ++i)
    for(int k=0; k<n; ++k)
        for(int j=0; j<n; ++j)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



Cache Misses: 2



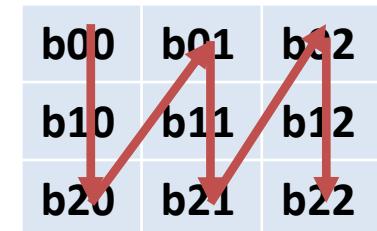
i	k	j
0	2	1

## Full example: matrix multiplication

Sub-optimal memory access will kill the performance...

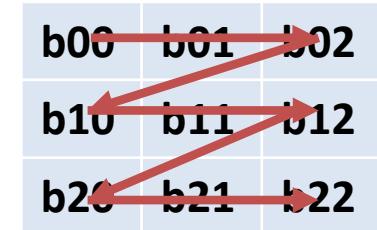
The less efficient way (column-major):

```
for(int i=0; i<n; ++i)
    for(int j=0; j<n; ++j)
        for(int k=0; k<n; ++k)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



The better way (row-major):

```
for(int i=0; i<n; ++i)
    for(int k=0; k<n; ++k)
        for(int j=0; j<n; ++j)
            c[i*n+j] += a[i*n+k] * b[k*n+j];
```



Cache Misses: 2

a00	a01	a02	a10	a11	a12	a20	a21	a22
b00	b01	b02	b10	b11	b12	b20	b21	b22
c00	c01	c02	c10	c11	c12	c20	c21	c22

i	k	j
0	2	2

## Full example: matrix multiplication with OpenMP / Numba

```
#pragma omp parallel for
for(int i=0; i<size; ++i)
    for(int k=0; k<size; ++k)
        for(int j=0; j<size; ++j)
            c[i*size+j] += a[i*size+k] * b[k*size+j];
```

C++

```
!$omp parallel private(i,j,k)
!$omp do schedule(static)
do i=0,size-1
    do k=0,size-1
        do j=0,size-1
            c(i*size+j) = c(i*size+j) + a(i*size+k) * b(k*size+j)
        enddo
    enddo
enddo
 !$omp end do
 !$omp end parallel
```

FORTRAN

```
@jit(Parallel=True)
def mat_mul(a,b,c,size):
    for i in prange(size):
        for k in range(size):
            for j in range(size):
                c[i*size+j] += a[i*size+k] + b[k*size+j]
```

Python

## Full example: matrix multiplication with GPU offloading (OpenMP 4.5+)

```
!$omp target map(tofrom:c) map(to:a,b)
!$omp teams distribute parallel do private(i,j,k) collapse(2)
do i=0,size-1
    do k=0,size-1
        do j=0,size-1
            c(i*size+j) = c(i*size+j) + a(i*size+k) * b(k*size+j)
        enddo
    enddo
enddo
 !$omp end target
```

FORTRAN

```
#pragma omp target map(to: a[0:n2], b[0:n2]), map(tofrom: c[0:n2])
#pragma omp teams distribute parallel for collapse(2)
for(int i=0; i<size; ++i)
    for(int j=0; j<size; ++j)
        for (int k = 0; k < size; k++)
            c[i*size+j] += a[i*size+k] * b[k*size+j];
```

C++

## Full example: matrix multiplication with CUDA

```
__global__ void cuda_mul(float* a, float* b, float* c, int size) {  
    int row = blockIdx.y*blockDim.y+threadIdx.y;  
    int col = blockIdx.x*blockDim.x+threadIdx.x;  
    for (int i = 0; i < size; i++)  
        c[row*size+col] += a[row*size+i] * b[i*size+col];  
}
```

CUDA Kernel

```
cudaMemcpy(dm1, a, sizeof(float)*size*size, cudaMemcpyHostToDevice);  
cudaMemcpy(dm2, b, sizeof(float)*size*size, cudaMemcpyHostToDevice);  
cudaMemcpy(dm3, c, sizeof(float)*size*size, cudaMemcpyHostToDevice);
```

```
dim3 blockSize = dim3(16, 16);  
dim3 gridSize = dim3(size / blockSize.x, size/ blockSize.y);
```

```
cuda_mul<<<gridSize, blockSize>>>(dm1, dm2, dm3, size);
```

```
cudaMemcpy(c, dm3, sizeof(float)*size*size, cudaMemcpyDeviceToHost);
```

Kernel Call

## Full example: matrix multiplication with CUDA

```
__global__ void cuda_mul(float* a, float* b, float* c, int size) {  
    int row = blockIdx.y*blockDim.y+threadIdx.y;  
    int col = blockIdx.x*blockDim.x+threadIdx.x;  
    for (int i = 0; i < size; i++)  
        c[row*size+col] += a[row*size+i] * b[i*size+col];  
}
```

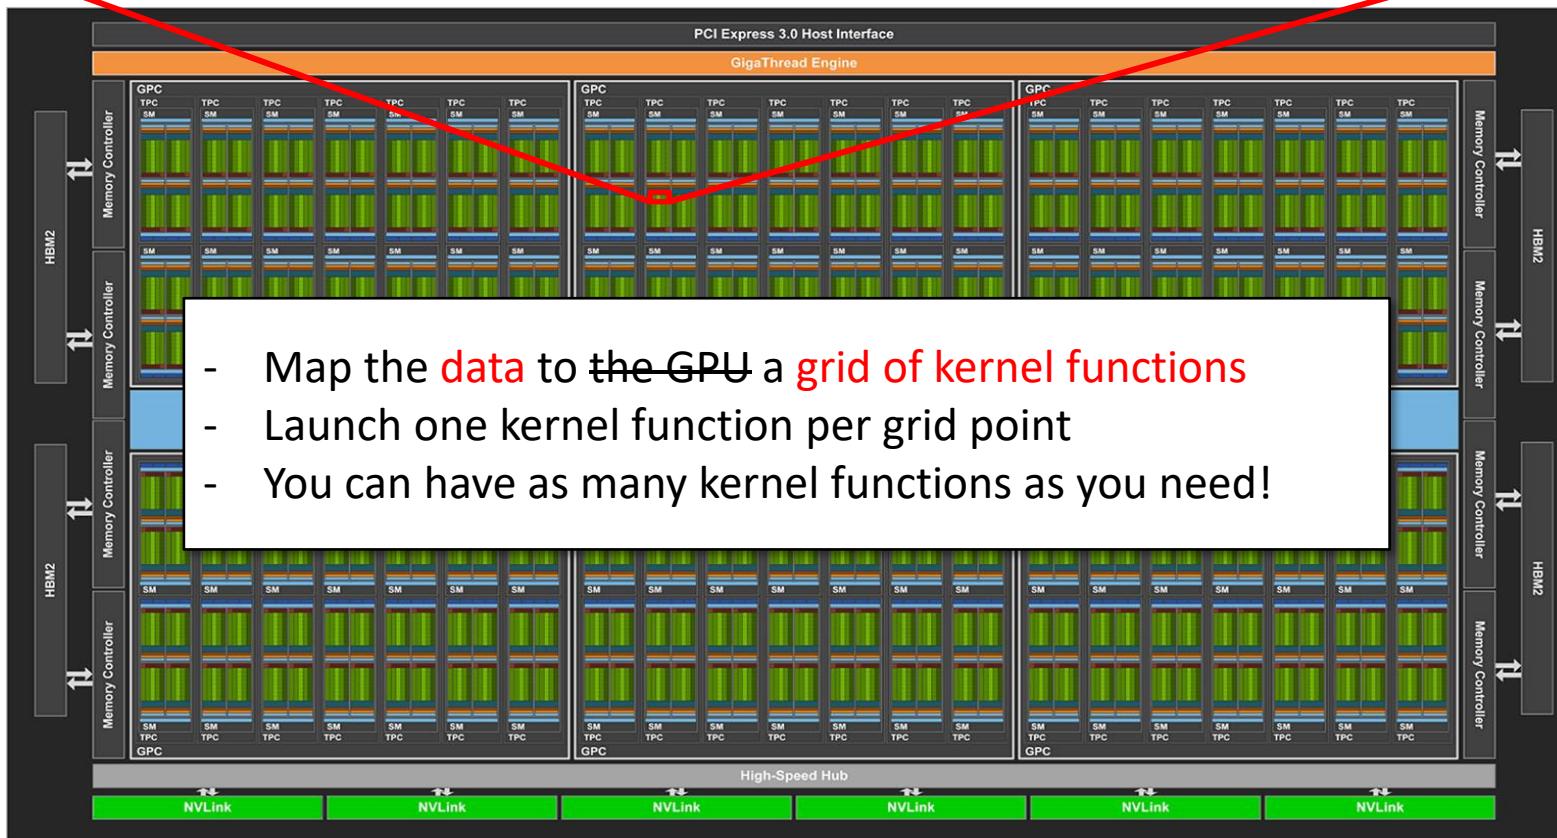
CUDA Kernel



## Full example: matrix multiplication with CUDA

```
__global__ void cuda_mul(float* a, float* b, float* c, int size) {  
    int row = blockIdx.y*blockDim.y+threadIdx.y;  
    int col = blockIdx.x*blockDim.x+threadIdx.x;  
    for (int i = 0; i < size; i++)  
        c[row*size+col] += a[row*size+i] * b[i*size+col];  
}
```

CUDA Kernel



## Full example: matrix multiplication with CUDA and Numba

```
@cuda.jit
def mat_mul(a,b,c,size):
    row=cuda.blockIdx.y*cuda.blockDim.y+cuda.threadIdx.y
    col=cuda.blockIdx.x*cuda.blockDim.x+cuda.threadIdx.x
    for i in range(size):
        c[row*size+col] += a[row*size+i] * b[i*size+col]
```

CUDA Kernel

```
threadsperblock = (16,16)
blockspergrid_x = int(np.ceil(size_array / threadsperblock[0]))
blockspergrid_y = int(np.ceil(size_array / threadsperblock[1]))
blockspergrid = (blockspergrid_x, blockspergrid_y)

mat_mul[blockspergrid, threadsperblock](a,b,c,size_array)
```

Kernel Call

## Full example: matrix multiplication with CUDA and Numba

```
@cuda.jit  
def mat_mul(a,b,c,size):  
    row=cuda.blockIdx.y*cuda.blockDim.y+cuda.threadIdx.y  
    col=cuda.blockIdx.x*cuda.blockDim.x+cuda.threadIdx.x  
    for i in range(size):  
        c[row*size+col] += a[row*size+i] * b[i*size+col]
```

CUDA Kernel

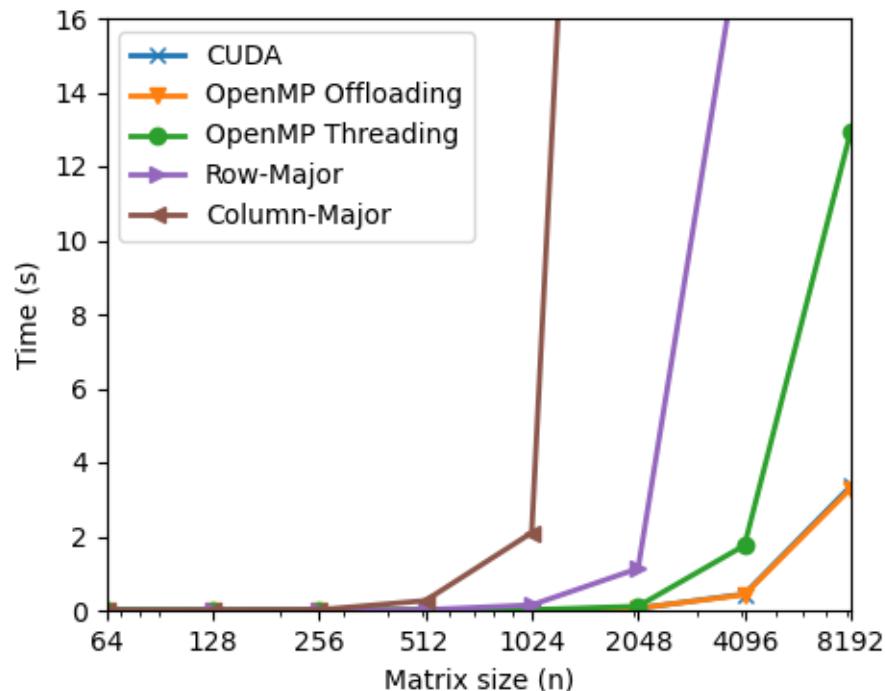
```
threadsperblock = (16,16)  
blockspergrid_x = np.ceil(size_array / threadsperblock[0]))  
blockspergrid_y = np.ceil(size_array / threadsperblock[1]))  
blockspergrid = (blockspergrid_x, blockspergrid_y)  
  
mat_mul[blockspergrid, threadsperblock](a,b,c,size_array)
```

Kernel Call

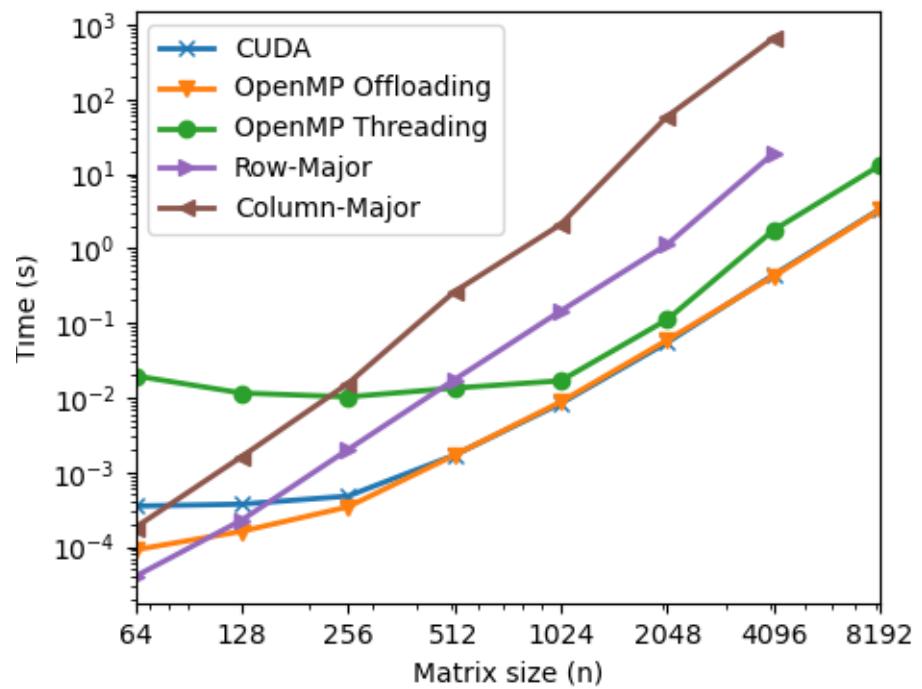
# Matrix Multiplication Demo

## Full example: matrix multiplication performance

Linear / log scale



Log / log scale



N = 4096

Column-Major: 658s

OpenMP 4.5/CUDA GPU: 0.43s

=> 1500x performance speedup!

There are two main ways of parallelizing a code



Open Multi-Processing

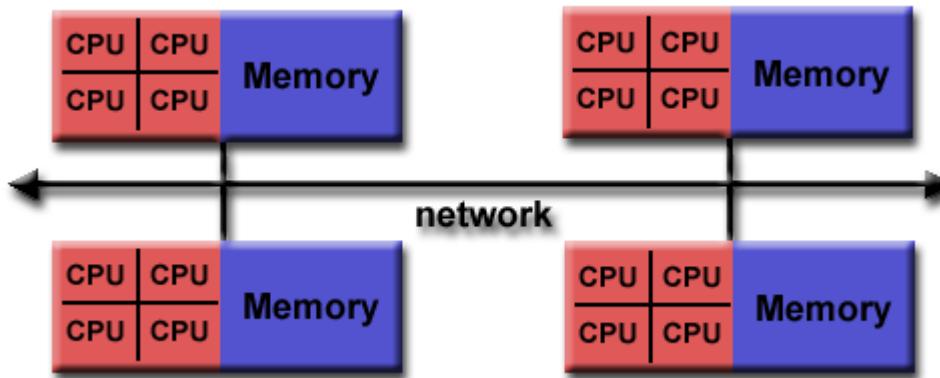


Message Passing Interface

Both are APIs for C/C++ and Fortran

Works in **shared** memory systems  
Can use as many cores as there are in the **compute node**  
Uses **threads** to split the work  
Threads have both **private and shared** variables  
Uses **preprocessor directives** in commented lines

Works in **shared and distributed** memory systems  
Can use as many cores as there are in the **cluster**  
Uses **processes** to split the work  
Processes have only **private** variables  
Loads a library and uses **specific commands**



There are two main ways of parallelizing a code



Open Multi-Processing

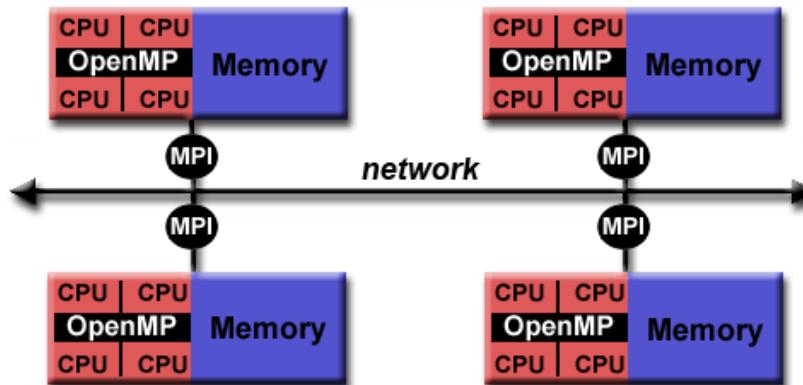


Message Passing Interface

Both are APIs for C/C++ and Fortran

Works in **shared** memory systems  
Can use as many cores as there are in the **compute node**  
Uses **threads** to split the work  
Threads have both **private and shared** variables  
Uses **preprocessor directives** in commented lines

Works in **shared and distributed** memory systems  
Can use as many cores as there are in the **cluster**  
Uses **processes** to split the work  
Processes have only **private** variables  
Loads a library and uses **specific commands**



## What about MPI?

Loads a library and uses **specific commands**

```
include 'mpif.h'

<variable declarations>

call MPI_INIT(ierr)
call MPI_COMM_RANK(MPI_COMM_WORLD,id,ierr)
call MPI_COMM_SIZE(MPI_COMM_WORLD,nproc,ierr)

<calculations>

call MPI_FINALIZE(ierr)
```



## What about MPI?

Loads a library and uses **specific commands**

```
include 'mpif.h'

<variable declarations>

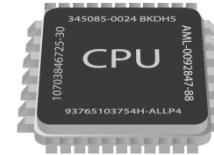
call MPI_INIT(ierr)
call MPI_COMM_RANK(MPI_COMM_WORLD,id,ierr)
call MPI_COMM_SIZE(MPI_COMM_WORLD,nproc,ierr)

do i=1,N
    <calculations>
enddo

call MPI_FINALIZE(ierr)
```

N iterations

N/4



N/4



N/4



N/4



...



## What about MPI?

Loads a library and uses **specific commands**

```
include 'mpif.h'

<variable declarations>

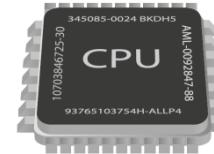
call MPI_INIT(ierr)
call MPI_COMM_RANK(MPI_COMM_WORLD,id,ierr)
call MPI_COMM_SIZE(MPI_COMM_WORLD,nproc,ierr)

do i=ini,fin
    <calculations>
enddo

call MPI_FINALIZE(ierr)
```

N iterations

$N/nproc \rightarrow$



0

$N/nproc \rightarrow$



1

$N/nproc \rightarrow$



2

$N/nproc \rightarrow$



3

...

$N/nproc \rightarrow$



**nproc-1**



OPEN MPI

## What about MPI?

Loads a library and uses **specific commands**

```
include 'mpif.h'

<variable declarations>

call MPI_INIT(ierr)
call MPI_COMM_RANK(MPI_COMM_WORLD,id,ierr)
call MPI_COMM_SIZE(MPI_COMM_WORLD,nproc,ierr)

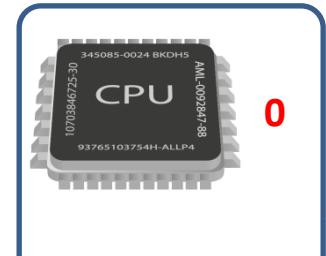
ini=id*(N/nproc)+1
fin=(id+1)*(N/nproc)

do i=ini,fin
    <calculations>
enddo

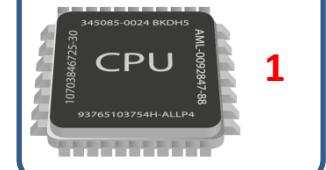
call MPI_FINALIZE(ierr)
```

N iterations

N/nproc →



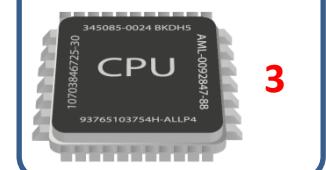
N/nproc →



N/nproc →

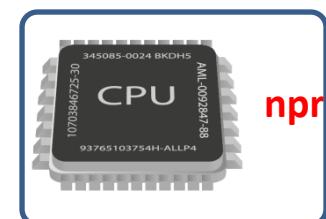


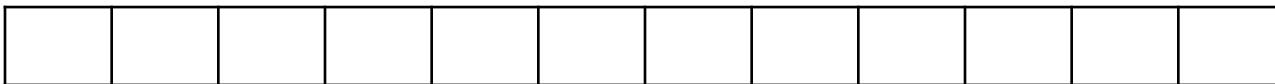
N/nproc →

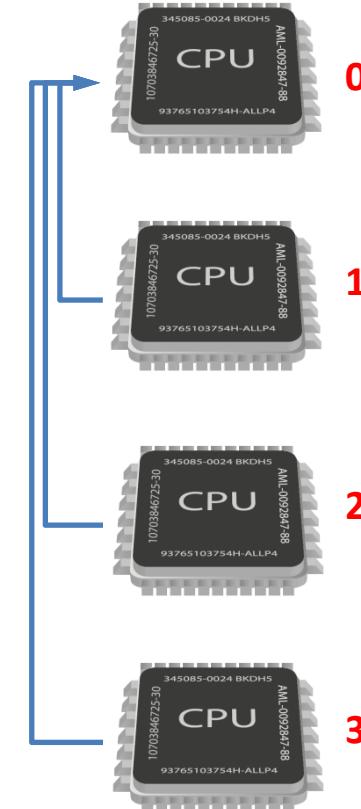


...

N/nproc →

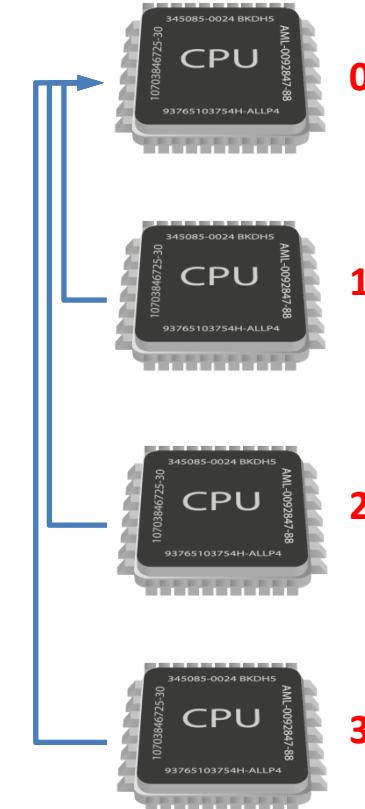
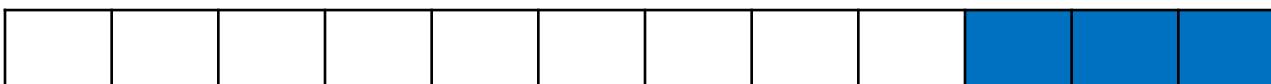






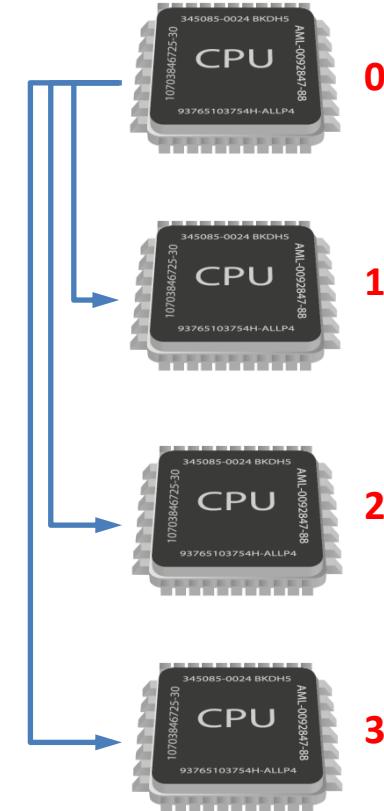
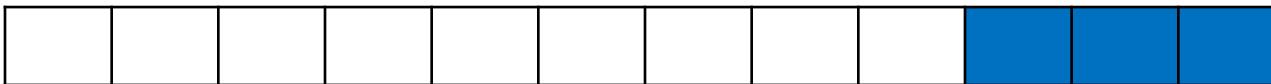
OPEN MPI

Only process 0 has the correct values!



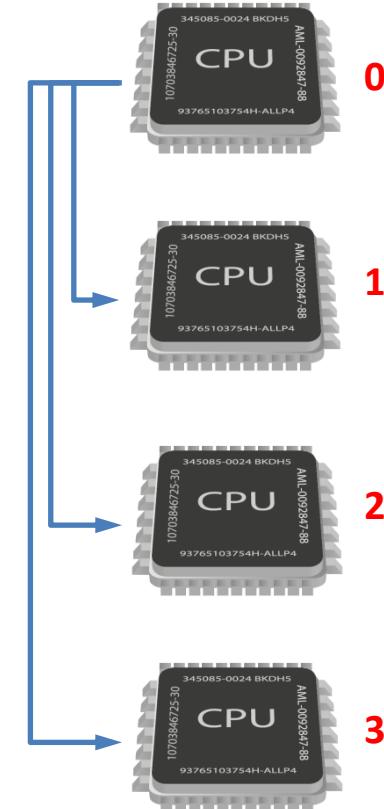
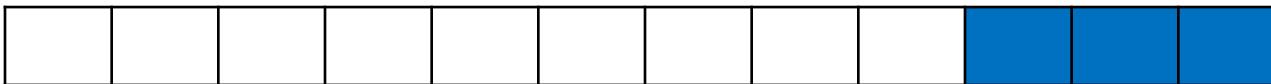
```
call MPI_REDUCE(arr,result,arraysize,<type>,<operation>,<dest_id>,MPI_WORLD_COMM,ierr)  
call MPI_REDUCE(arr,result,12,MPI_INT,MPI_SUM,0,MPI_WORLD_COMM,ierr)
```

Only process 0 has the correct values!



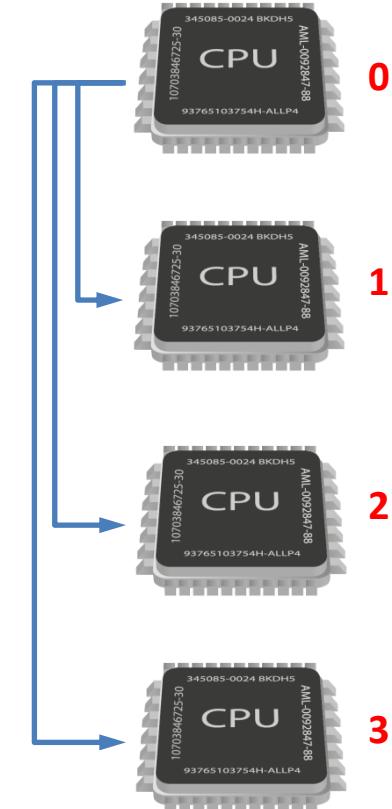
```
call MPI_REDUCE(arr,result,arraysize,<type>,<operation>,<dest_id>,MPI_WORLD_COMM,ierr)
call MPI_REDUCE(arr,result,12,MPI_INT,MPI_SUM,0,MPI_WORLD_COMM,ierr)
```

Only process 0 has the correct values!

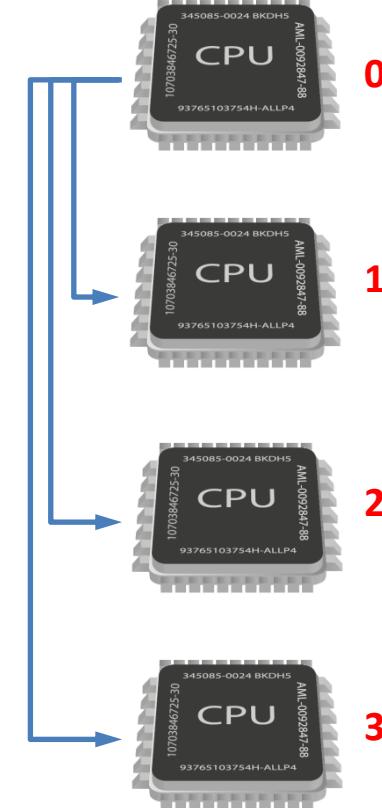


```
call MPI_REDUCE(arr,result,arraysize,<type>,<operation>,<dest_id>,MPI_WORLD_COMM,ierr)
call MPI_REDUCE(arr,result,12,MPI_INT,MPI_SUM,0,MPI_WORLD_COMM,ierr)
call MPI_BCAST(result,12,MPI_INT,0,MPI_WORLD_COMM,ierr)
```

Only process 0 has the correct values!



```
call MPI_REDUCE(arr,result,arraysize,<type>,<operation>,<dest_id>,MPI_WORLD_COMM,ierr)
call MPI_REDUCE(arr,result,12,MPI_INT,MPI_SUM,0,MPI_WORLD_COMM,ierr)
call MPI_BCAST(result,12,MPI_INT,0,MPI_WORLD_COMM,ierr)
```



```
call MPI_ALLREDUCE(arr,result,12,MPI_INT,MPI_SUM,MPI_WORLD_COMM,ierr)
```



OPEN MPI

## Example: Numerical integration

```
include 'mpif.h'

call MPI_INIT(ierr); comm = MPI_COMM_WORLD
call MPI_COMM_RANK(comm, rank, ierr)
call MPI_COMM_SIZE(comm, size, ierr)
```

Initialize MPI

FORTRAN



OPEN MPI

## Example: Numerical integration

```
include 'mpif.h'

call MPI_INIT(ierr); comm = MPI_COMM_WORLD
call MPI_COMM_RANK(comm, rank, ierr)
call MPI_COMM_SIZE(comm, size, ierr)

totalsteps = 1000000
steps = totalsteps/size

ini = rank * steps
fin = (rank+1) * steps

if(fin.ge.totalsteps) then
    fin = totalsteps
endif
```

Initialize MPI

Define the range

**FORTRAN**



OPEN MPI

## Example: Numerical integration

```
include 'mpif.h'

call MPI_INIT(ierr); comm = MPI_COMM_WORLD
call MPI_COMM_RANK(comm, rank, ierr)
call MPI_COMM_SIZE(comm, size, ierr)

totalsteps = 1000000
steps = totalsteps/size

ini = rank * steps
fin = (rank+1) * steps

if(fin.ge.totalsteps) then
    fin = totalsteps
endif

dx = 1.d0 / totalsteps
x = -0.5d0 * dx

localSum = 0.d0
do i = ini, fin
    x = (i-0.5d0)*dx
    localSum = localSum + 4.d0 / (1.d0 + x*x)
enddo
localSum = localSum * dx
```

Initialize MPI

Define the range

**FORTRAN**

Do the computation

## Example: Numerical integration

```
include 'mpif.h'

call MPI_INIT(ierr); comm = MPI_COMM_WORLD
call MPI_COMM_RANK(comm, rank, ierr)
call MPI_COMM_SIZE(comm, size, ierr)

totalsteps = 1000000
steps = totalsteps/size

ini = rank * steps
fin = (rank+1) * steps

if(fin.ge.totalsteps) then
    fin = totalsteps
endif

dx = 1.d0 / totalsteps
x = -0.5d0 * dx

localSum = 0.d0
do i = ini, fin
    x = (i-0.5d0)*dx
    localSum = localSum + 4.d0 / (1.d0 + x*x)
enddo
localSum = localSum * dx

call MPI_ALLREDUCE(localSum,globalSum,1,MPI_DOUBLE_PRECISION,MPI_SUM,comm,ierr)

write(*,*) localSum, globalSum
```

Initialize MPI

Define the range

**FORTRAN**

Do the computation

Sum the data across nodes

## Example: Numerical integration

```
include 'mpif.h'

call MPI_INIT(ierr); comm = MPI_COMM_WORLD
call MPI_COMM_RANK(comm, rank, ierr)
call MPI_COMM_SIZE(comm, size, ierr)

totalsteps = 1000000
steps = totalsteps/size

ini = rank * steps
fin = (rank+1) * steps

if(fin.ge.totalsteps) then
    fin = totalsteps
endif

dx = 1.d0 / totalsteps
x = -0.5d0 * dx

localSum = 0.d0
do i = ini, fin
    x = (i-0.5d0)*dx
    localSum = localSum + 4.d0 / (1.d0 + x*x)
enddo
localSum = localSum * dx

call MPI_ALLREDUCE(localSum,globalSum,1,MPI_DOUBLE_PRECISION,MPI_SUM,comm,ierr)

write(*,*) localSum, globalSum

call MPI_FINALIZE(ierr)
```

Initialize MPI

Define the range

**FORTRAN**

Do the computation

Sum the data across nodes

Don't forget to Finalize MPI

## Example: Numerical integration

```
from mpi4py import MPI
import numpy

comm = MPI.COMM_WORLD
rank = comm.Get_rank()
size = comm.Get_size()

totalsteps = 1000000
steps = totalsteps/size

ini = rank * steps
fin = (rank+1) * steps

if fin > totalsteps:
    fin = totalsteps

dx = 1./totalsteps
x = -0.5*dx

localSum = 0.
for i in range(int(ini), int(fin)):
    x = (i-0.5)*dx
    localSum = localSum + 4./(1. + x*x)

localSum = localSum * dx

globalSum = comm.allreduce(localSum, op=MPI.SUM)

print(localSum, globalSum)
```

Initialize MPI

Define the range

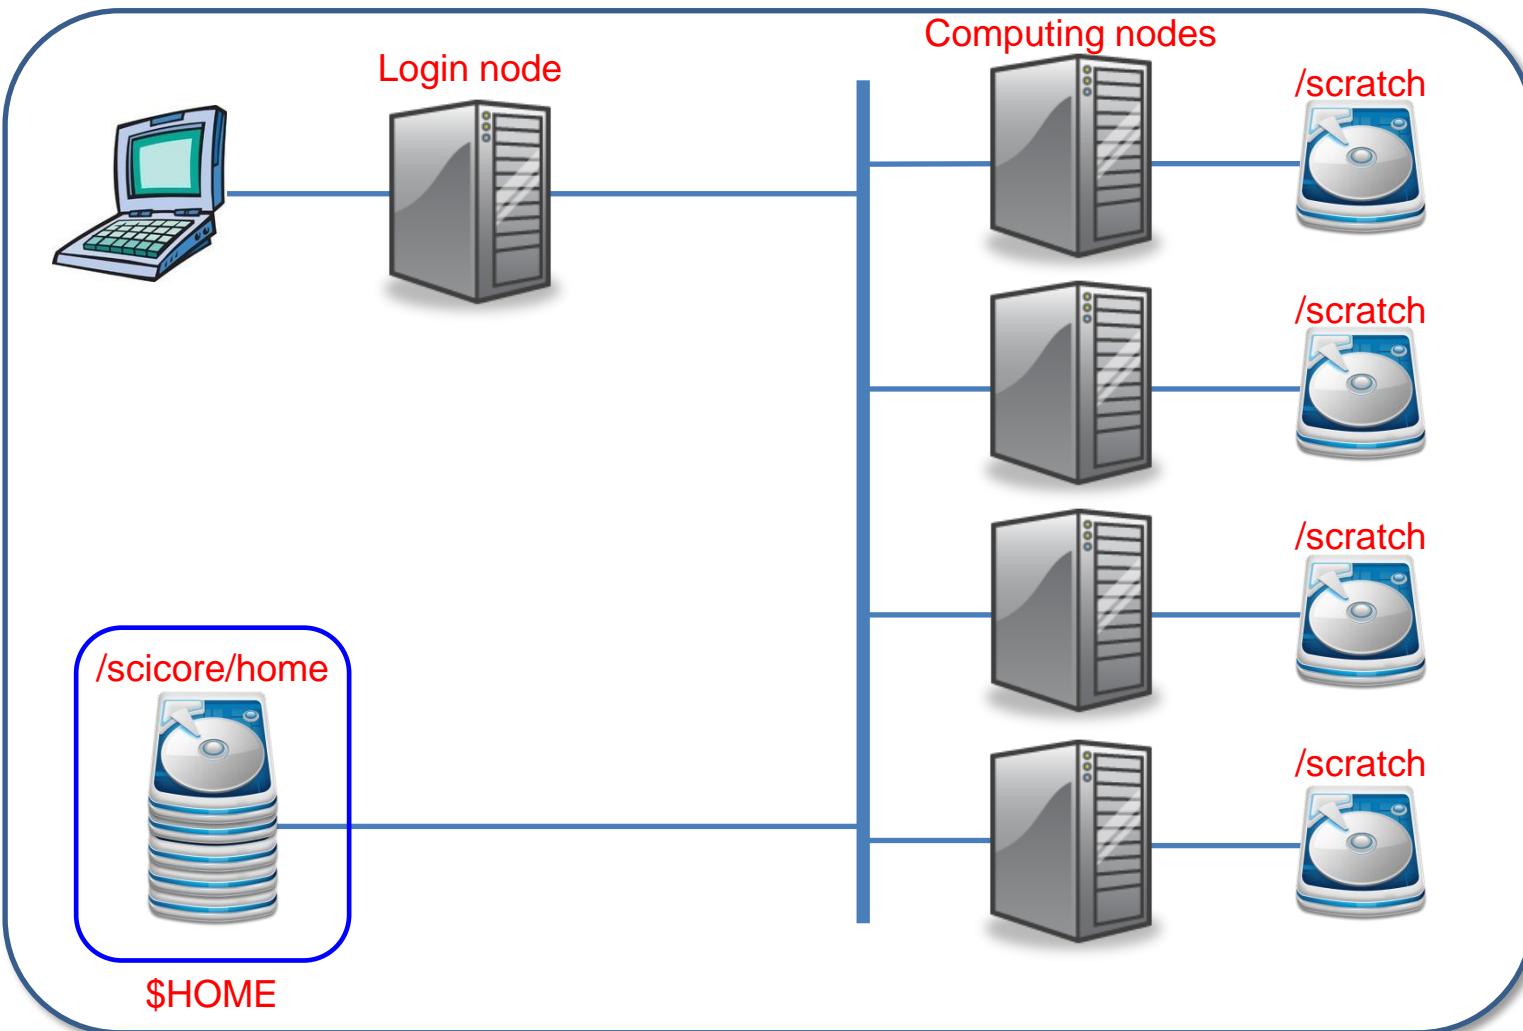
Python

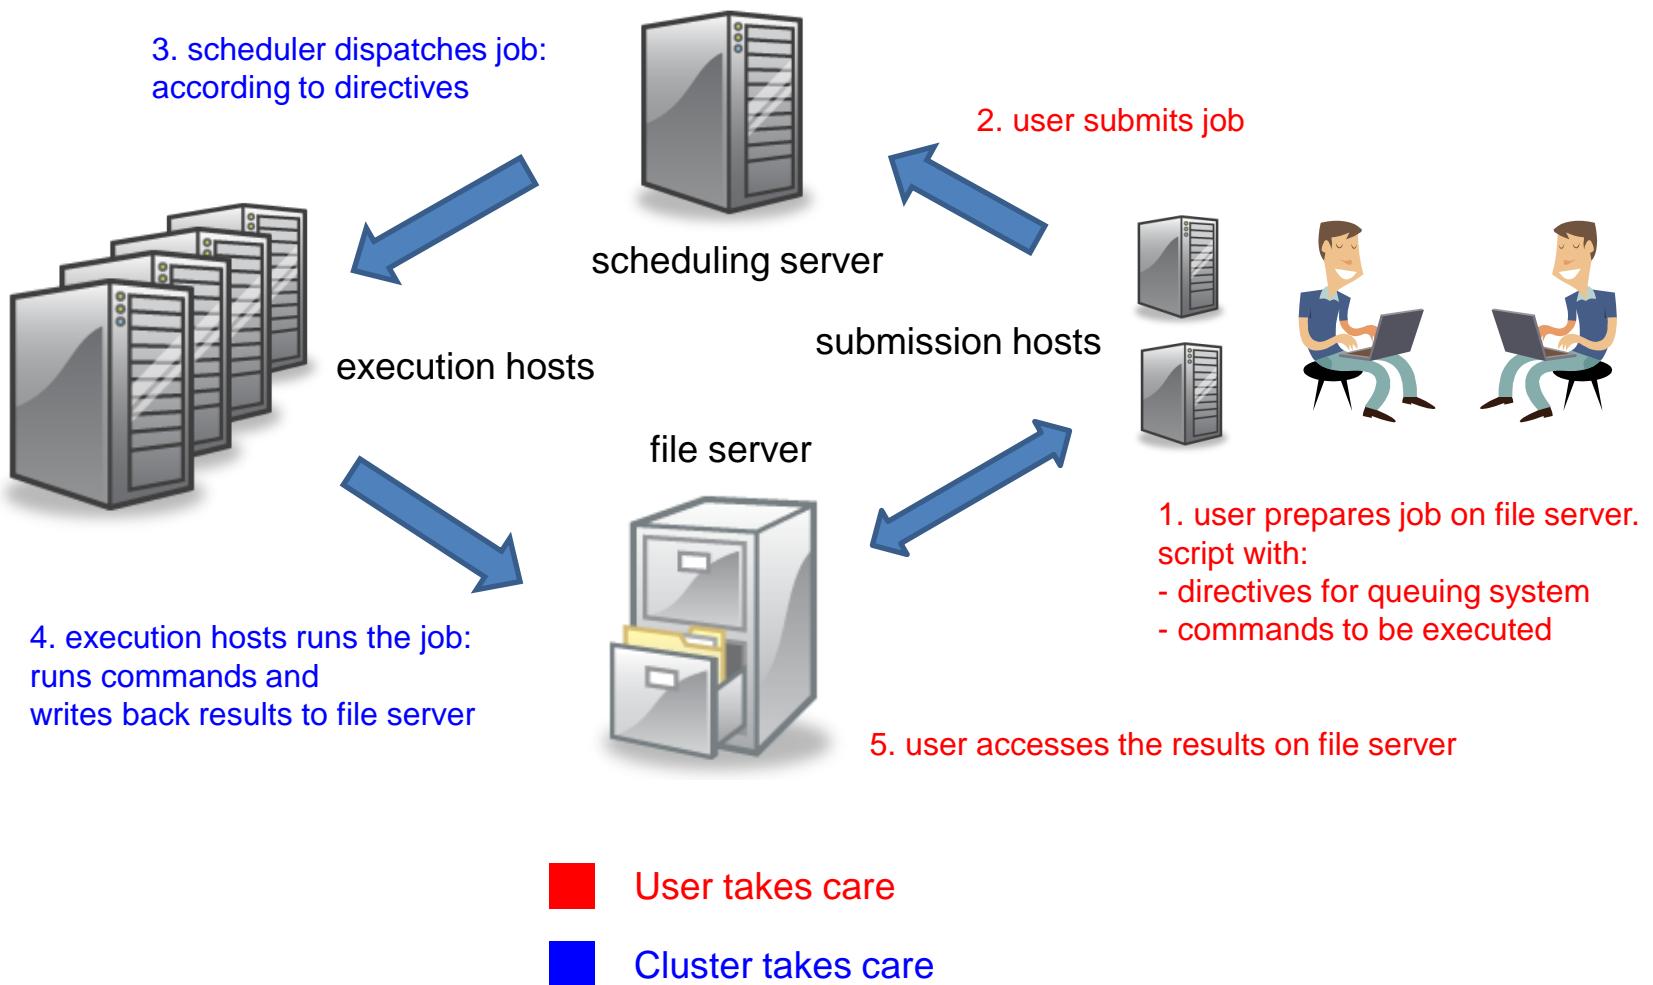
Do the computation

Sum the data across nodes



Mpi4py tutorial at: <https://mpi4py.readthedocs.io/en/stable/tutorial.html>





Lest's play "Spot the Differences"

## Lest's play "Spot the Differences"

Which script launches an OpenMP job and which one an MPI?

```
#!/bin/bash

#SBATCH --job-name=myJob
#SBATCH --cpus-per-task=8
#SBATCH --mem=3.5G
#SBATCH --time=05:00:00
#SBATCH --qos=6hour
#SBATCH --output=/path/to/stdout/folder
#SBATCH --error=/path/to/stderr/folder
#SBATCH --mail-type=END,FAIL,TIME_LIMIT
#SBATCH --mail-user=mailaddress@unibas.ch

# Tell OpenMP how many threads to use
#####
export OMP_NUM_THREADS=$SLURM_CPUS_PER_TASK

# load your required modules
#####
module load <module name>
...

# and here goes your command line
./my_openmp_program

# Some third party programs include a threading
# option, configuration or environment variable
```

```
#!/bin/bash

#SBATCH --job-name=myJob
#SBATCH --ntasks=8
#SBATCH --mem=3.5G
#SBATCH --time=05:00:00
#SBATCH --qos=6hour
#SBATCH --output=/path/to/stdout/folder
#SBATCH --error=/path/to/stderr/folder
#SBATCH --mail-type=END,FAIL,TIME_LIMIT
#SBATCH --mail-user=mailaddress@unibas.ch

# Tell OpenMP how many threads to use
#####
export OMP_NUM_THREADS=$SLURM_CPUS_PER_TASK

# load your required modules
#####
module load <module name>
...

# and here goes your command line
srun ./my_openmp_program

# Some third party programs include a threading
# option, configuration or environment variable
```

```
#!/bin/bash                                         How do I launch

#SBATCH --job-name=myJob
#SBATCH --ntasks=?
#SBATCH --tasks-per-node=?
#SBATCH --cpus-per-task=?
#SBATCH --mem=3.5G
#SBATCH --time=05:00:00
#SBATCH --qos=6hour
#SBATCH --output=/path/to/stdout/folder
#SBATCH --error=/path/to/stderr/folder
#SBATCH --mail-type=END,FAIL,TIME_LIMIT
#SBATCH --mail-user=mailaddress@unibas.ch

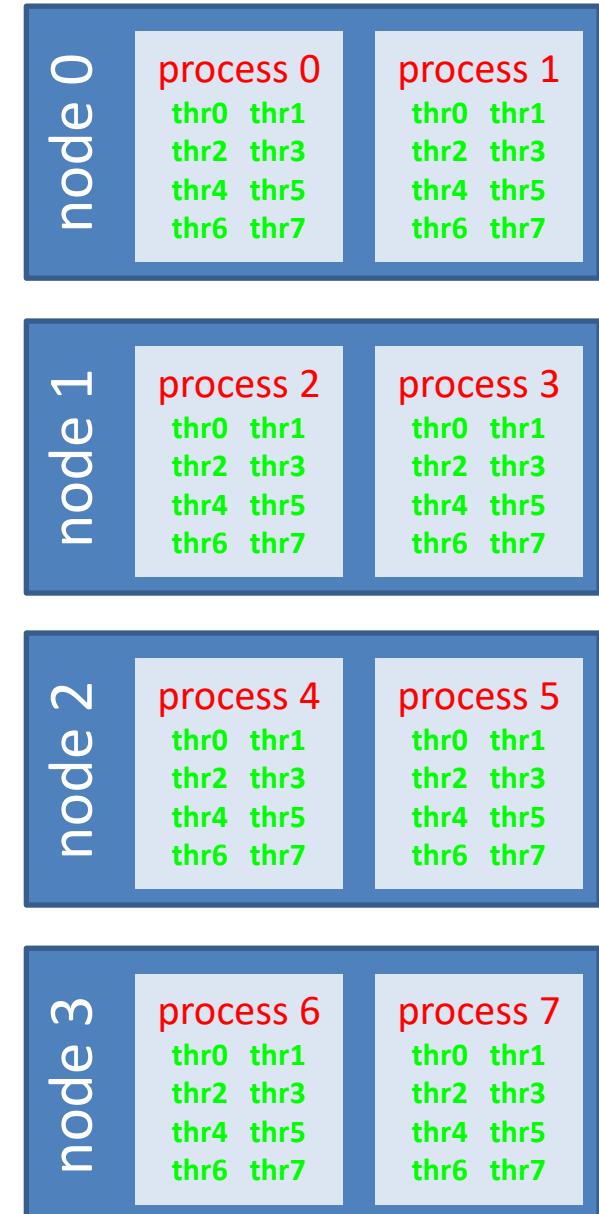
# Tell OpenMP how many threads to use
#####
export OMP_NUM_THREADS=$SLURM_CPUS_PER_TASK

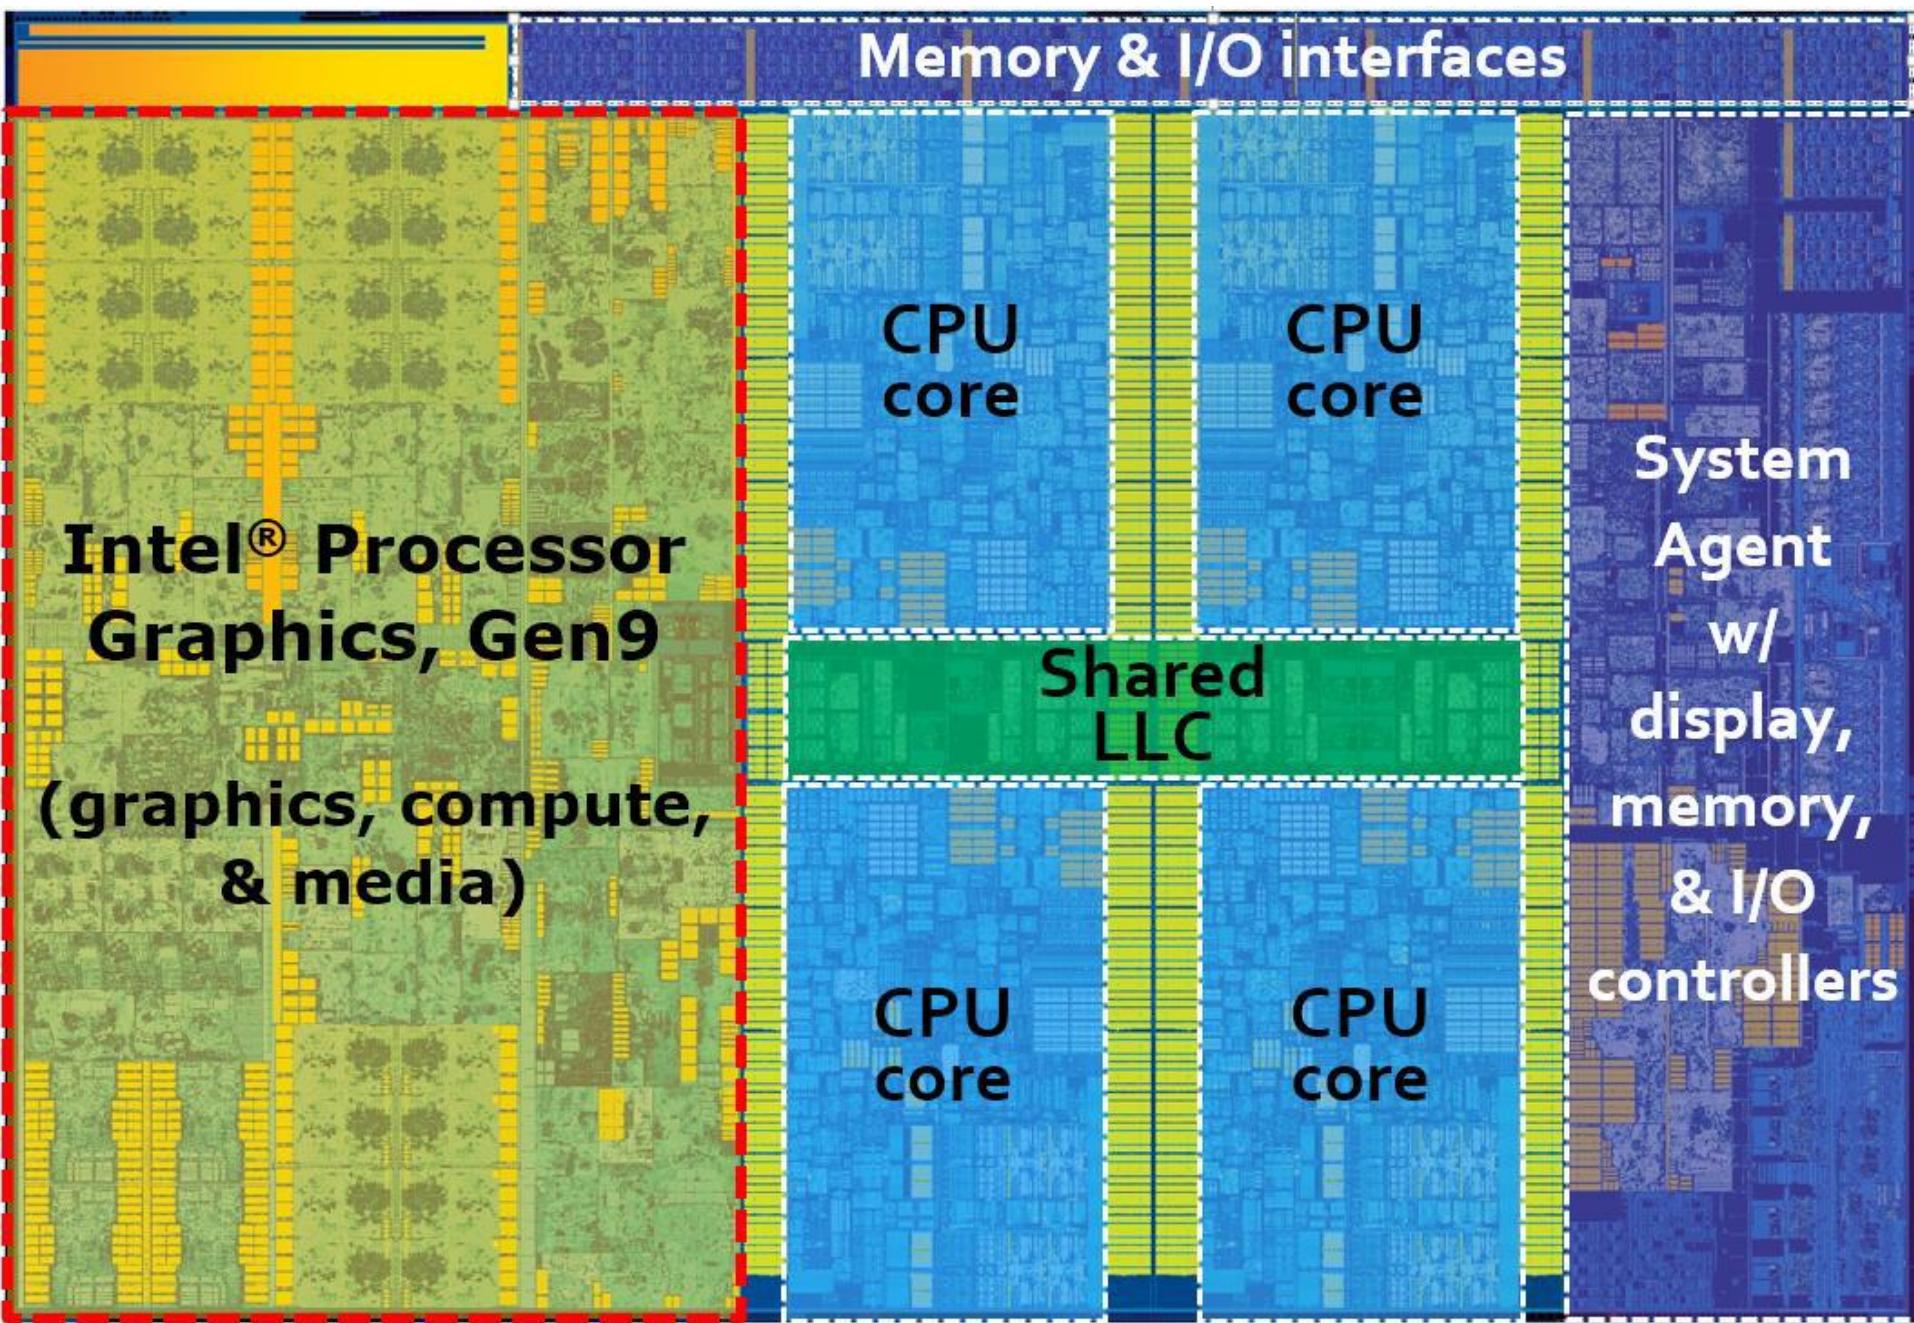
# load your required modules
#####
module load <module name>
...

# and here goes your command line
srun my_openmp_program

# Some third party programs include a threading
# option, configuration or environment variable
```

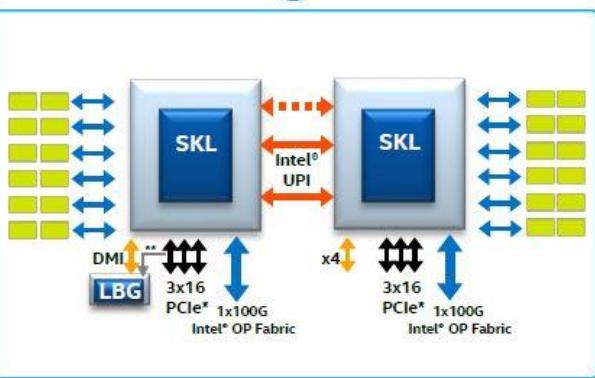
# How do I launch an hybrid job?





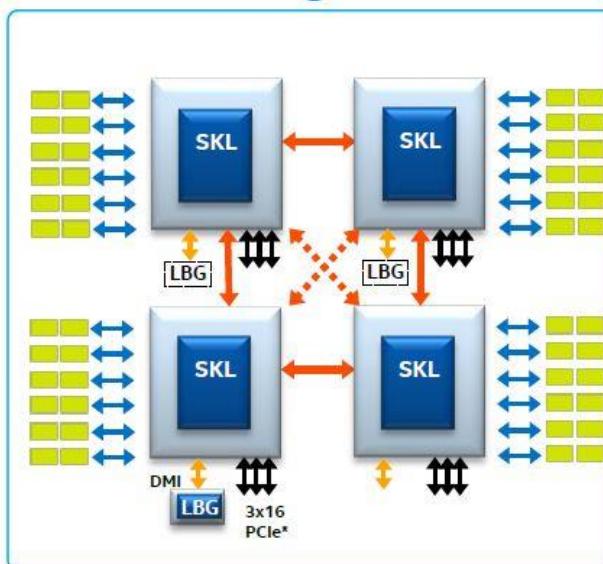
# Platform Topologies

2S Configurations



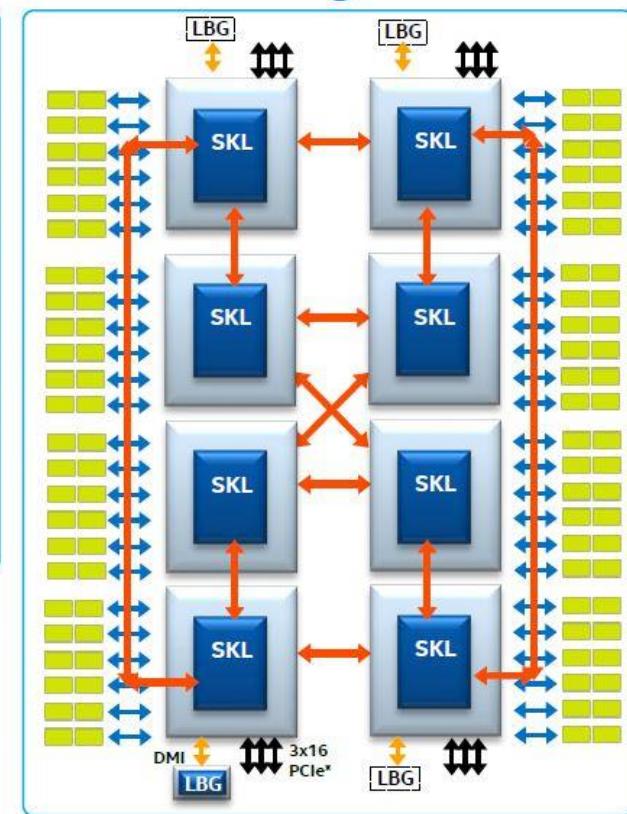
(2S-2UPI & 2S-3UPI shown)

4S Configurations



(4S-2UPI & 4S-3UPI shown)

8S Configuration



INTEL® XEON® SCALABLE PROCESSOR SUPPORTS CONFIGURATIONS RANGING FROM 2S-2UPI TO 8S

When using parallel programs you must be more aware of the hardware topology!

## numactl --hardware

```
[cabezón@login10 ~]$ numactl --hardware
available: 2 nodes (0-1)
node 0 cpus: 0 1 2 3 4 5 6 7
node 0 size: 32742 MB
node 0 free: 188 MB
node 1 cpus: 8 9 10 11 12 13 14 15
node 1 size: 32768 MB
node 1 free: 10518 MB
node distances:
node    0    1
  0: 10 11 ← Relative distances
  1: 11 10
```

## lscpu

```
[cabezón@login10 ~]$ lscpu
Architecture:          x86_64
CPU op-mode(s):        32-bit, 64-bit
Byte Order:            Little Endian
CPU(s):                16
On-line CPU(s) list:  0-15
Thread(s) per core:   1
Core(s) per socket:   8
Socket(s):             2
NUMA node(s):          2
Vendor ID:             GenuineIntel
CPU family:            6
Model:                 45
Model name:            Intel(R) Xeon(R) CPU E5-2670 0 @ 2.60GHz
Stepping:               7
CPU MHz:                1298.578
BogoMIPS:              5205.23
Virtualization:        VT-x
L1d cache:              32K
L1i cache:              32K
L2 cache:                256K
L3 cache:                20480K
NUMA node0 CPU(s):      0-7
NUMA node1 CPU(s):      8-15
```

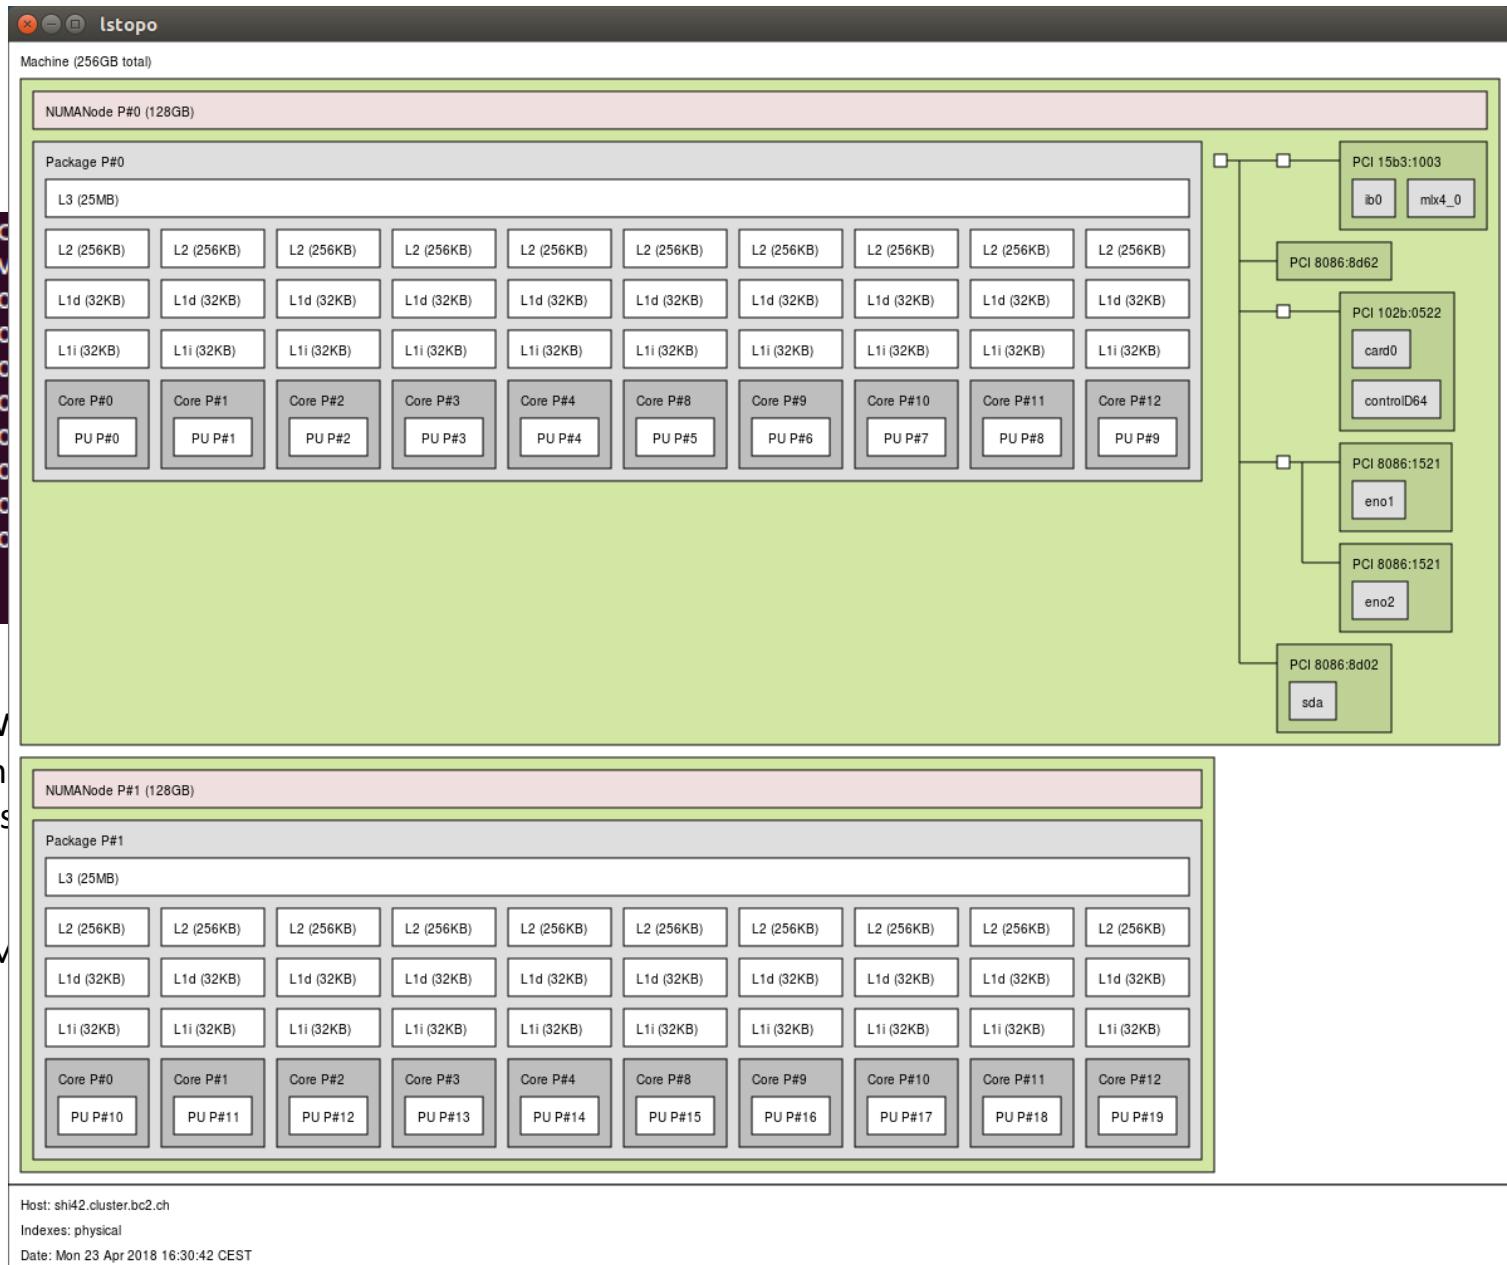
Two sockets!

This is a NUMA (non-uniform memory access) system.

vs.

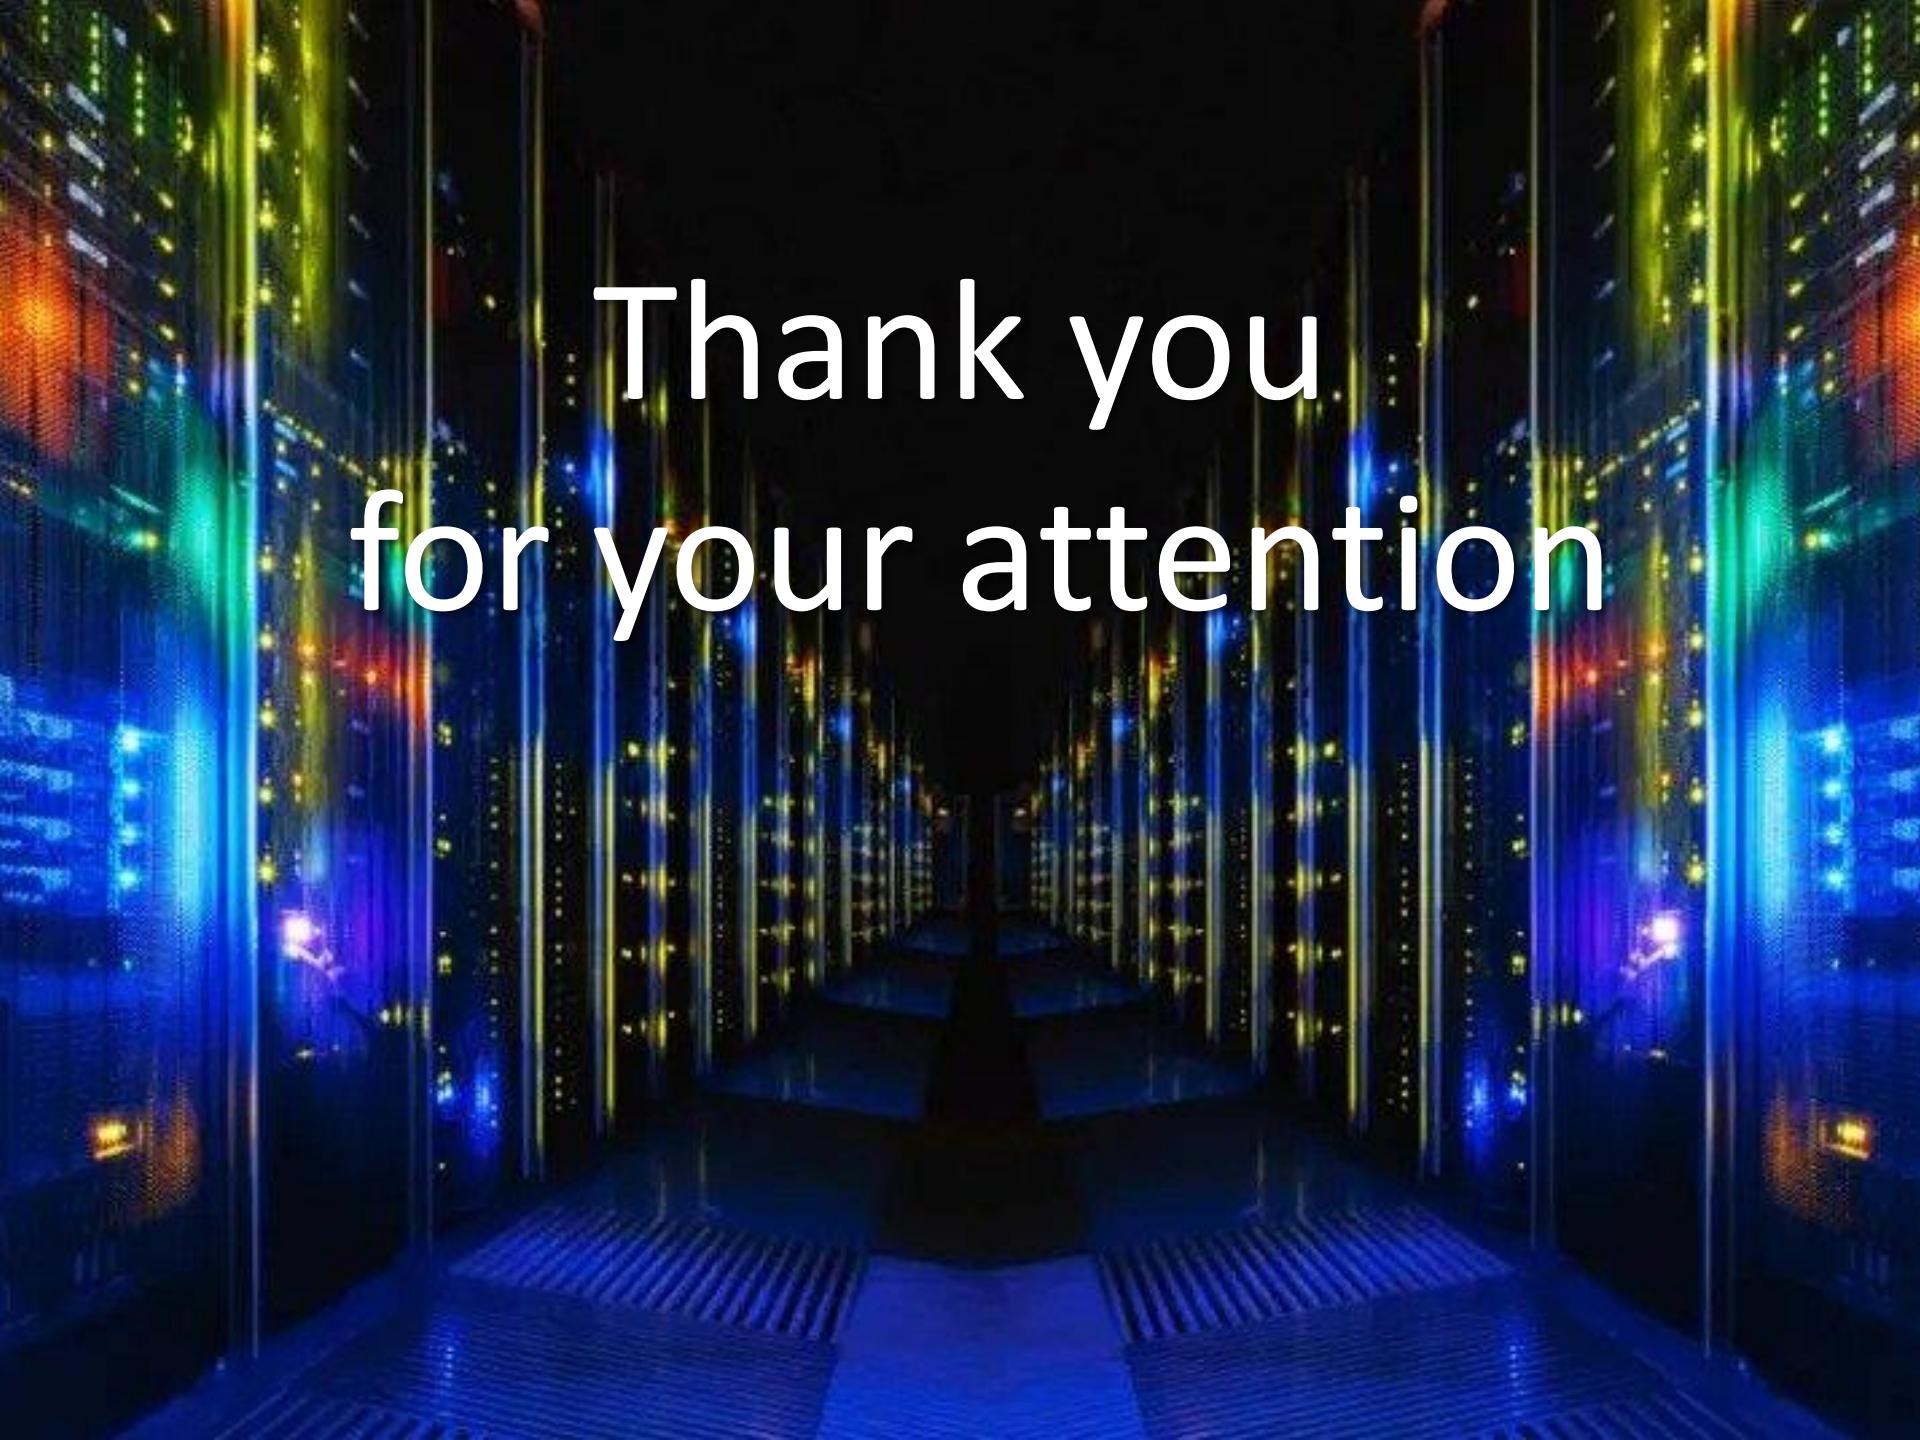
SMP (symmetric multi-processor):

```
ruben@jarvis:~$ numactl --hardware
available: 1 nodes (0)
node 0 cpus: 0 1 2 3 4 5 6 7
node 0 size: 7885 MB
node 0 free: 252 MB
node distances:
node    0
  0: 10
```



**Istopo**  
(from hwloc package)

E5-2670 0 @ 2.60GHz



Thank you  
for your attention



**Resume at:  
13:30**

Coffee  
break

